

# Bottlenecks Uncovered: A Component-Wise Breakdown of the Runtime of an OLTP System

Max Gilbert

Technische Universität Kaiserslautern

*m\_gilbert13@cs.uni-kl.de*

September 30, 2020

# Table of Contents

## Page Eviction by the Buffer Pool Management

- Introduction

- Least Recently Used

- LeanStore Replacement [Lei+]

- All Page Replacement Algorithms

- Conclusion

## Component-Wise Performance Evaluation of OLTP Systems

- OLTP through the Looking Glass, and What We Found There [Har+]

- Results

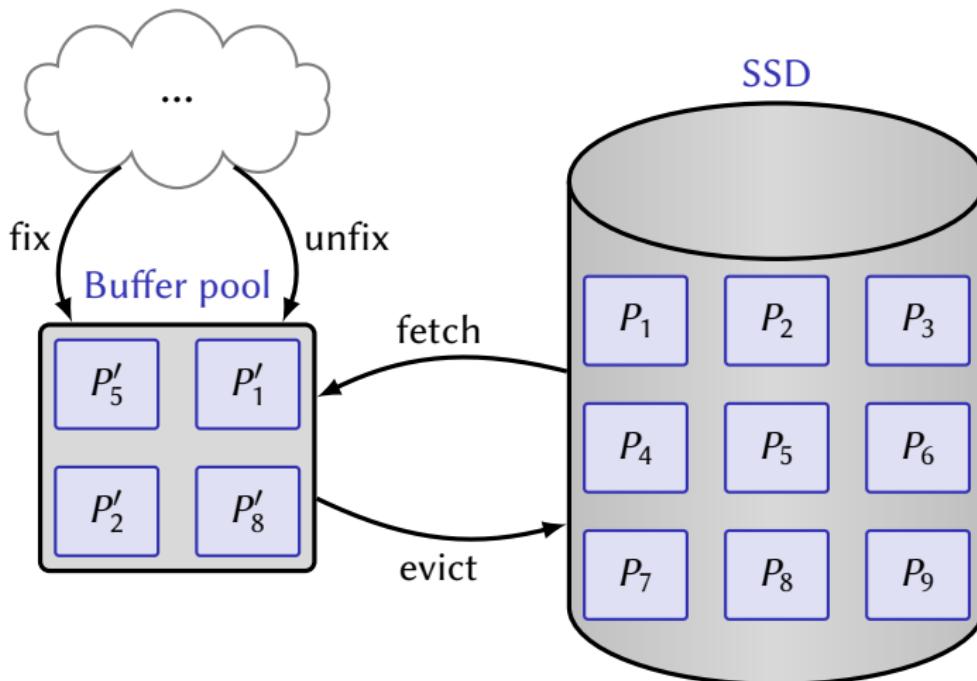
- Conclusion

- Optimizations of OLTP Systems

## Section 1

# Page Eviction by the Buffer Pool Management

# The Buffer Pool



# Extended Page Replacement Algorithm Classification by Effelsberg and Härdter From [EH]

Consideration during selection decision		Age			
		No consideration	Since most recent reference		Since first reference
References	No consideration				
	Most recent reference				
	All references				

# Extended Page Replacement Algorithm Classification by Effelsberg and Härdter From [EH]

Consideration during selection decision		Age			
		No consideration	Since most recent reference	Since some recent reference	Since first reference
References	No consideration				
	Most recent reference				
	All references				

# Extended Page Replacement Algorithm Classification by Effelsberg and Härdter From [EH]

Consideration during selection decision		Age			
		No consideration	Since most recent reference	Since some recent reference	Since first reference
References	No consideration				
	Most recent reference				
	Some recent references				
	All references				

# Extended Page Replacement Algorithm Classification by Effelsberg and Härdter From [EH]

Consideration during selection decision		Age			
		No consideration	Since most recent reference	Since some recent reference	Since first reference
References	No consideration	RANDOM			
	Most recent reference				
	Some recent references				
	All references				

# Extended Page Replacement Algorithm Classification by Effelsberg and Härdter From [EH]

Consideration during selection decision		Age			
		No consideration	Since most recent reference	Since some recent reference	Since first reference
References	No consideration	RANDOM			FIFO
	Most recent reference				
	Some recent references				
	All references				

# Extended Page Replacement Algorithm Classification by Effelsberg and Härdter From [EH]

Consideration during selection decision		Age			
		No consideration	Since most recent reference	Since some recent reference	Since first reference
References	No consideration	RANDOM			FIFO FILO
	Most recent reference				
	Some recent references				
	All references				

# Extended Page Replacement Algorithm Classification by Effelsberg and Härdter From [EH]

Consideration during selection decision		Age			
		No consideration	Since most recent reference	Since some recent reference	Since first reference
References	No consideration	RANDOM			FIFO FILO
	Most recent reference		LRU		
	Some recent references				
	All references				

# Extended Page Replacement Algorithm Classification by Effelsberg and Härdter From [EH]

Consideration during selection decision		Age			
		No consideration	Since most recent reference	Since some recent reference	Since first reference
References	No consideration	RANDOM			FIFO FILO
	Most recent reference		LRU MRU		
	Some recent references				
	All references				

# Extended Page Replacement Algorithm Classification by Effelsberg and Härdter From [EH]

Consideration during selection decision		Age			
		No consideration	Since most recent reference	Since some recent reference	Since first reference
References	No consideration	RANDOM			FIFO FILO
	Most recent reference		LRU MRU		
	Some recent references			LRU-K	
	All references				

# Extended Page Replacement Algorithm Classification by Effelsberg and Härdter From [EH]

Consideration during selection decision		Age			
		No consideration	Since most recent reference	Since some recent reference	Since first reference
References	No consideration	RANDOM			FIFO FILO
	Most recent reference		LRU MRU		
	Some recent references		SLRU	LRU-K	
	All references				

# Extended Page Replacement Algorithm Classification by Effelsberg and Härdter From [EH]

Consideration during selection decision		Age			
		No consideration	Since most recent reference	Since some recent reference	Since first reference
References	No consideration	RANDOM			FIFO FILO
	Most recent reference		LRU MRU CLOCK		
	Some recent references		SLRU	LRU-K	
	All references				

# Extended Page Replacement Algorithm Classification by Effelsberg and Härdter From [EH]

Consideration during selection decision		Age			
		No consideration	Since most recent reference	Since some recent reference	Since first reference
References	No consideration	RANDOM			FIFO FILO
	Most recent reference	ZCLOCK	LRU MRU CLOCK		
	Some recent references		SLRU	LRU-K	
	All references				

# Extended Page Replacement Algorithm Classification by Effelsberg and Härdter From [EH]

Consideration during selection decision		Age			
		No consideration	Since most recent reference	Since some recent reference	Since first reference
References	No consideration	RANDOM			FIFO FILO
	Most recent reference	ZCLOCK	LRU MRU CLOCK		
	Some recent references		SLRU	LRU-K	
	All references		GCLOCK-V1		

# Extended Page Replacement Algorithm Classification by Effelsberg and Härdter From [EH]

Consideration during selection decision		Age			
		No consideration	Since most recent reference	Since some recent reference	Since first reference
References	No consideration	RANDOM			FIFO FILO
	Most recent reference	ZCLOCK	LRU MRU CLOCK GCLOCK-V2		
	Some recent references		SLRU	LRU-K	
	All references		GCLOCK-V1		

# Extended Page Replacement Algorithm Classification by Effelsberg and Härdter From [EH]

Consideration during selection decision		Age			
		No consideration	Since most recent reference	Since some recent reference	Since first reference
References	No consideration	RANDOM			FIFO FILO
	Most recent reference	ZCLOCK	LRU MRU CLOCK GCLOCK-V2		
	Some recent references		SLRU	LRU-K	
	All references		GCLOCK-V1 DGCLOCK-V1		

# Extended Page Replacement Algorithm Classification by Effelsberg and Härdter From [EH]

Consideration during selection decision		Age			
		No consideration	Since most recent reference	Since some recent reference	Since first reference
References	No consideration	RANDOM			FIFO FILO
	Most recent reference	ZCLOCK	LRU MRU CLOCK GCLOCK-V2 DGCLOCK-V2		
	Some recent references		SLRU	LRU-K	
	All references		GCLOCK-V1 DGCLOCK-V1		

# Extended Page Replacement Algorithm Classification by Effelsberg and Härdter From [EH]

Consideration during selection decision		Age			
		No consideration	Since most recent reference	Since some recent reference	Since first reference
References	No consideration	RANDOM			FIFO FILO
	Most recent reference	ZCLOCK	LRU MRU CLOCK GCLOCK-V2 DGCLOCK-V2		
	Some recent references		SLRU	LRU-K	
	All references	LFU	GCLOCK-V1 DGCLOCK-V1		

# Extended Page Replacement Algorithm Classification by Effelsberg and Härdter From [EH]

Consideration during selection decision		Age			
		No consideration	Since most recent reference	Since some recent reference	Since first reference
References	No consideration	RANDOM			FIFO FILO
	Most recent reference	ZCLOCK	LRU MRU CLOCK GCLOCK-V2 DGCLOCK-V2		
	Some recent references		SLRU	LRU-K	
	All references	LFU	GCLOCK-V1 DGCLOCK-V1		LFUDA

# Extended Page Replacement Algorithm Classification by Effelsberg and Härdter From [EH]

Consideration during selection decision		Age			
		No consideration	Since most recent reference	Since some recent reference	Since first reference
References	No consideration	RANDOM			FIFO FILO
	Most recent reference	ZCLOCK	LRU MRU CLOCK GCLOCK-V2 DGCLOCK-V2		
	Some recent references		SLRU	LRU-K	
	All references	LFU	GCLOCK-V1 DGCLOCK-V1		LRD-V1 LFUDA

# Extended Page Replacement Algorithm Classification by Effelsberg and Härdter From [EH]

Consideration during selection decision		Age			
		No consideration	Since most recent reference	Since some recent reference	Since first reference
References	No consideration	RANDOM			FIFO FILO
	Most recent reference	ZCLOCK	LRU MRU CLOCK GCLOCK-V2 DGCLOCK-V2		
	Some recent references		SLRU	LRU-K LRD-V2	
	All references	LFU	GCLOCK-V1 DGCLOCK-V1		LRD-V1 LFUDA

# Extended Page Replacement Algorithm Classification by Effelsberg and Härdter From [EH]

Consideration during selection decision		Age			
		No consideration	Since most recent reference	Since some recent reference	Since first reference
References	No consideration	RANDOM			FIFO FILO
	Most recent reference	ZCLOCK	LRU MRU CLOCK GCLOCK-V2 DGCLOCK-V2 LeanStore		
	Some recent references		SLRU	LRU-K LRD-V2	
	All references	LFU	GCLOCK-V1 DGCLOCK-V1		LRD-V1 LFUDA

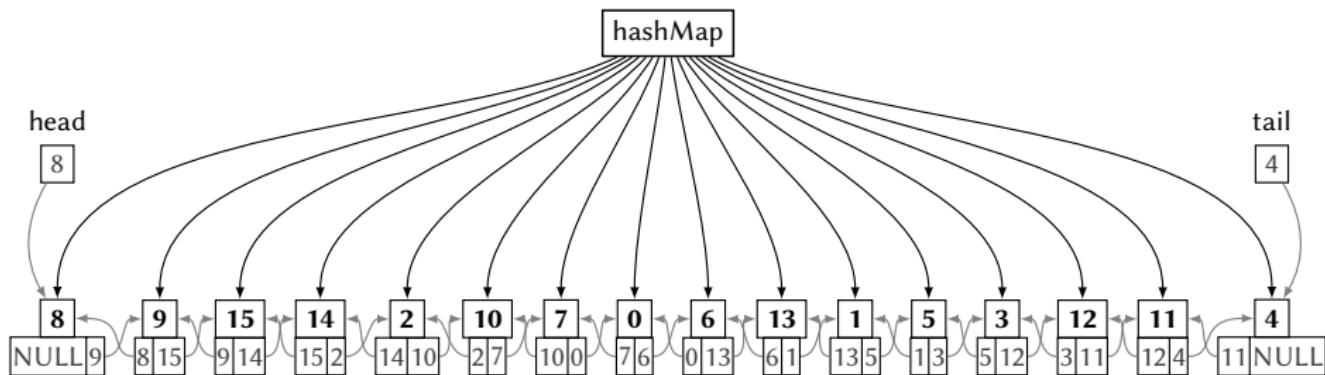
## Subsection 2

### Least Recently Used

### Subsubsection 1

## Hash-Map-Doubly-Linked-List Implementation

# Hash-Map-Doubly-Linked-List



## Subsubsection 2

### Timestamp-Sorting Implementation

# Page Reference Statistics

	liveTimestamps	LRUqueue0	LRUqueue1
0	2020-09-14 12:34:56.197548	2   2020-09-14 12:34:56 071309	8   2020-09-14 12:34:55 185406
1	2020-09-14 12:34:56 821122	0   2020-09-14 12:34:56.197548	9   2020-09-14 12:34:55.218616
2	2020-09-14 12:34:57 062836	12   2020-09-14 12:34:56 213471	1   2020-09-14 12:34:55 222490
3	2020-09-14 12:34:56 227794	3   2020-09-14 12:34:56 227794	13   2020-09-14 12:34:55 439575
4	2020-09-14 12:34:57 289365	14   2020-09-14 12:34:56 346659	11   2020-09-14 12:34:55 582807
5	2020-09-14 12:34:56 849729	6   2020-09-14 12:34:56 377511	12   2020-09-14 12:34:55 607186
6	2020-09-14 12:34:56 377511	8   2020-09-14 12:34:56 498737	5   2020-09-14 12:34:55 611986
7	2020-09-14 12:34:57 179364	10   2020-09-14 12:34:56 720335	4   2020-09-14 12:34:55 688895
8	2020-09-14 12:34:56 498737	9   2020-09-14 12:34:56 776239	15   2020-09-14 12:34:55 726161
9	2020-09-14 12:34:56 776239	15   2020-09-14 12:34:56 793698	7   2020-09-14 12:34:55 742479
10	2020-09-14 12:34:56 720335	1   2020-09-14 12:34:56 821122	10   2020-09-14 12:34:55 809492
11	2020-09-14 12:34:56 855806	5   2020-09-14 12:34:56 849729	6   2020-09-14 12:34:55 953010
12	2020-09-14 12:34:56 213471	11   2020-09-14 12:34:56 855806	3   2020-09-14 12:34:55 955357
13	2020-09-14 12:34:56 881209	4   2020-09-14 12:34:56 865145	14   2020-09-14 12:34:55 978132
14	2020-09-14 12:34:56 346659	13   2020-09-14 12:34:56 881209	2   2020-09-14 12:34:56 071309
15	2020-09-14 12:34:56 793698	7   2020-09-14 12:34:56 987742	0   2020-09-14 12:34:56.197548

Diagram illustrating the state of the buffer manager:

- lastChecked**: A circular icon with a small circle inside, connected by a curved arrow to the **useLRUqueue0** field.
- useLRUqueue0**: Value: true
- useLRUqueue1**: Value: false
- sortingInProgress**: Value: false
- waitForSortedMutex**: Value:

## Subsection 3

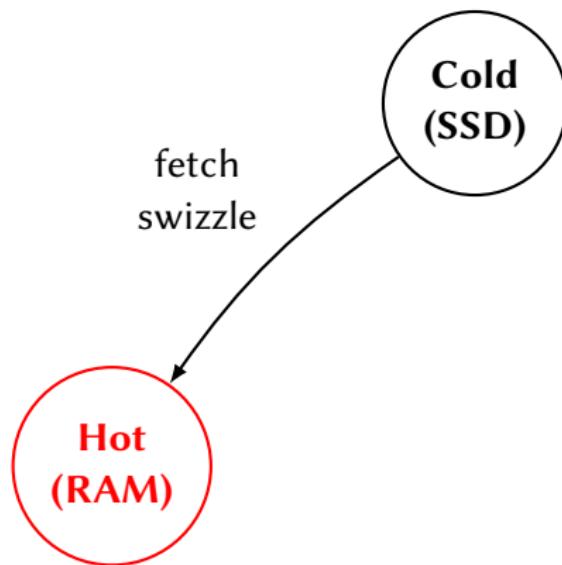
### LeanStore Replacement [Lei+]

# Basic Concept (Pointer Swizzling)

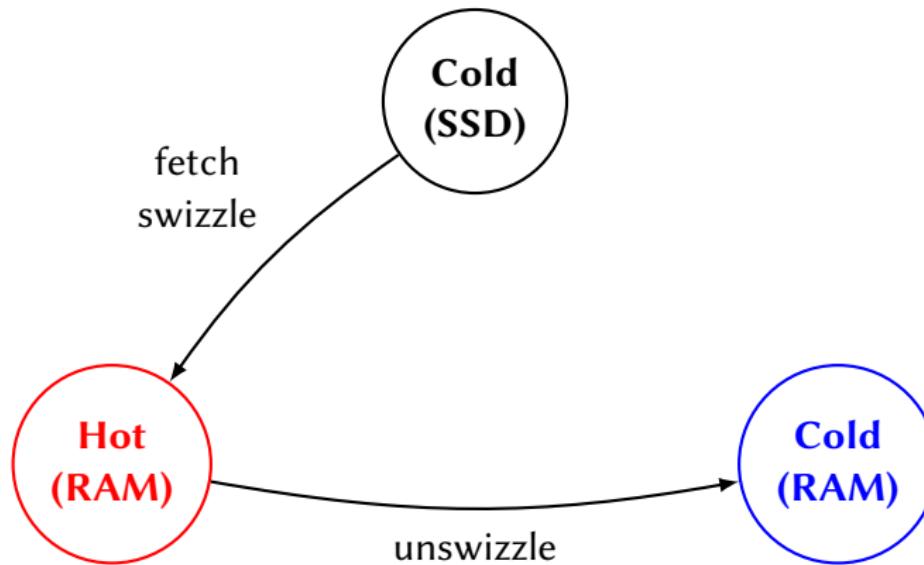
# Basic Concept (Pointer Swizzling)



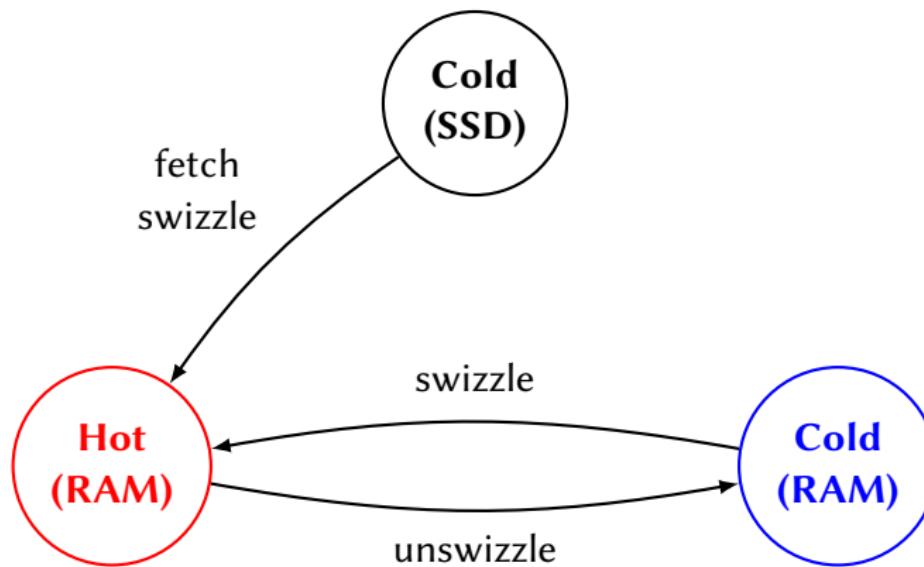
## Basic Concept (Pointer Swizzling)



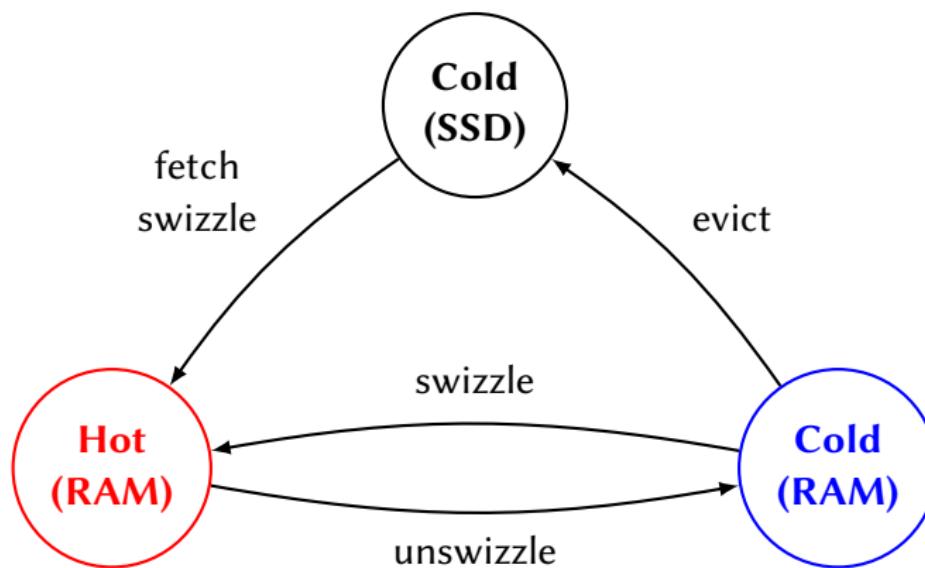
## Basic Concept (Pointer Swizzling)



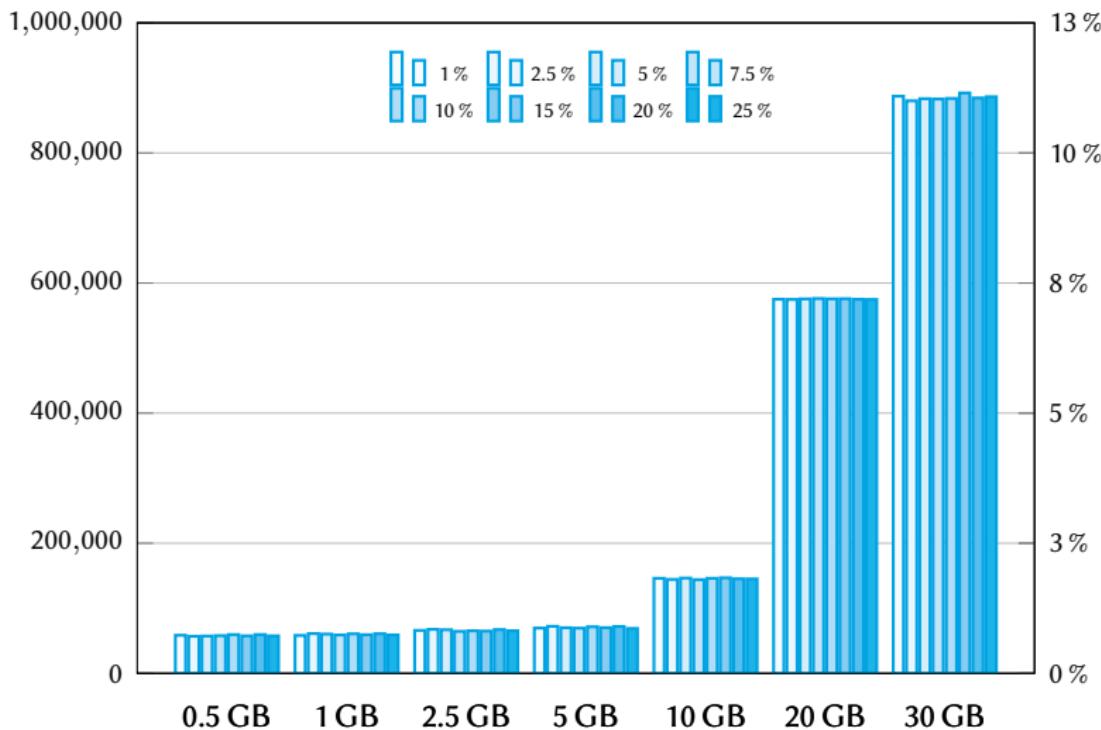
## Basic Concept (Pointer Swizzling)



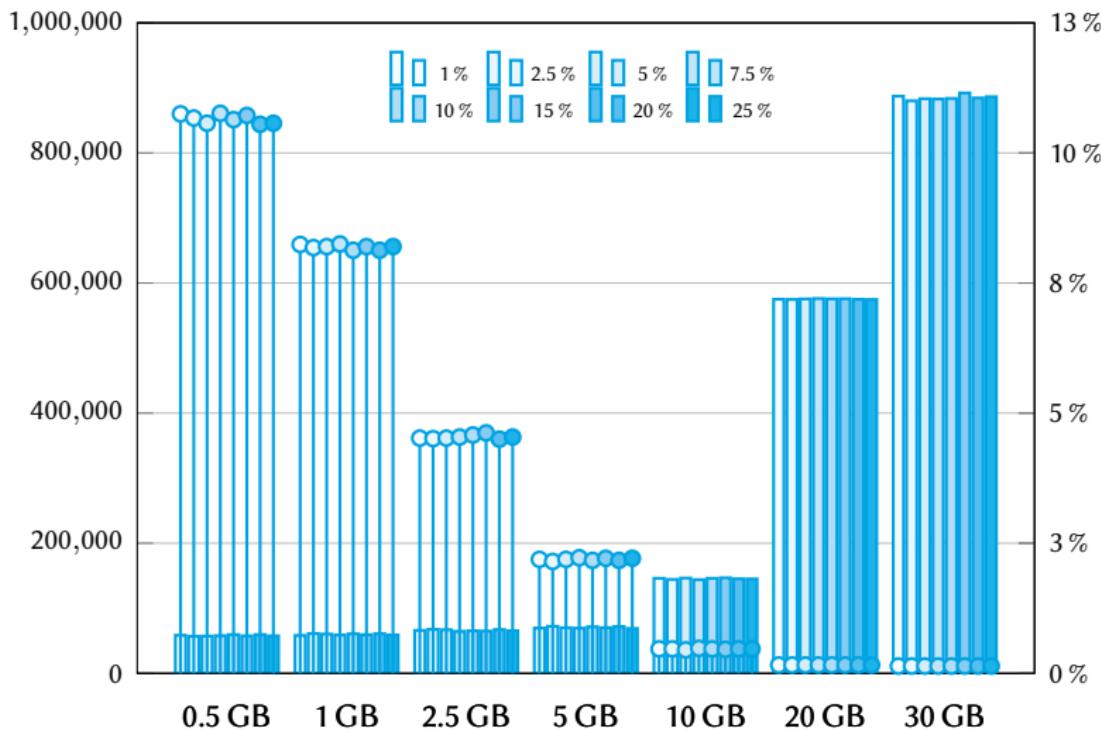
## Basic Concept (Pointer Swizzling)



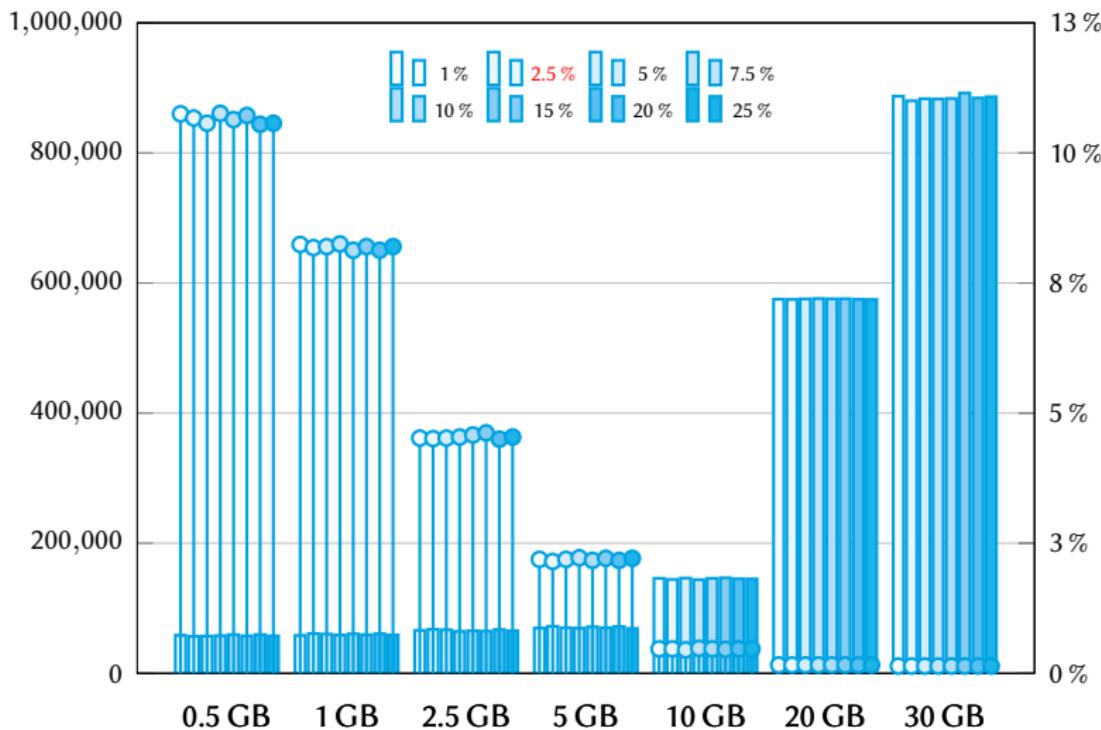
# Performance Evaluation



# Performance Evaluation



# Performance Evaluation



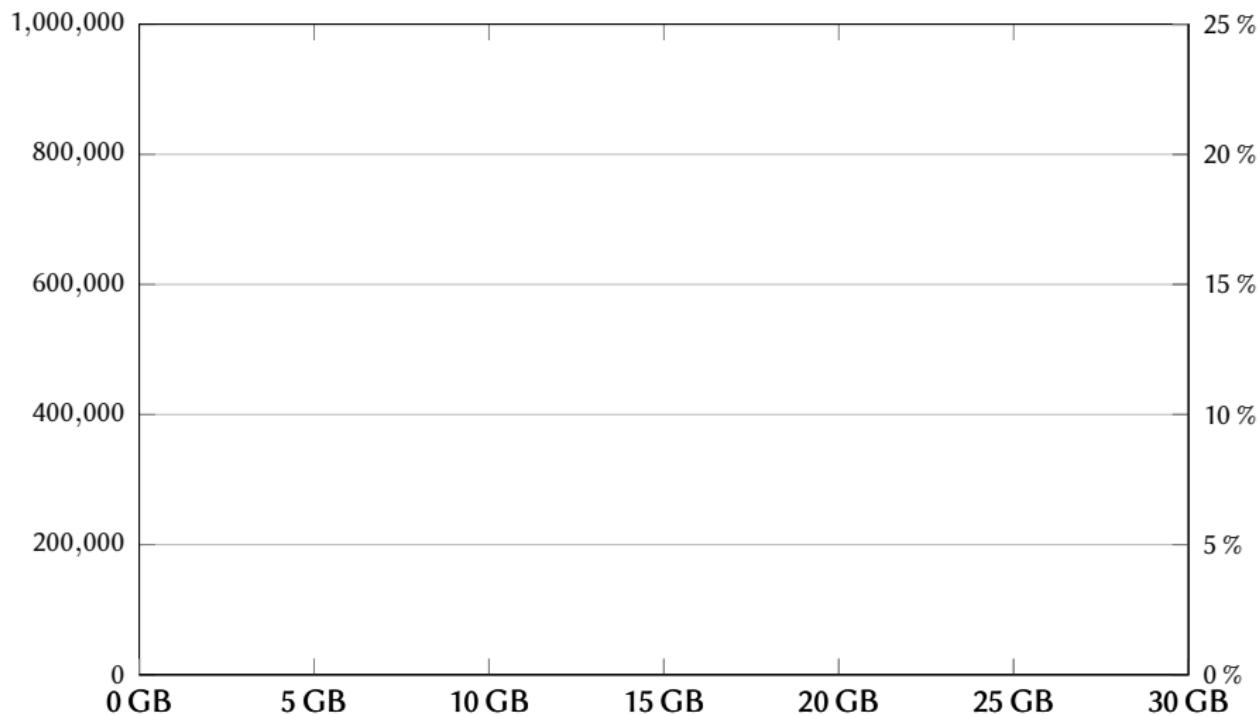
## Subsection 4

### All Page Replacement Algorithms

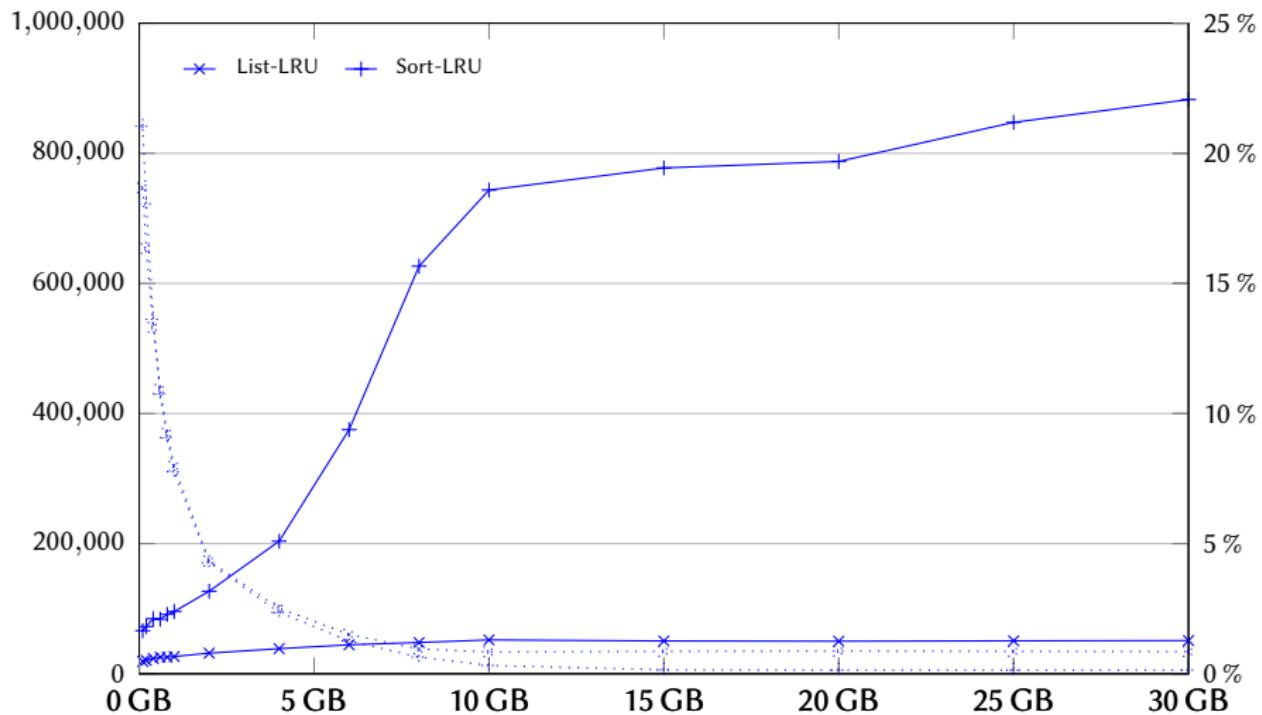
# Overview

Consideration during selection decision		Age			
		No consideration	Since most recent reference	Since some recent reference	Since first reference
References	No consideration	RANDOM			FIFO FILO
	Most recent reference	ZCLOCK	LRU MRU CLOCK GCLOCK-V2 DGCLOCK-V2 LeanStore		
	Some recent references		SLRU	LRU-K LRD-V2	
	All references	LFU	GCLOCK-V1 DGCLOCK-V1		LRD-V1 LFUDA

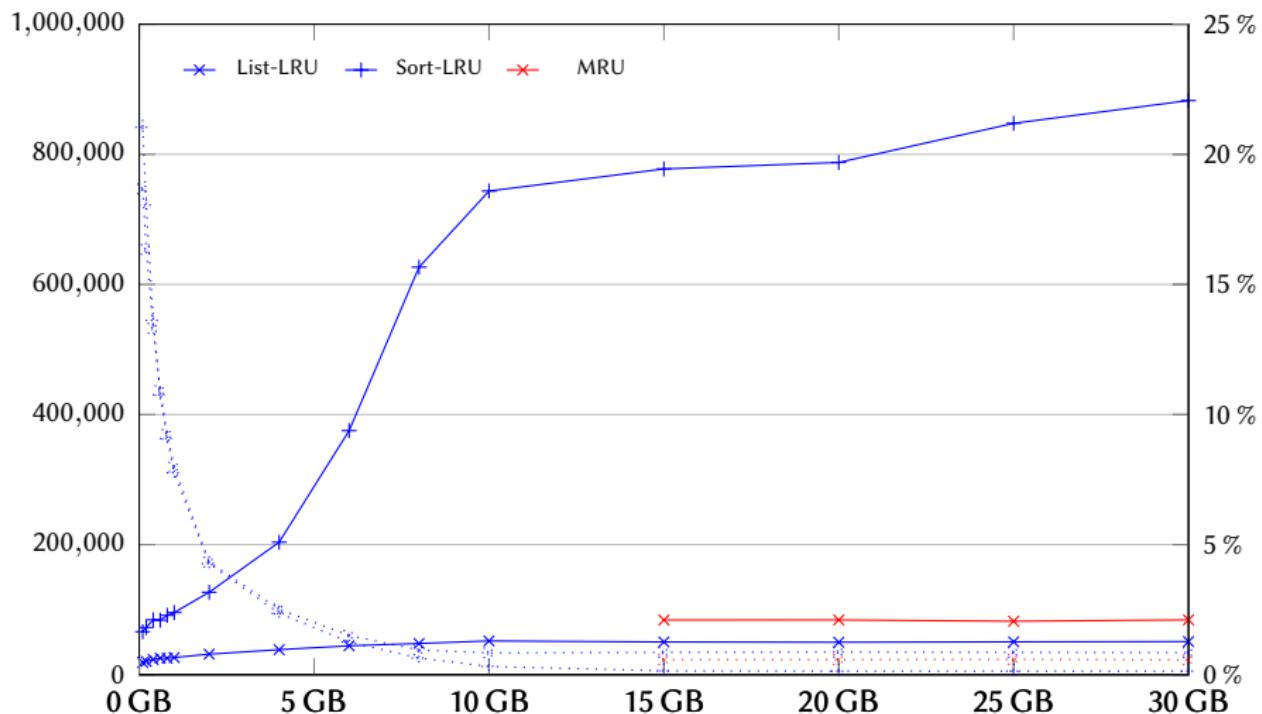
# Performance Evaluation



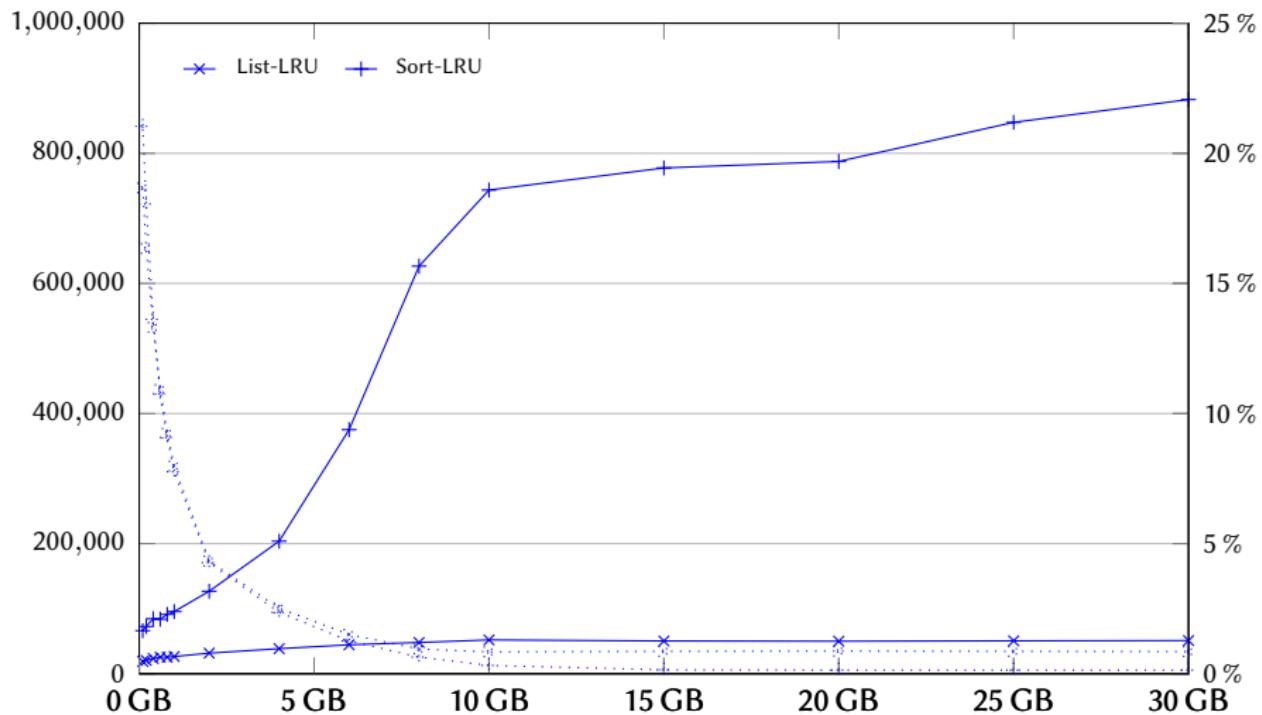
# Performance Evaluation



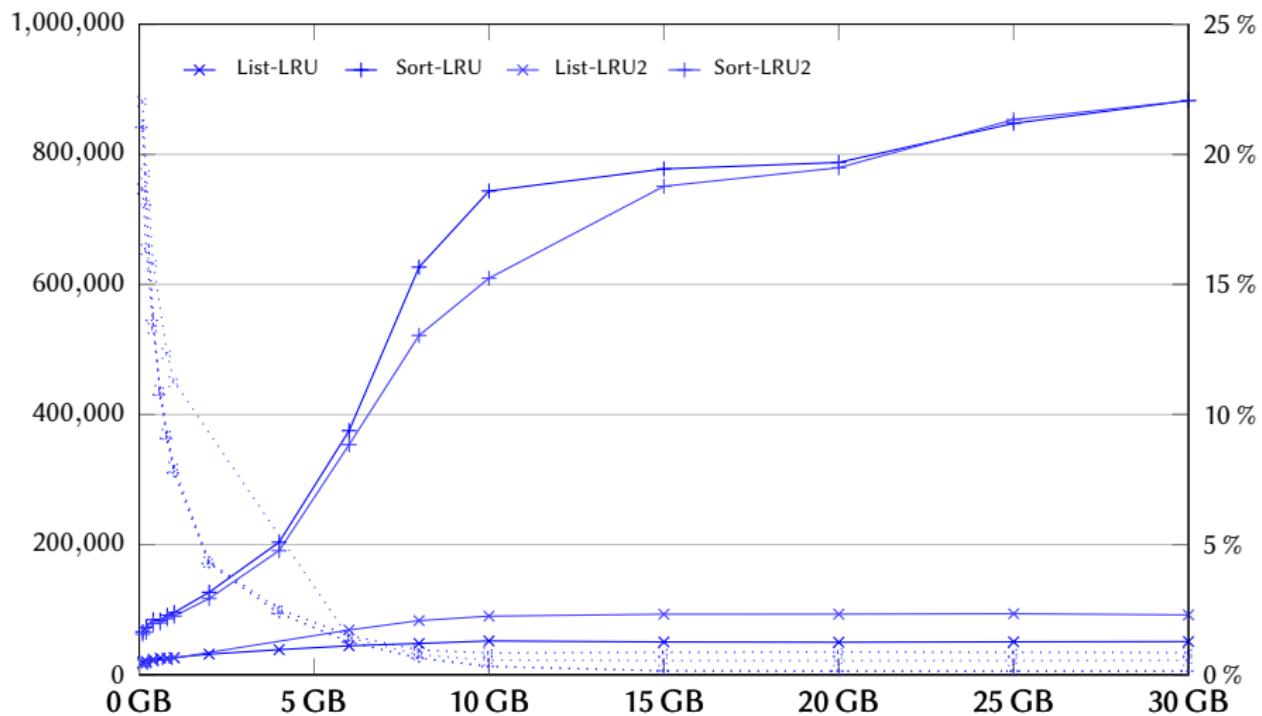
# Performance Evaluation



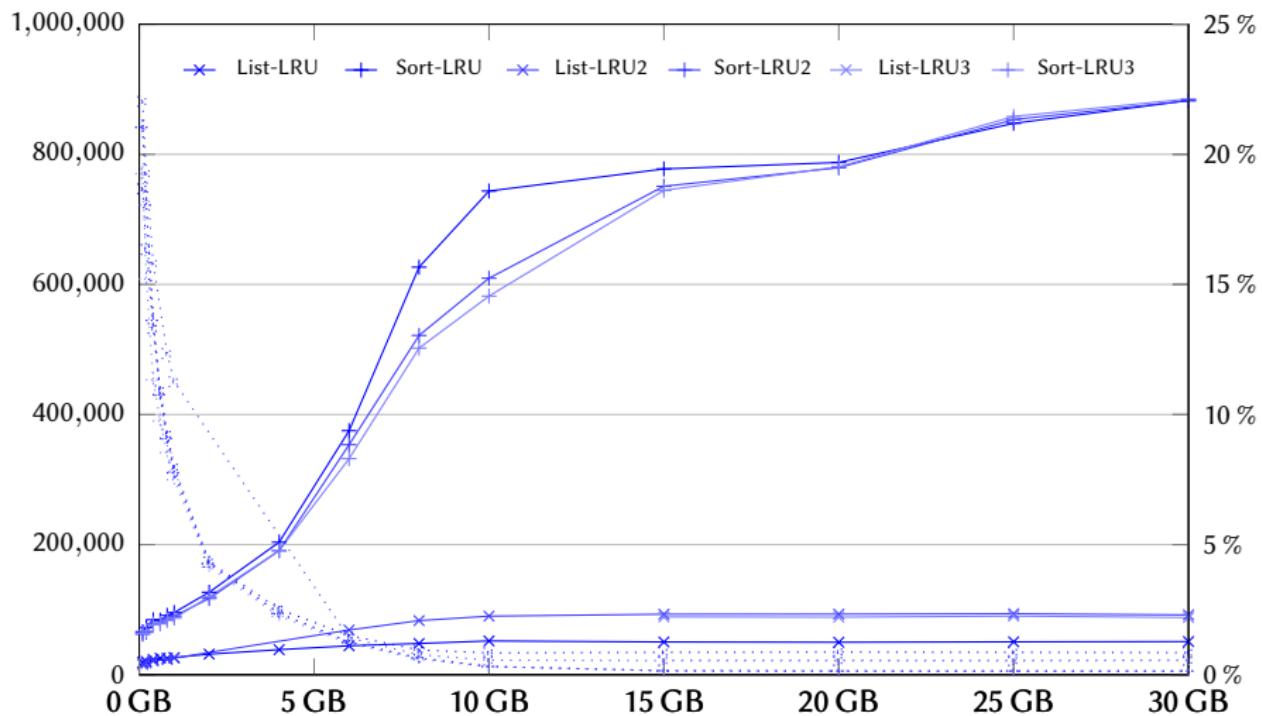
# Performance Evaluation



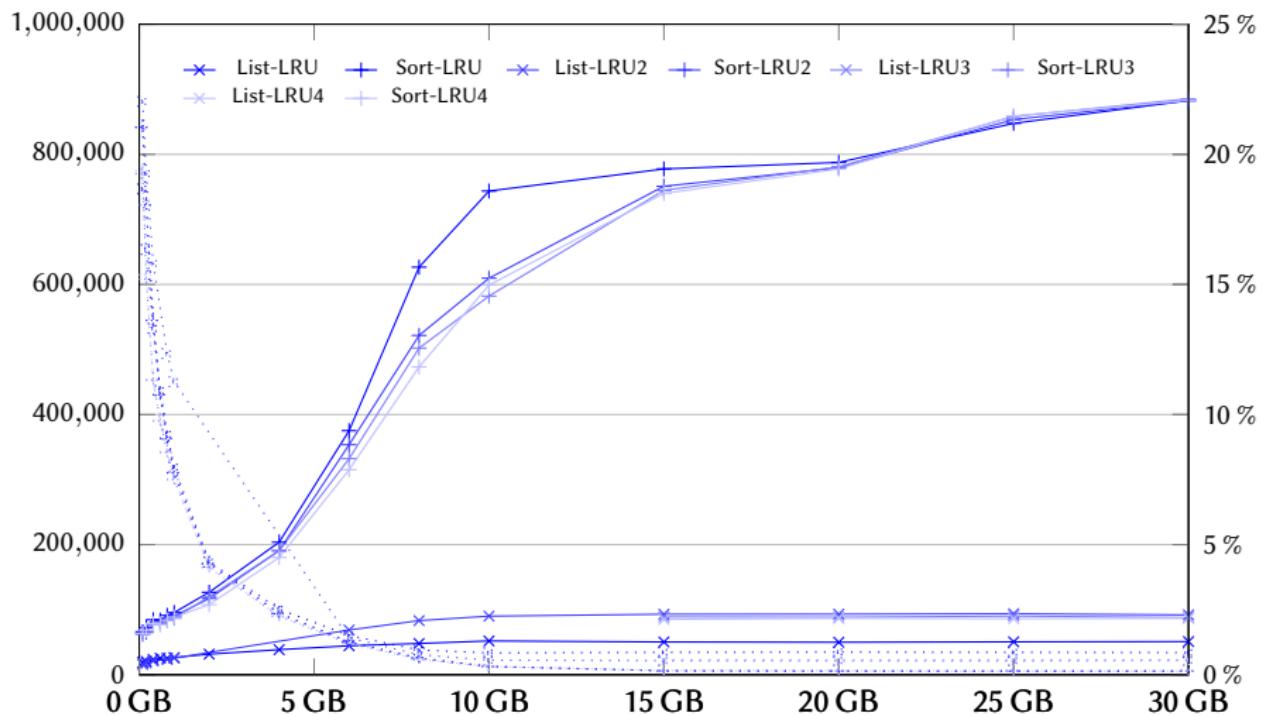
# Performance Evaluation



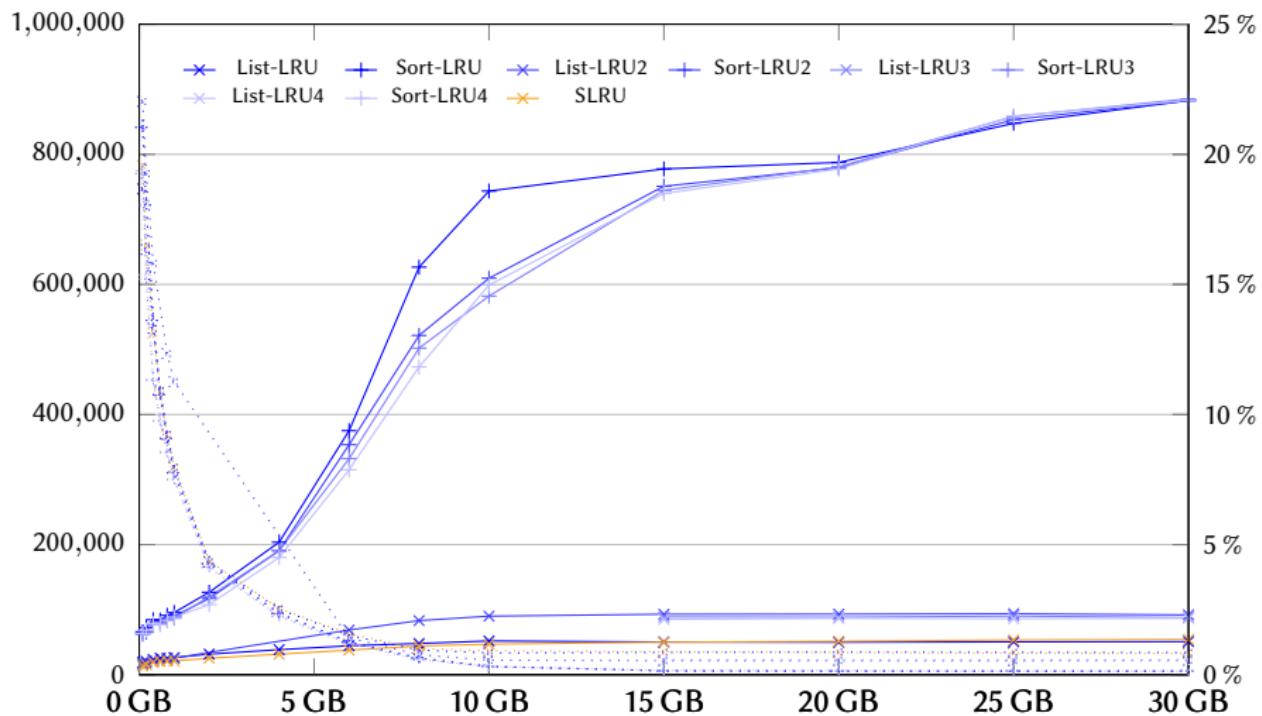
# Performance Evaluation



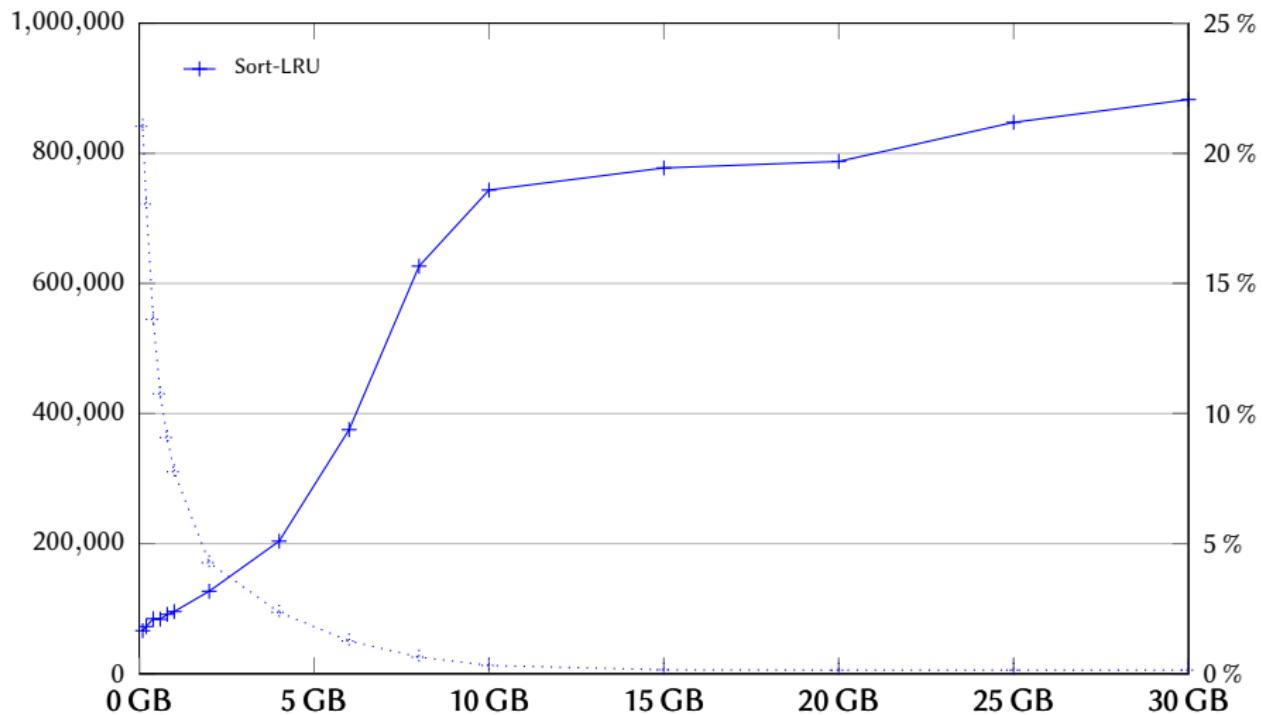
# Performance Evaluation



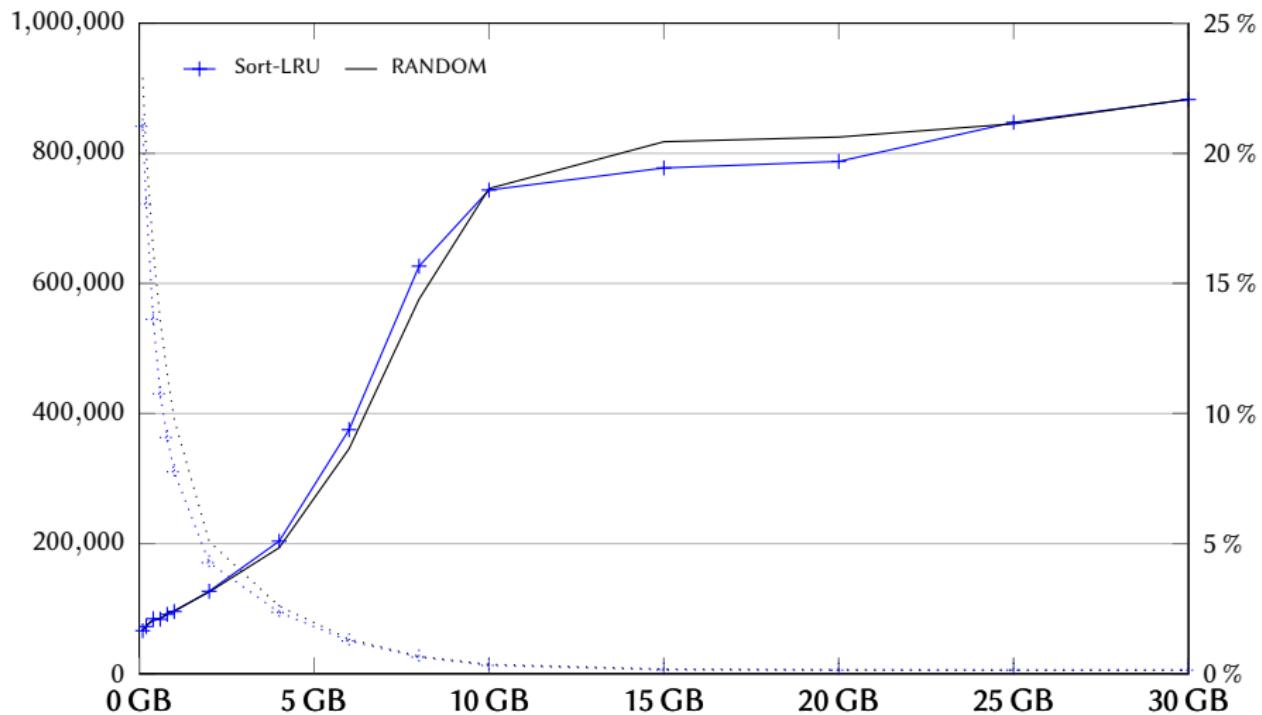
# Performance Evaluation



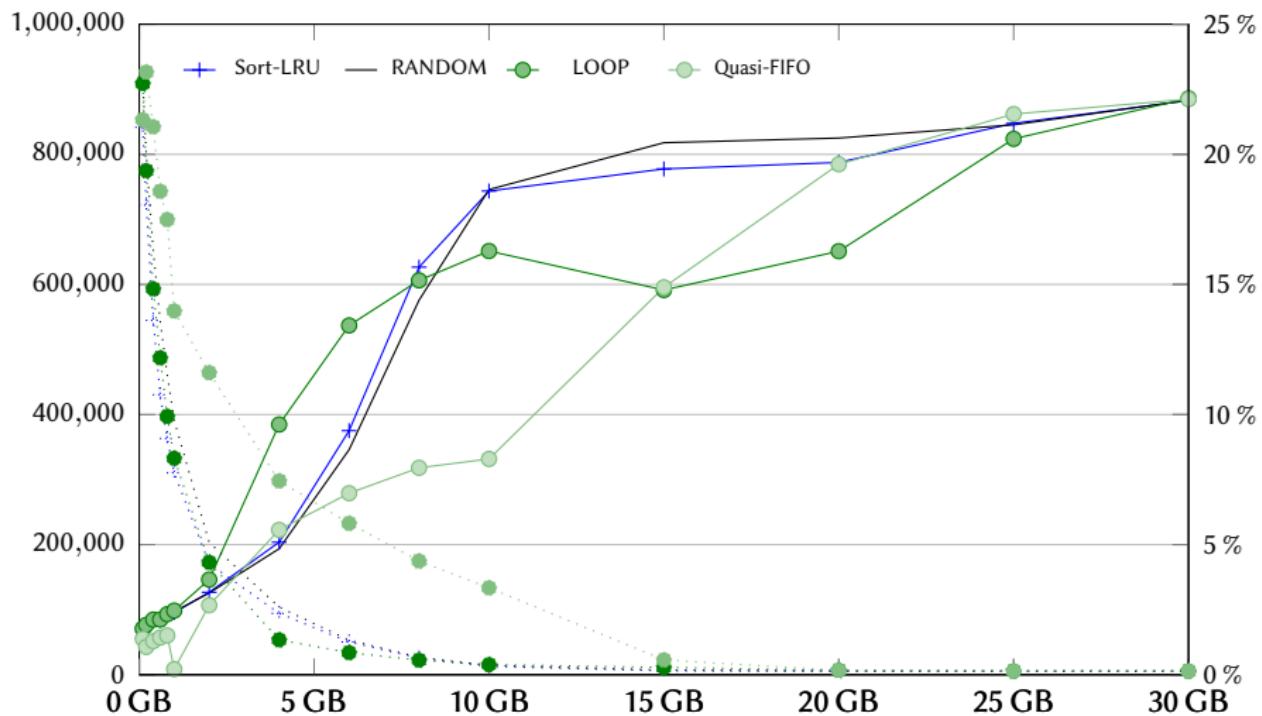
# Performance Evaluation



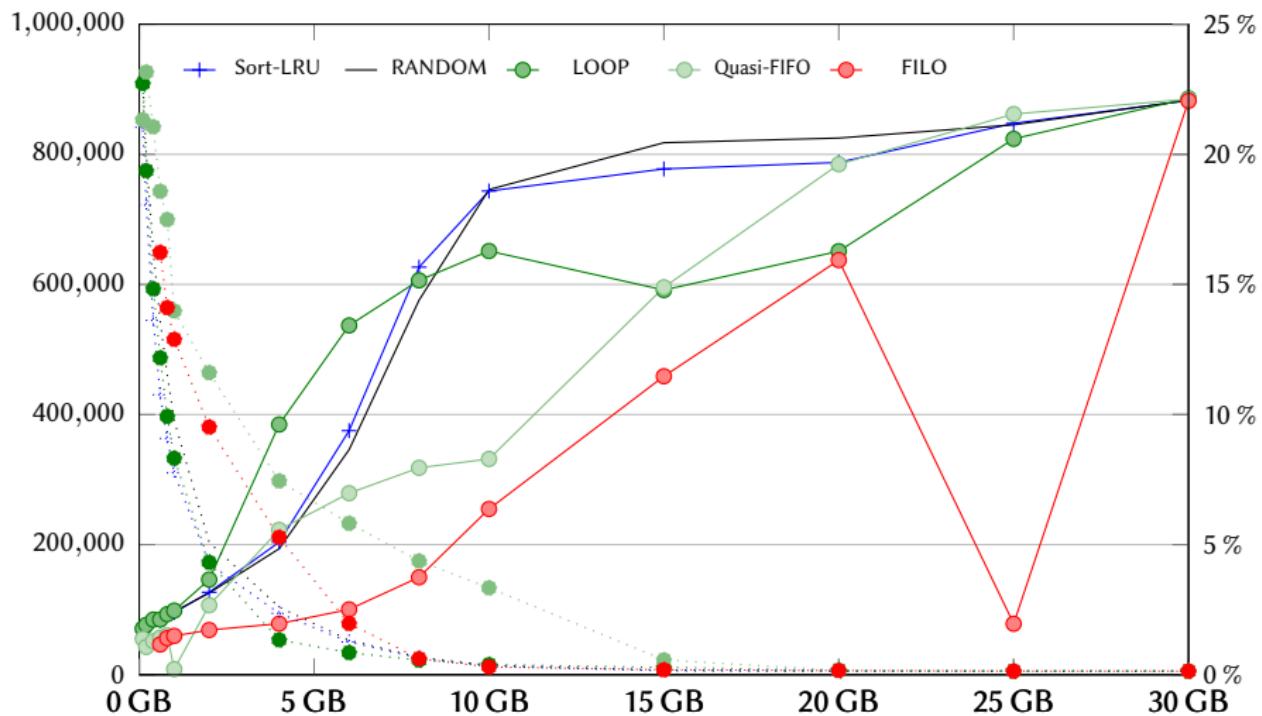
# Performance Evaluation



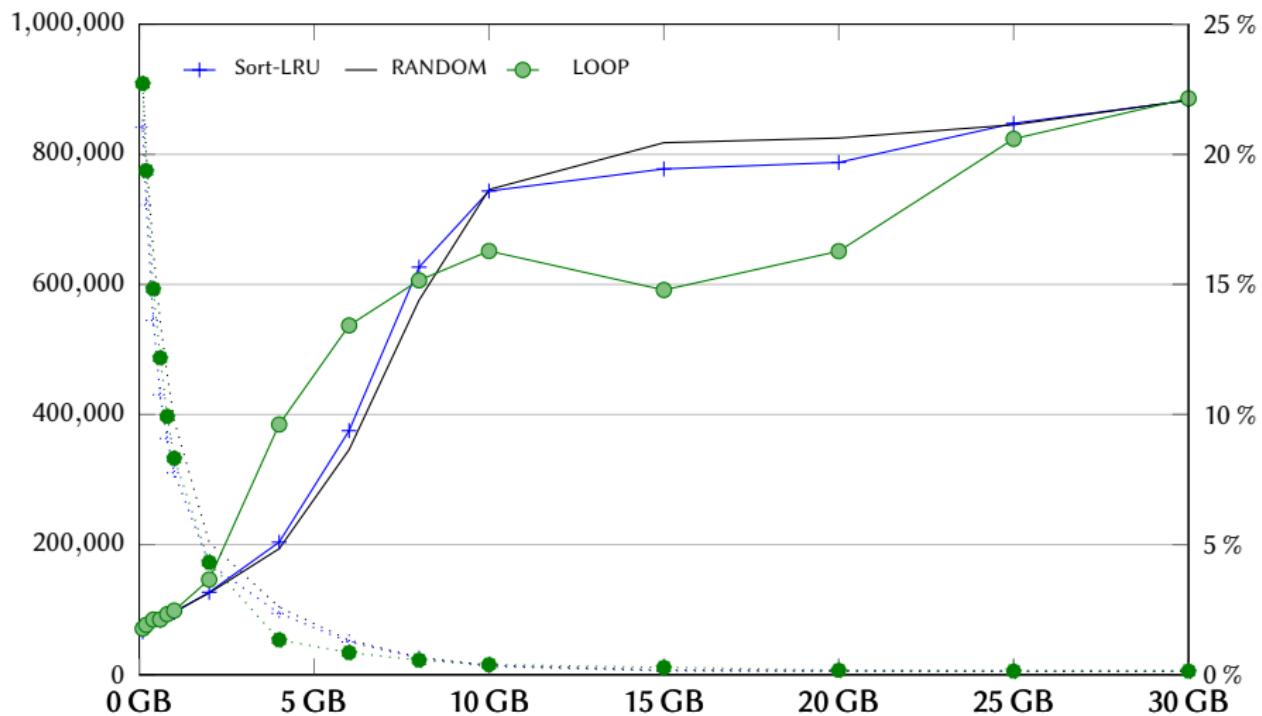
# Performance Evaluation



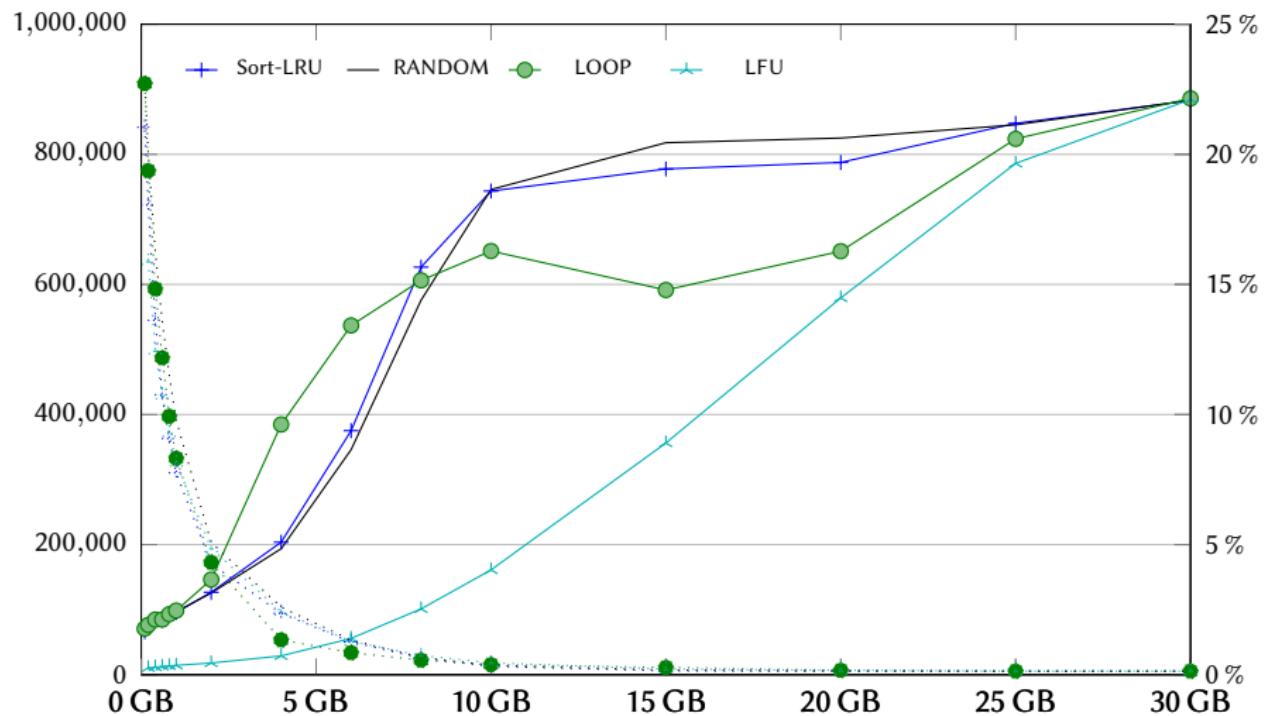
# Performance Evaluation



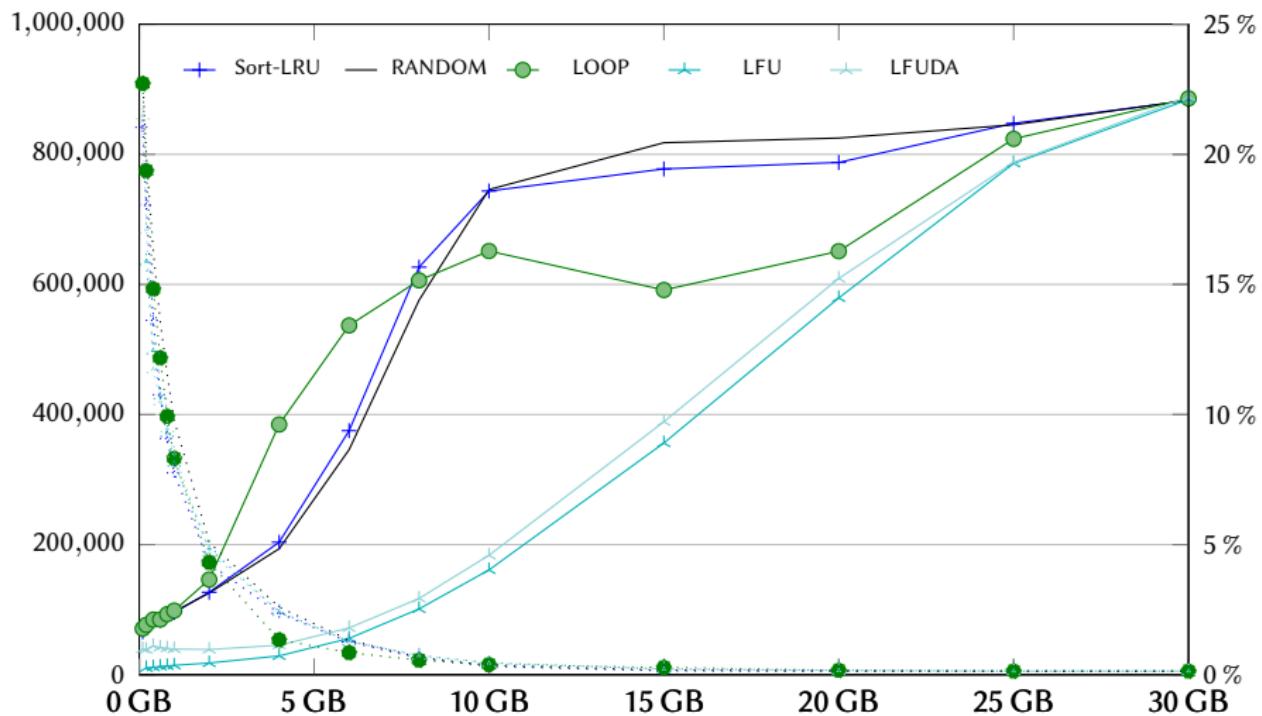
# Performance Evaluation



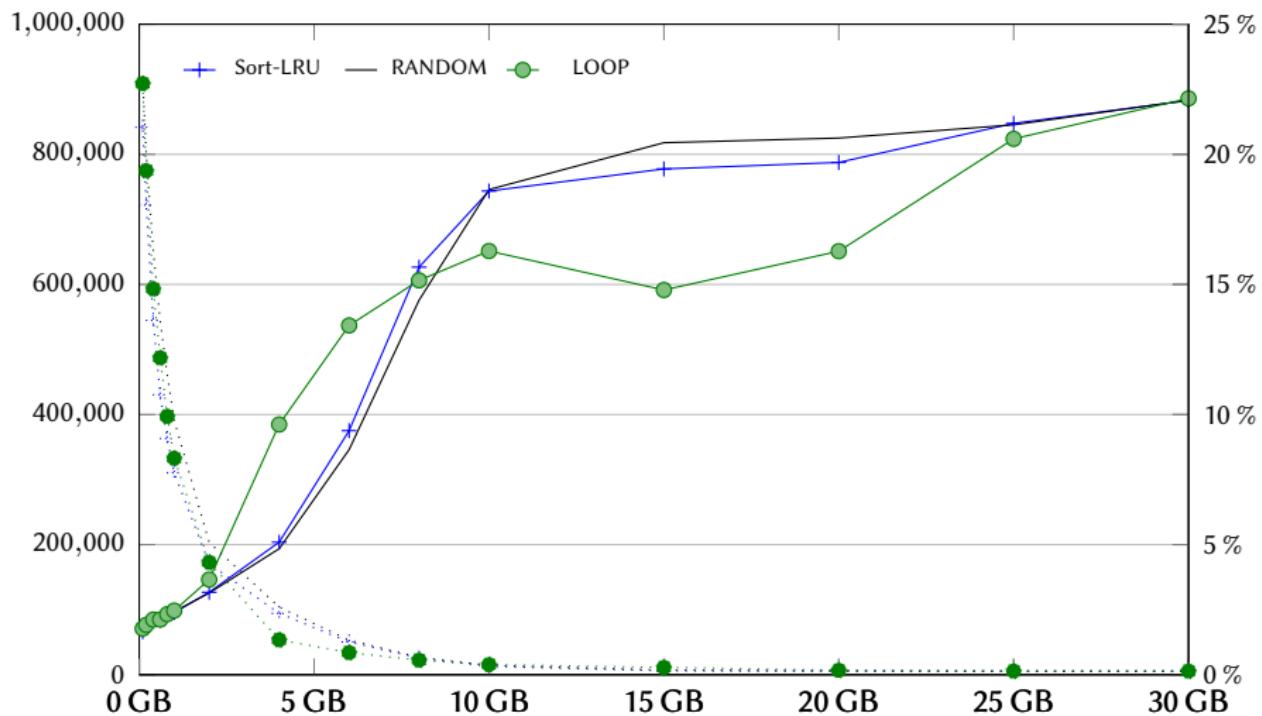
# Performance Evaluation



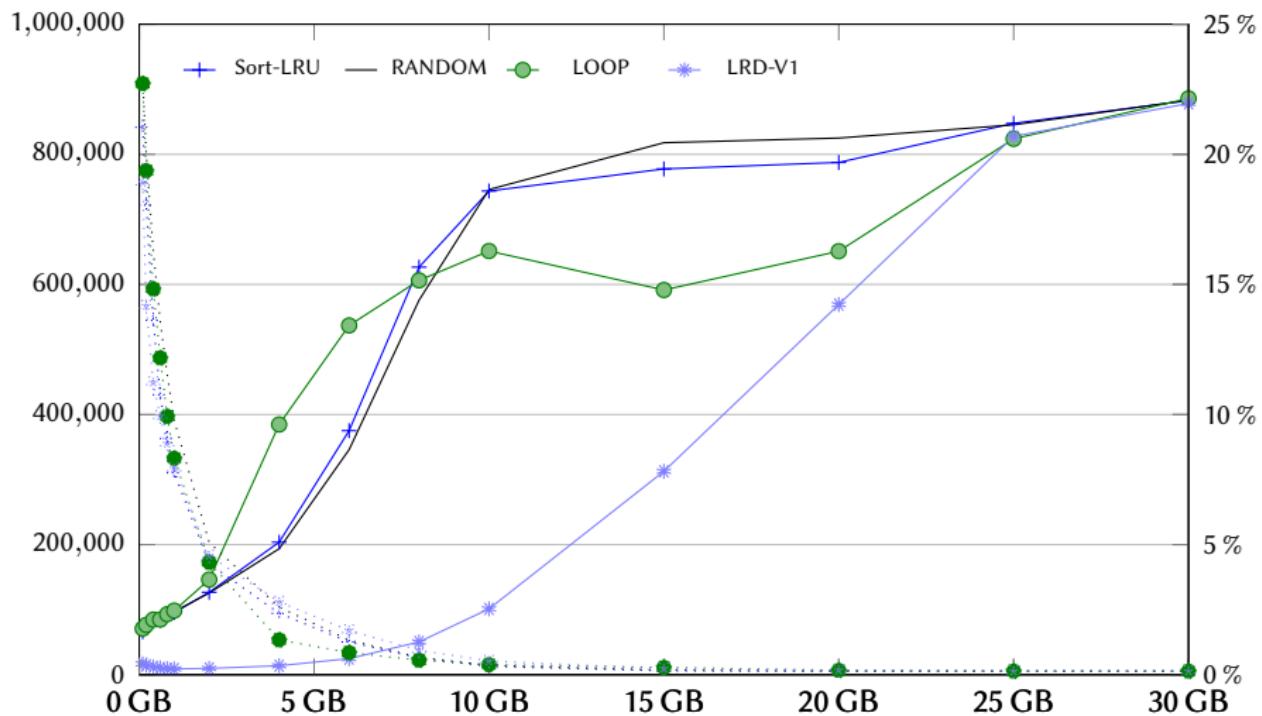
# Performance Evaluation



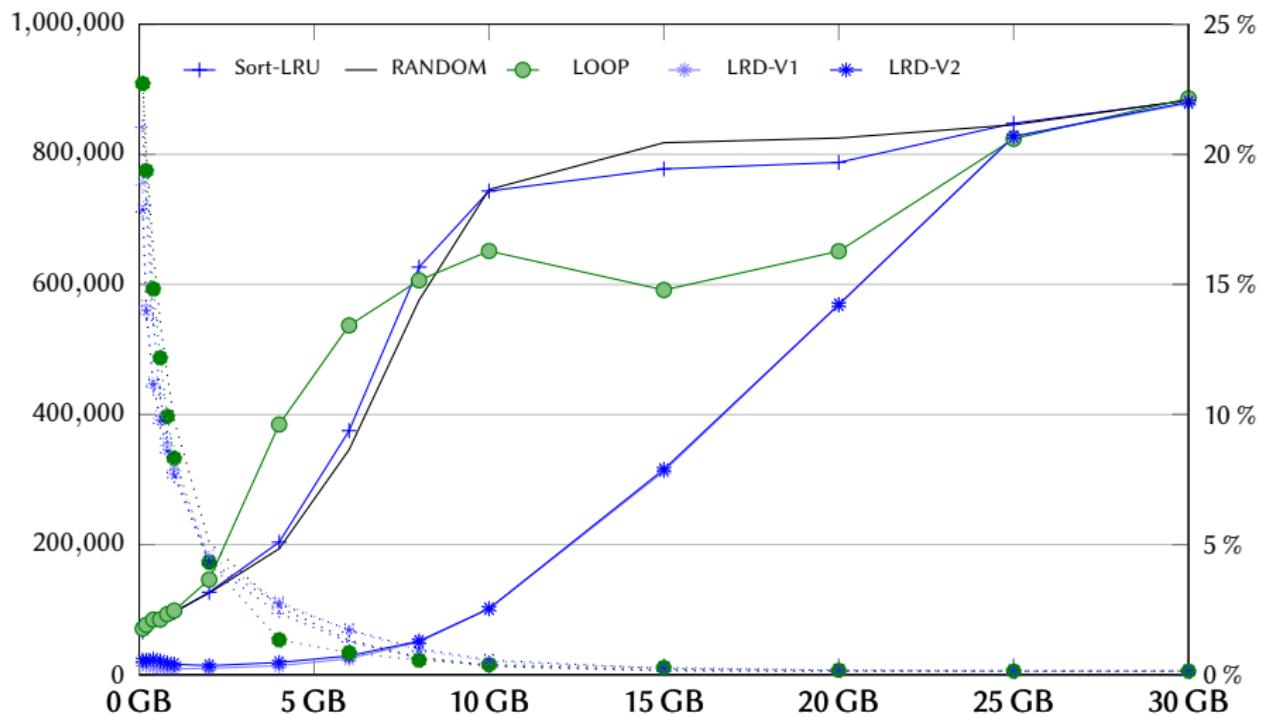
# Performance Evaluation



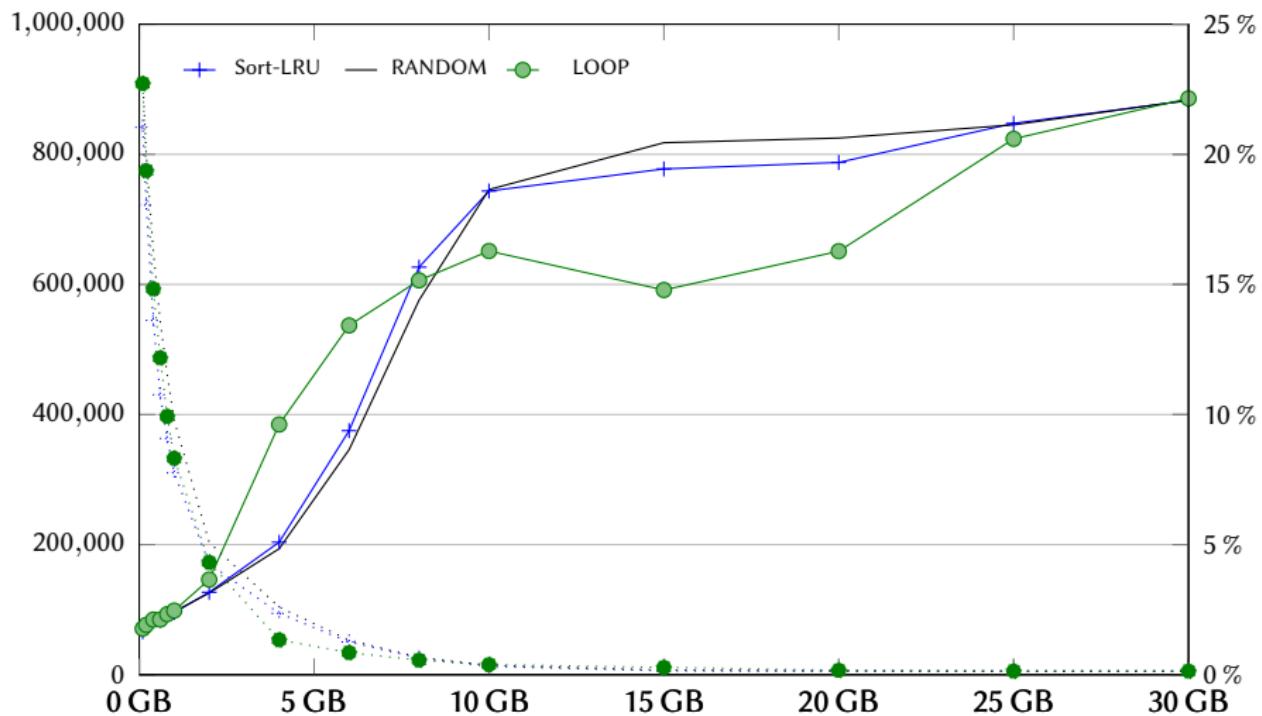
# Performance Evaluation



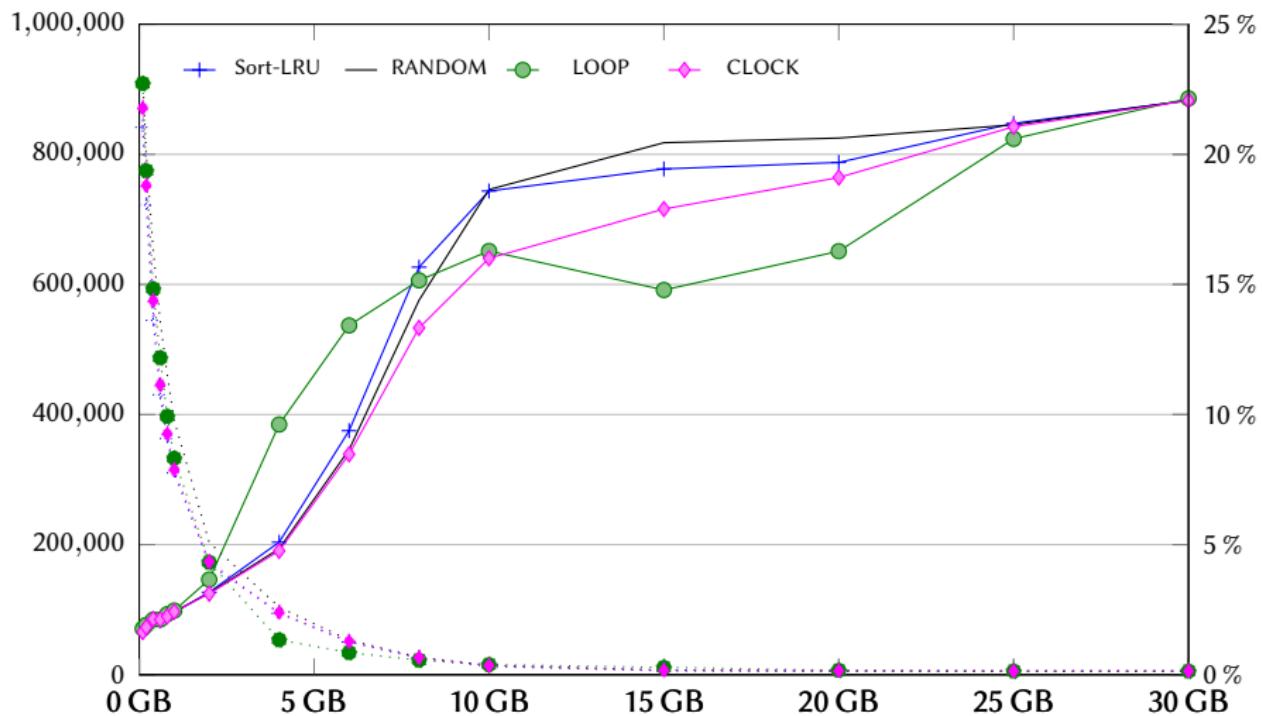
# Performance Evaluation



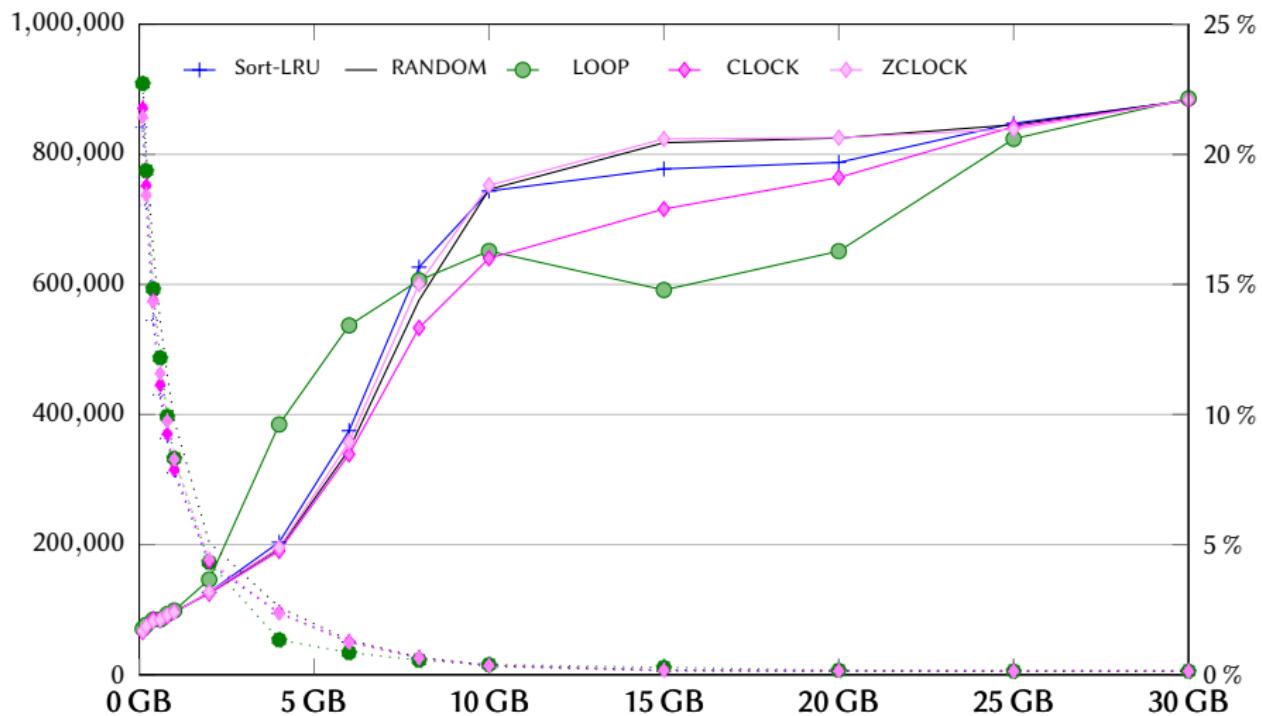
# Performance Evaluation



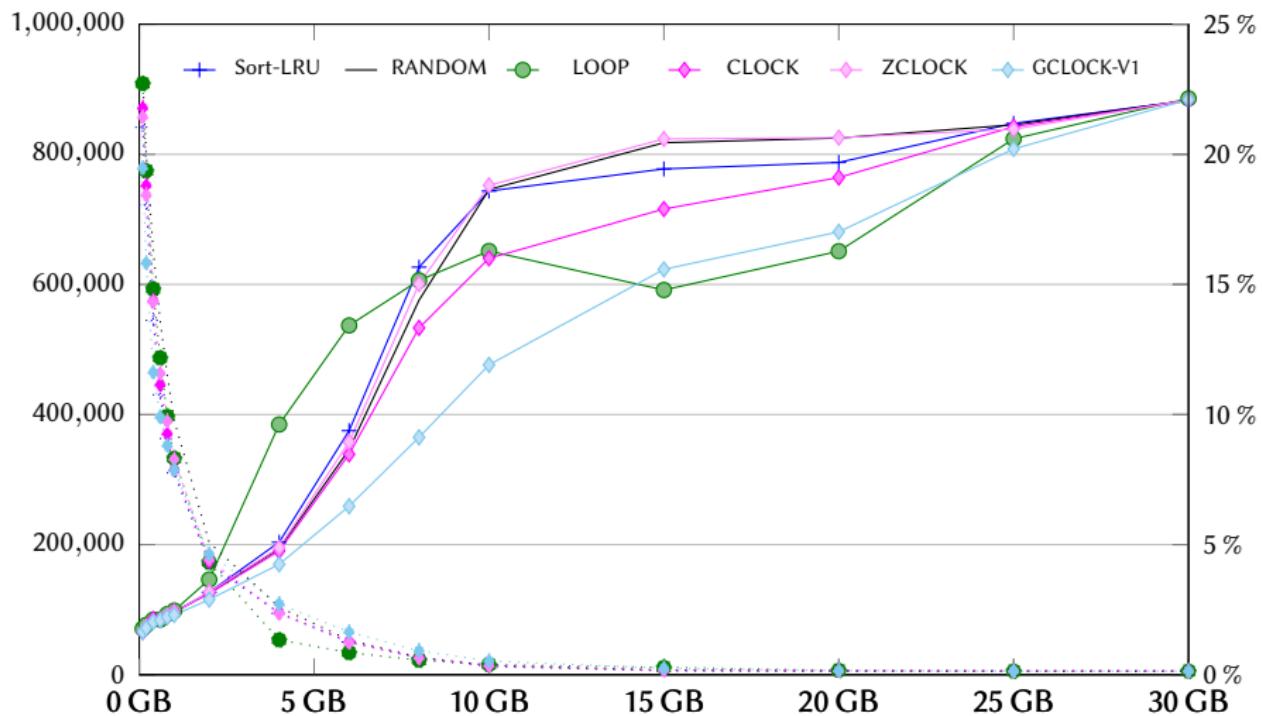
# Performance Evaluation



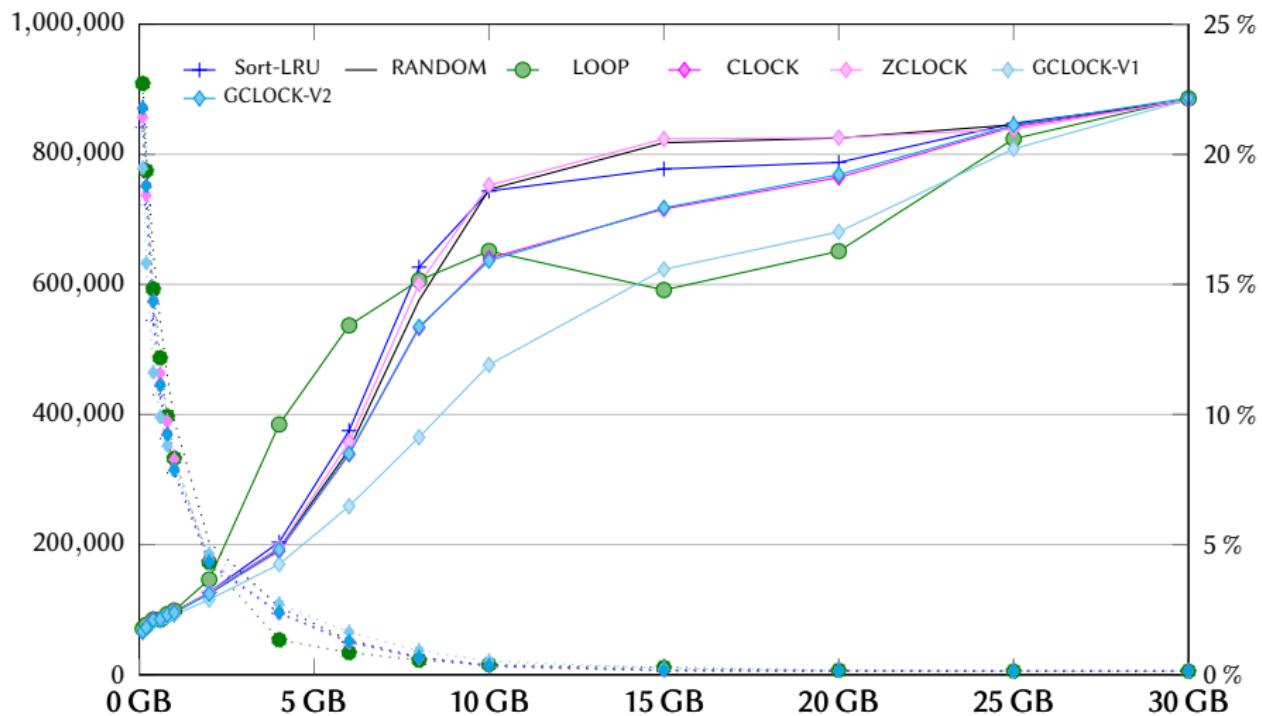
# Performance Evaluation



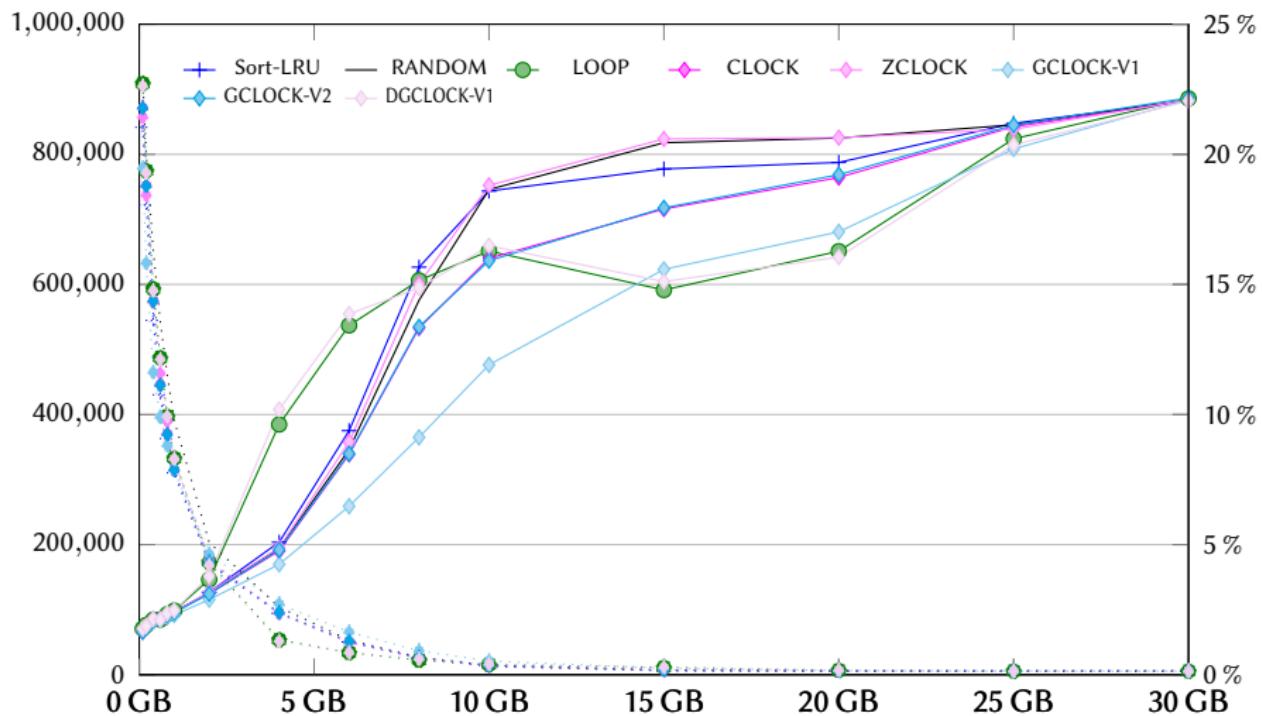
# Performance Evaluation



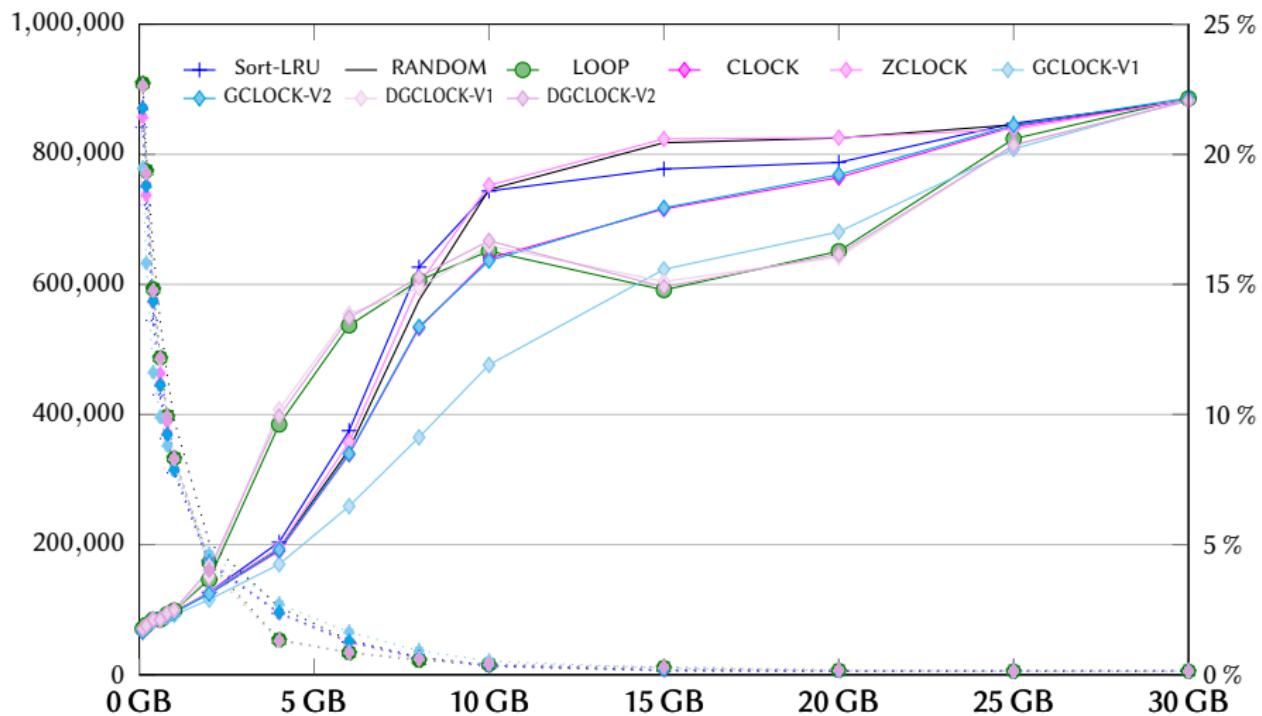
# Performance Evaluation



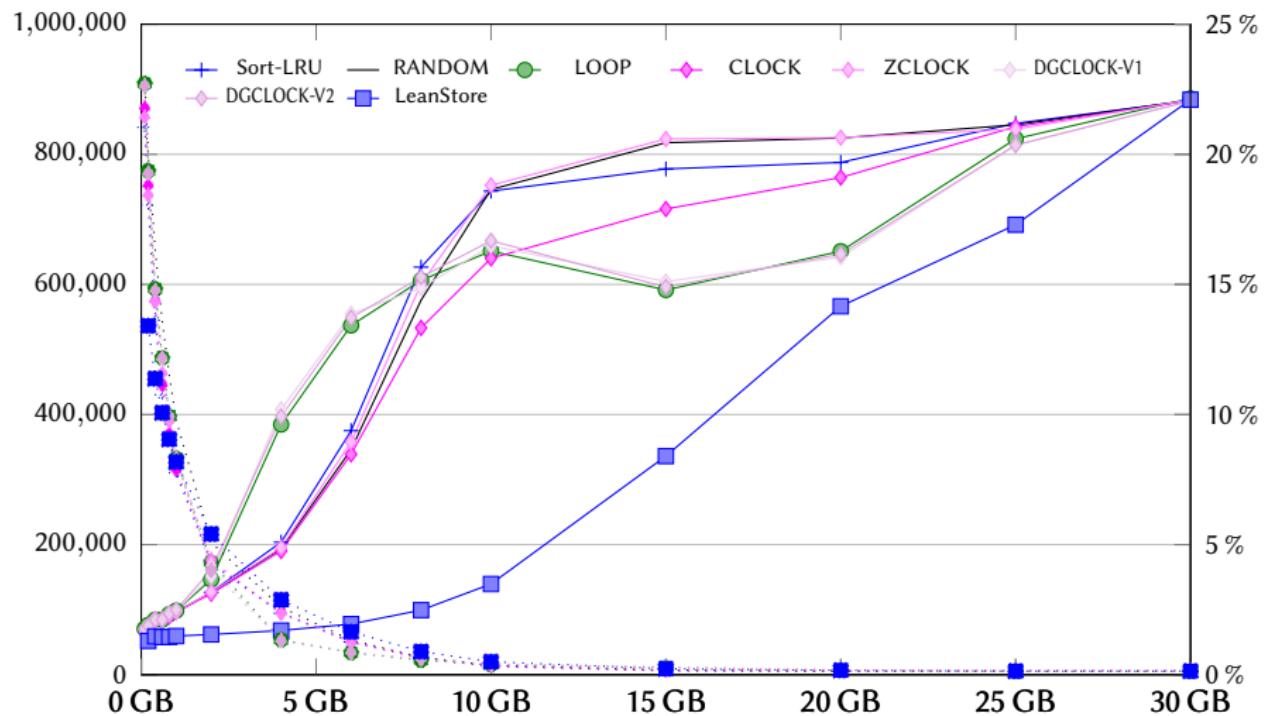
# Performance Evaluation



# Performance Evaluation



# Performance Evaluation



# Conclusion

Simplicity is the Key

RANDOM and LOOP outperform the more complex competitors

# Conclusion

Simplicity is the Key

RANDOM and LOOP outperform the more complex competitors (at least in combination with *Zero*, *TPC-C* and the specific benchmark configuration)

## Section 2

# Component-Wise Performance Evaluation of OLTP Systems

## Subsection 1

OLTP through the Looking Glass, and What We Found  
There [Har+]

## OLTP Through the Looking Glass, and What We Found There

Stavros Harizopoulos  
HP Labs  
Palo Alto, CA  
stavros@hp.com

Daniel J. Abadi  
Yale University  
New Haven, CT  
dna@cs.yale.edu

Samuel Madden  
Massachusetts Institute of Technology  
Cambridge, MA  
{madden, stonebraker}@csail.mit.edu

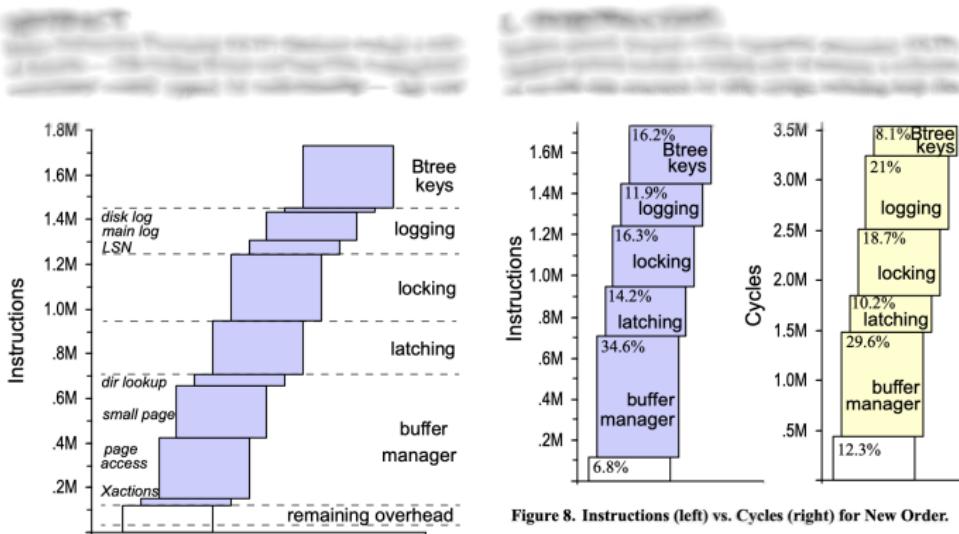


Figure 7. Expanding breakdown for New Order (see Section 3.2 for the labels on the left column).

Figure 8. Instructions (left) vs. Cycles (right) for New Order.

# The Author's Assumptions

# The Author's Assumptions

- ▶ No more larger-than-memory DBs

# The Author's Assumptions

- ▶ No more larger-than-memory DBs
- ▶ Thus, multithreaded and interleaved transaction execution no more needed

## The Author's Assumptions

- ▶ No more larger-than-memory DBs
- ▶ Thus, multithreaded and interleaved transaction execution no more needed
- ▶ Many applications do not need ACID transactions

## The Author's Assumptions

- ▶ No more larger-than-memory DBs
- ▶ Thus, multithreaded and interleaved transaction execution no more needed
- ▶ Many applications do not need ACID transactions
- ▶ Replication instead of logging for recovery

# The Author's Methodology

# The Author's Methodology

Gradually ...

# The Author's Methodology

Gradually ...

- ▶ ... make the DBS in-memory—removing buffer pool—, ...

# The Author's Methodology

Gradually ...

- ▶ ... make the DBS in-memory—removing buffer pool, ...
- ▶ ... remove latching

# The Author's Methodology

Gradually ...

- ▶ ... make the DBS in-memory—removing buffer pool, ...
- ▶ ... remove latching and locking, ...

# The Author's Methodology

Gradually ...

- ▶ ... make the DBS in-memory—removing buffer pool, ...
- ▶ ... remove latching and locking, ...
- ▶ ... remove transaction logging and LSN maintenance, ...

# The Author's Methodology

Gradually ...

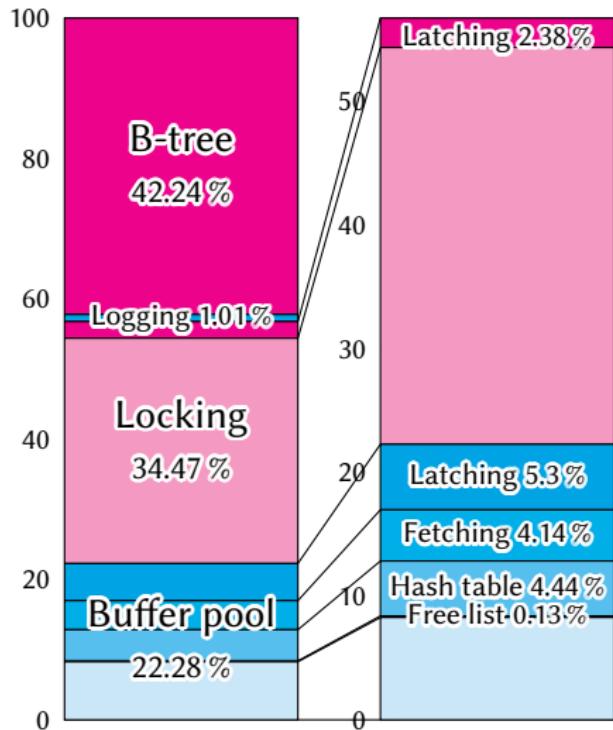
- ▶ ... make the DBS in-memory—removing buffer pool, ...
- ▶ ... remove latching and locking, ...
- ▶ ... remove transaction logging and LSN maintenance, ...
- ▶ ... and hand-optimize B-tree code for in-memory DBS.

## Subsection 2

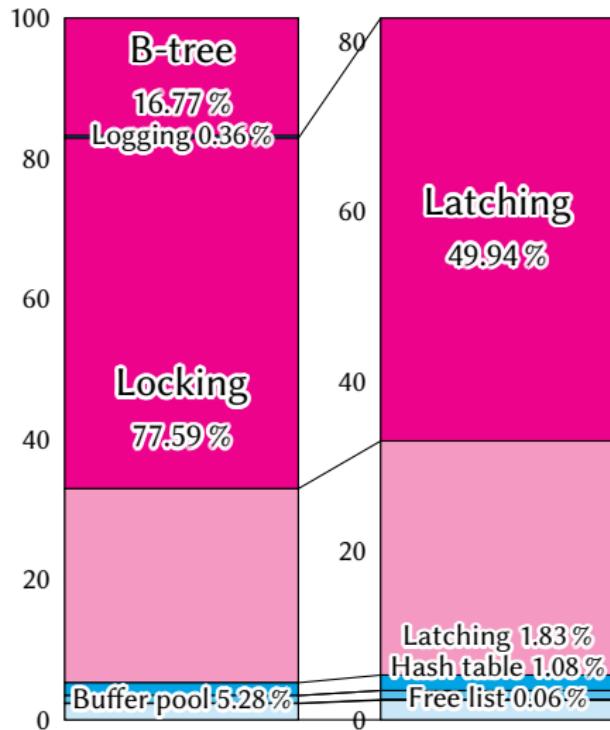
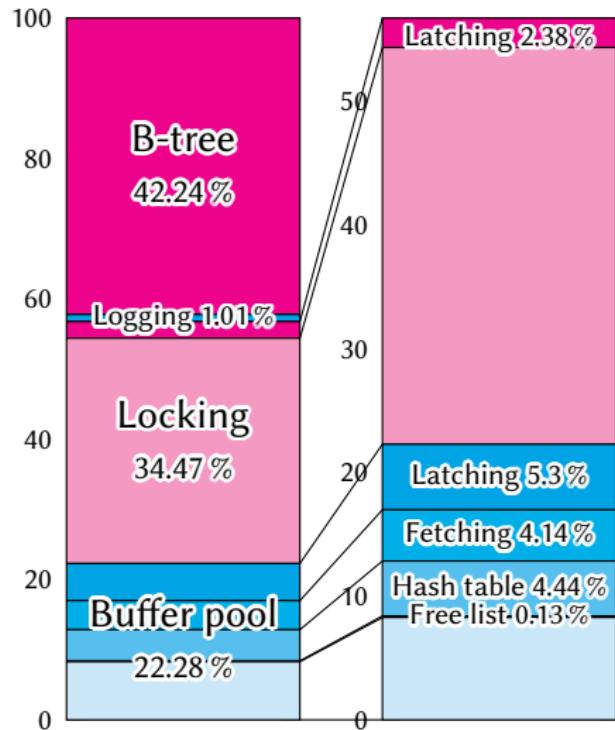
### Results

# Read-Only YCSB

## Read-Only YCSB

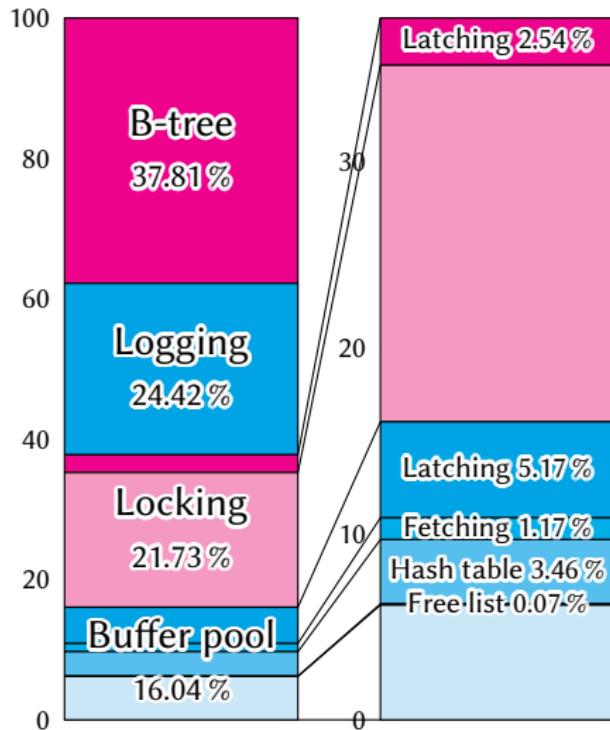


## Read-Only YCSB

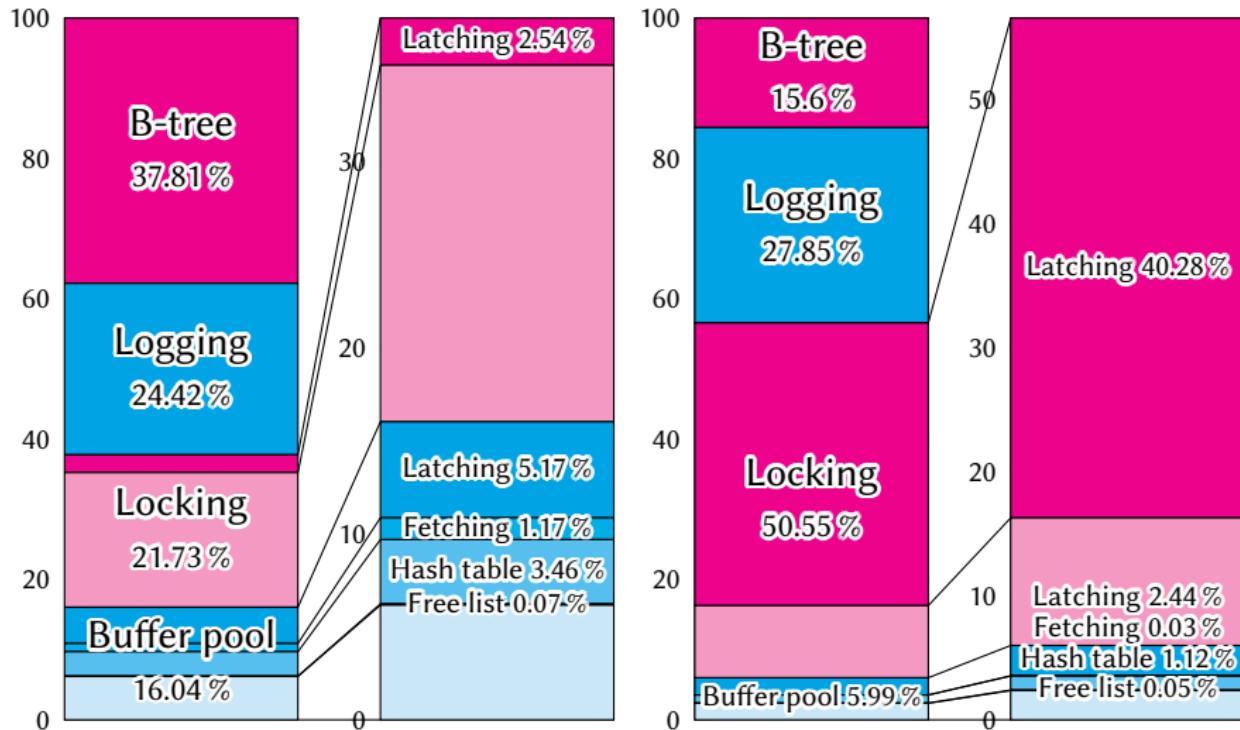


# Update-Only YCSB

## Update-Only YCSB

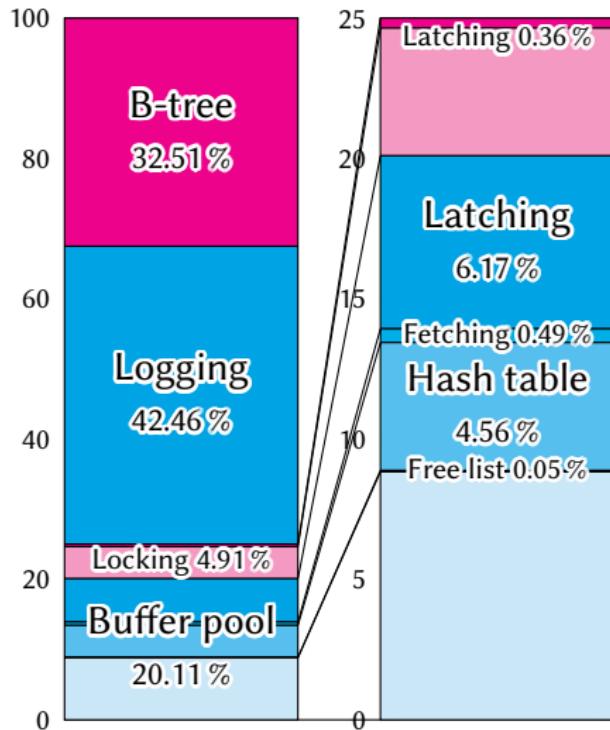


# Update-Only YCSB

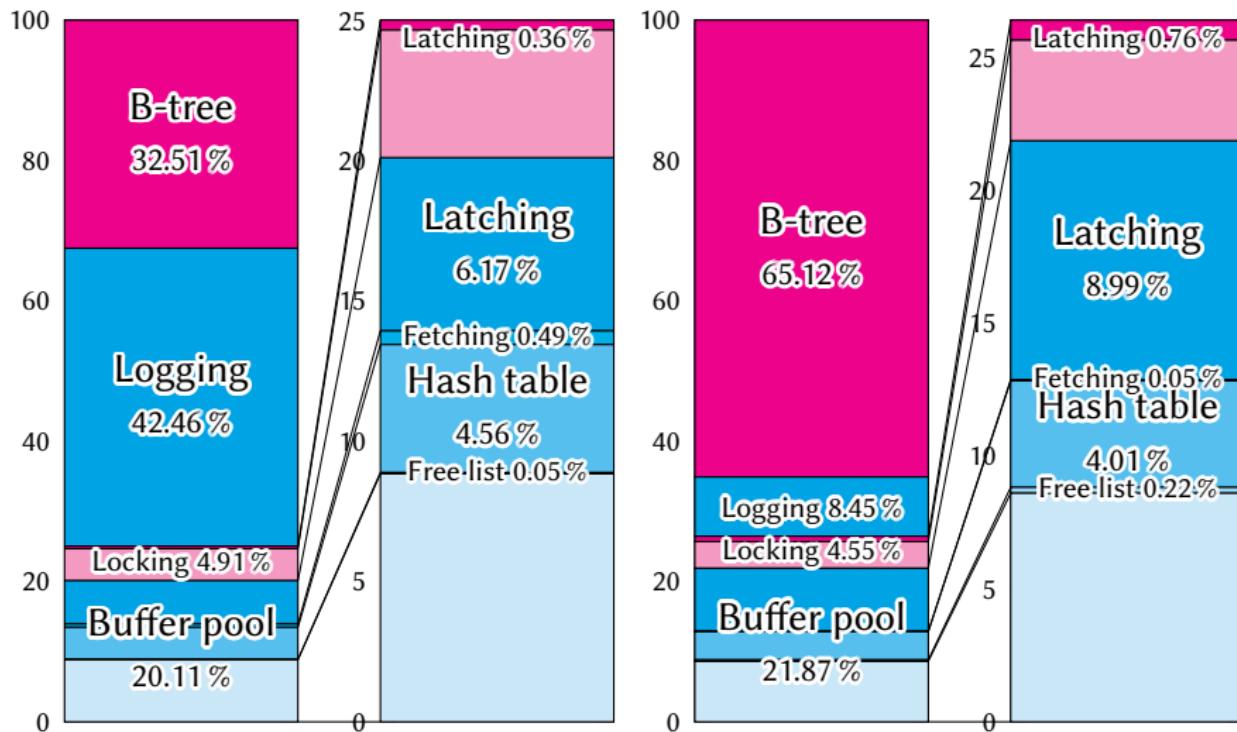


# TPC-C

# TPC-C



## TPC-C



# Conclusion

# Conclusion

## Complex OLTP (and OLAP)

The index structures (here Foster B-tree) are the most performance-critical components.

# Conclusion

## Complex OLTP (and OLAP)

The index structures (here Foster B-tree) are the most performance-critical components.

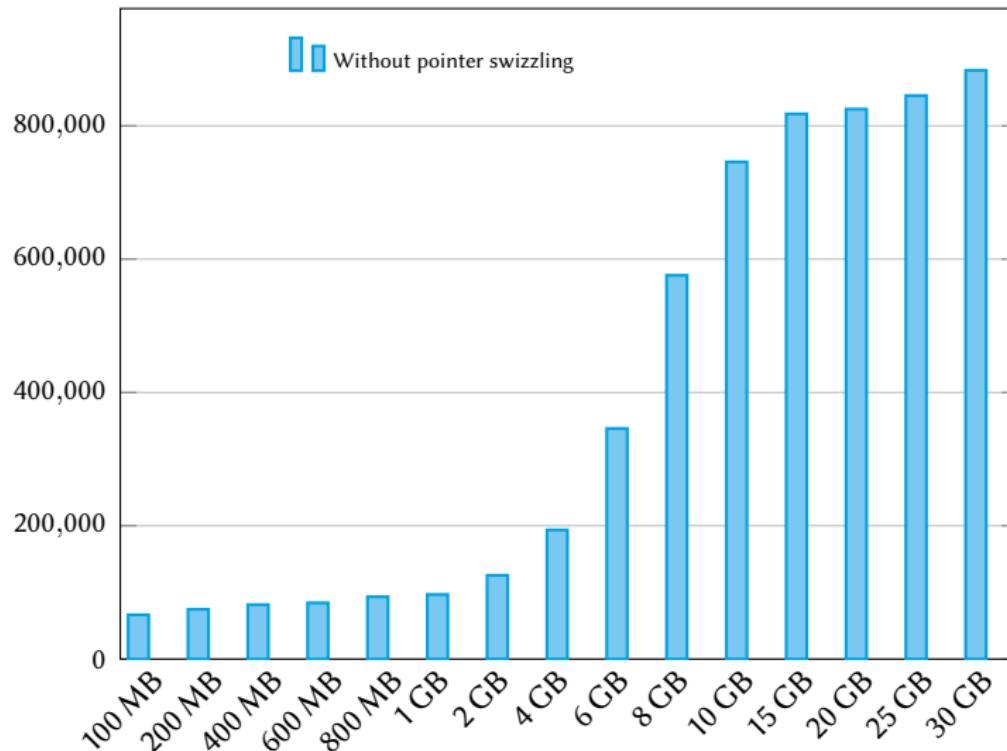
## Simple Workloads

A global latch in the transaction manager limits the performance.

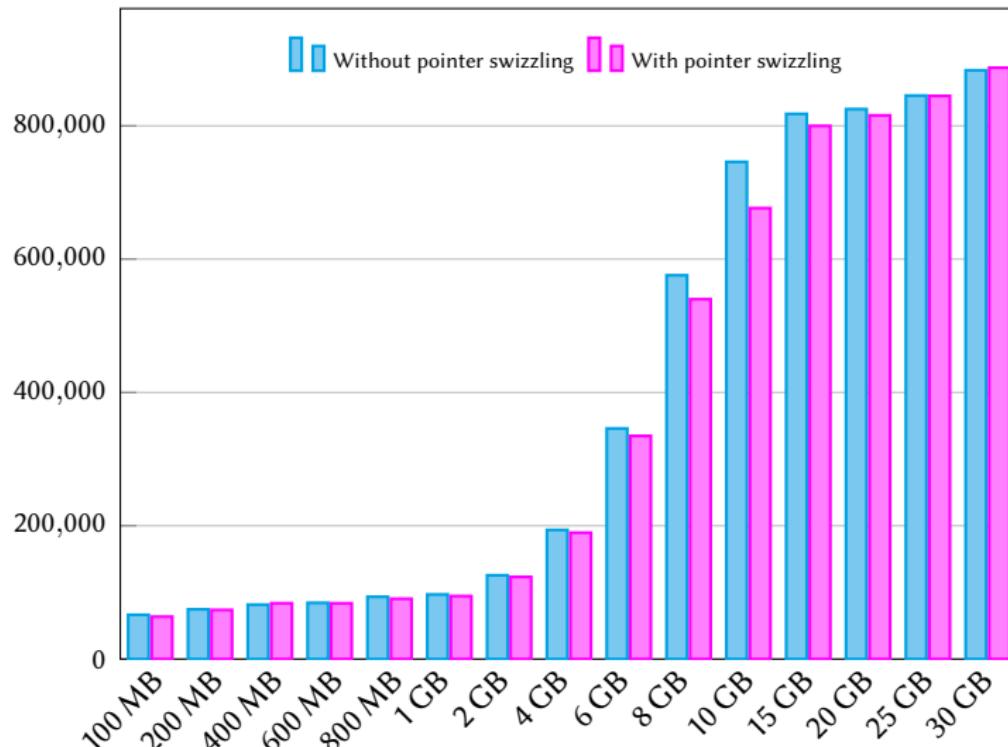
## Subsection 4

### Optimizations of OLTP Systems

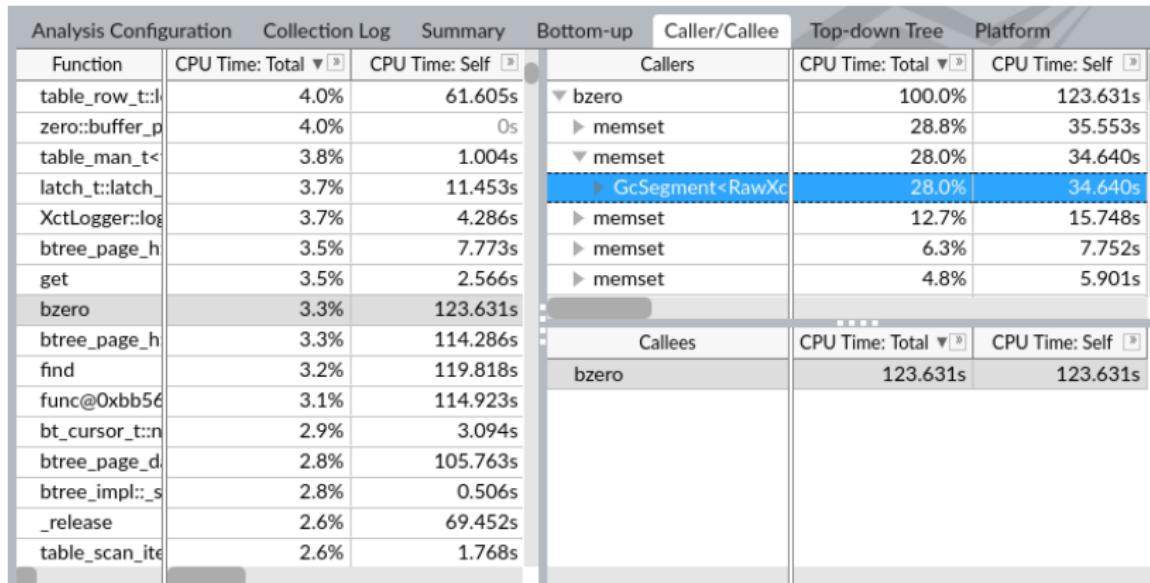
## Pointer Swizzling in the Buffer Pool [Gra+]



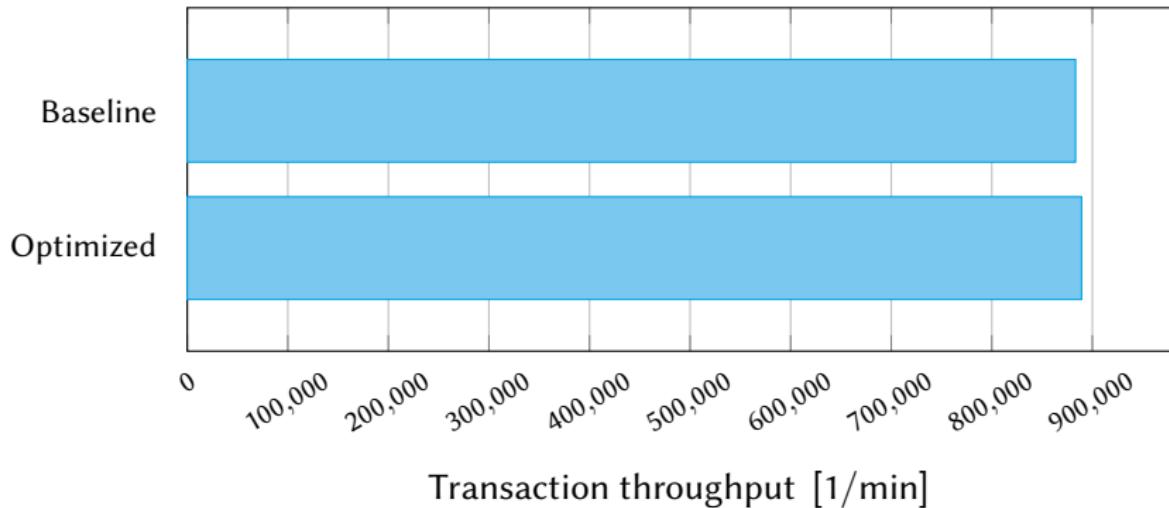
## Pointer Swizzling in the Buffer Pool [Gra+]



# System Call: bzero



## System Call: bzero



# References I

-  [László A. Bélády.](#) “A Study of Replacement Algorithms for Virtual-Storage Computer”.
-  [F. J. Corbató.](#) “A Paging Experiment with the Multics System”. In.
-  [Wolfgang Effelsberg and Theo Härdter.](#) “Principles of Database Buffer Management”.
-  [Goetz Graefe et al.](#) “In-Memory Performance for Big Data”.
-  [Stavros Harizopoulos et al.](#) “OLTP through the Looking Glass, and What We Found There”.
-  [Viktor Leis et al.](#) “LeanStore: In-Memory Data Management beyond Main Memory”.

## References II



Elizabeth J. O’Neil, Patrick E. O’Neil, and Gerhard Weikum.  
“The LRU-K Page Replacement Algorithm for Database Disk Buffering”.

# Your Turn to Ask ...

# Page Replacement Algorithm Classification by Bélády From [Bél]

# Page Replacement Algorithm Classification by Bélády From [Bél]

Class 1 No statistics

# Page Replacement Algorithm Classification by Bélády From [Bél]

Class 1 RANDOM, FIFO and FILO

# Page Replacement Algorithm Classification by Bélády From [Bél]

Class 1 RANDOM, FIFO and FILO

Class 2 Statistics about the latest references of *pages in the buffer pool*

# Page Replacement Algorithm Classification by Bélády From [Bél]

Class 1 RANDOM, FIFO and FILO

Class 2 LRU, MRU, LRU-K, SLRU, CLOCK, ZCLOCK, GCLOCK,  
DGCLOCK, LRD, LFU, LFU-Aging, LFUDA, 2Q, MQ  
and LeanStore

# Page Replacement Algorithm Classification by Bélády From [Bél]

Class 1 RANDOM, FIFO and FILO

Class 2 LRU, MRU, LRU-K, SLRU, CLOCK, ZCLOCK, GCLOCK,  
DGCLOCK, LRD, LFU, LFU-Aging, LFUDA, 2Q, MQ  
and LeanStore

Class 3 Statistics about each time *any page* was fetched from  
the database and each time it was evicted from the  
buffer pool

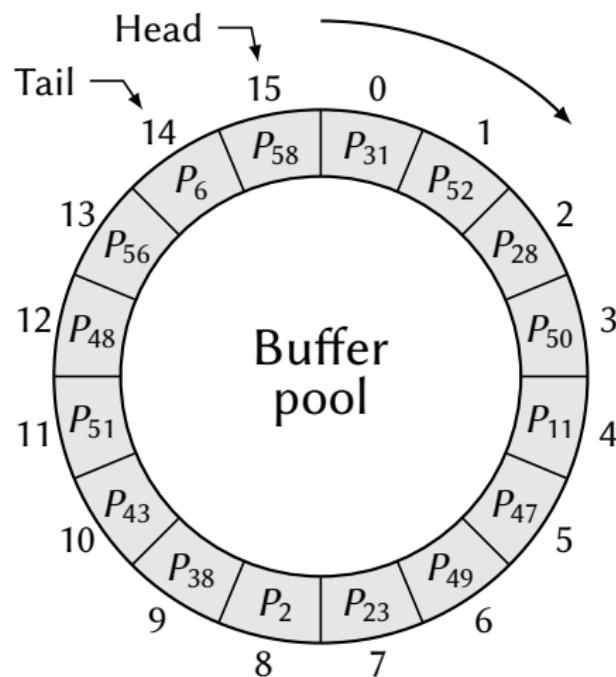
# Page Replacement Algorithm Classification by Bélády From [Bél]

Class 1 RANDOM, FIFO and FILO

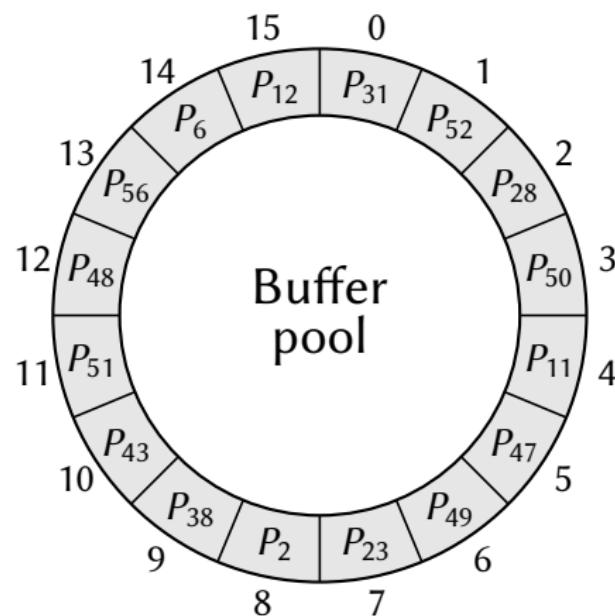
Class 2 LRU, MRU, LRU-K, SLRU, CLOCK, ZCLOCK, GCLOCK,  
DGCLOCK, LRD, LFU, LFU-Aging, LFUDA, 2Q, MQ  
and LeanStore

Class 3 ARC, CAR, CART, LIRS, CLOCK-Pro, DLIRS and  
CLOCK-Pro+

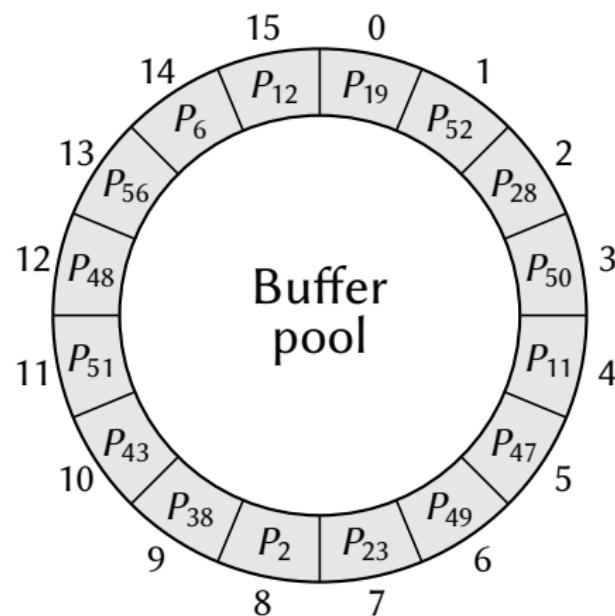
## LOOP Algorithm



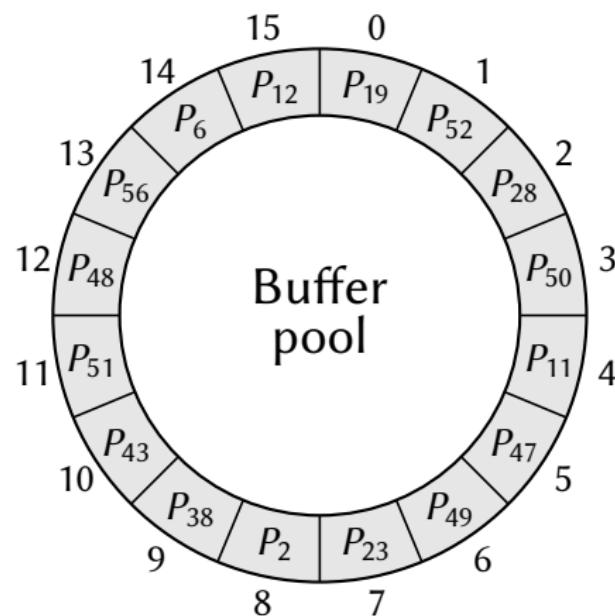
# LOOP Algorithm



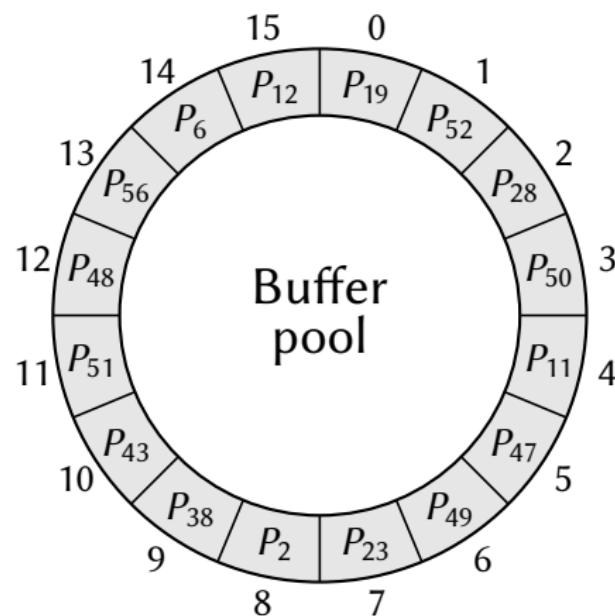
# LOOP Algorithm



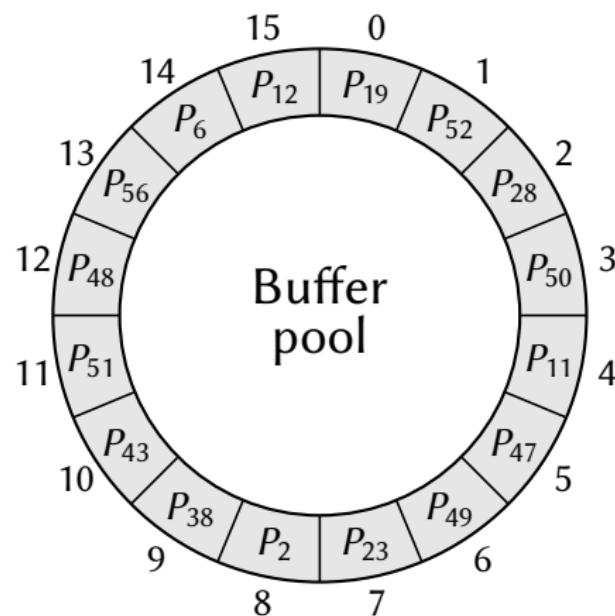
# LOOP Algorithm



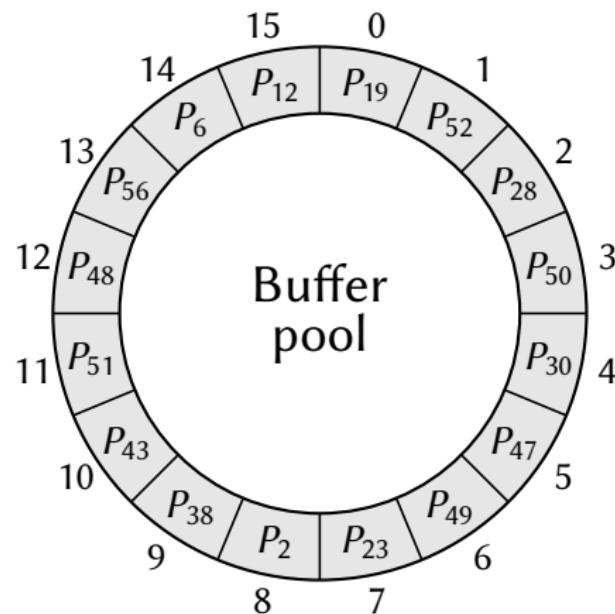
# LOOP Algorithm



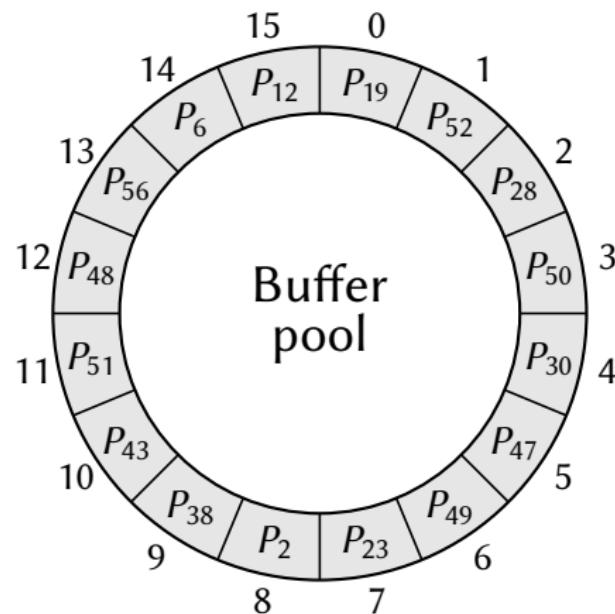
# LOOP Algorithm



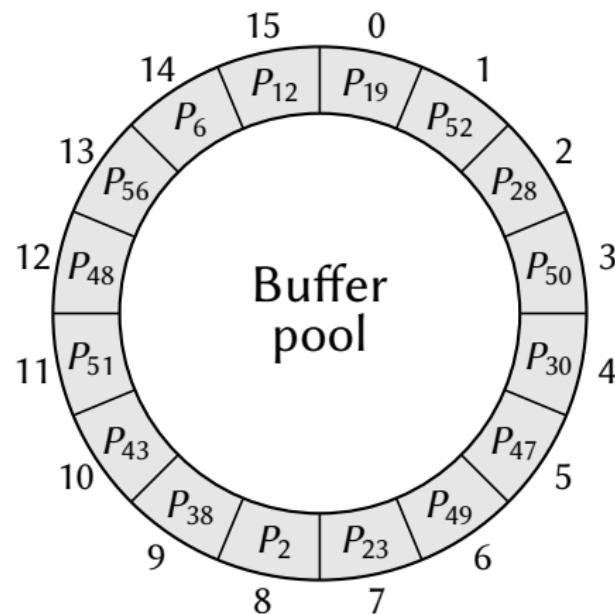
# LOOP Algorithm



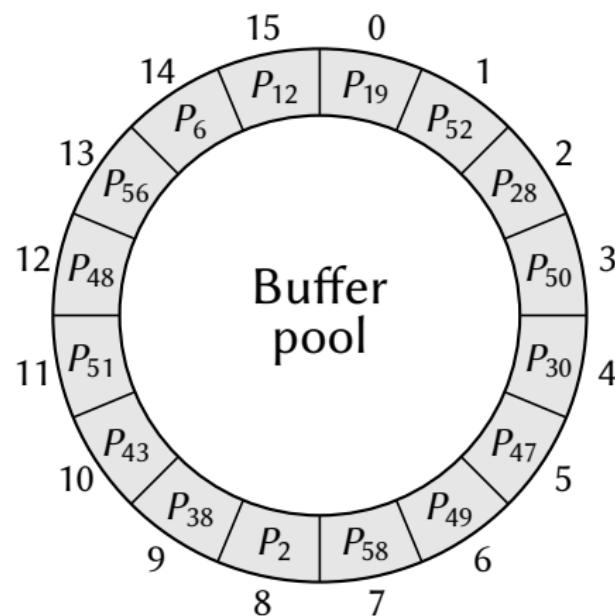
# LOOP Algorithm



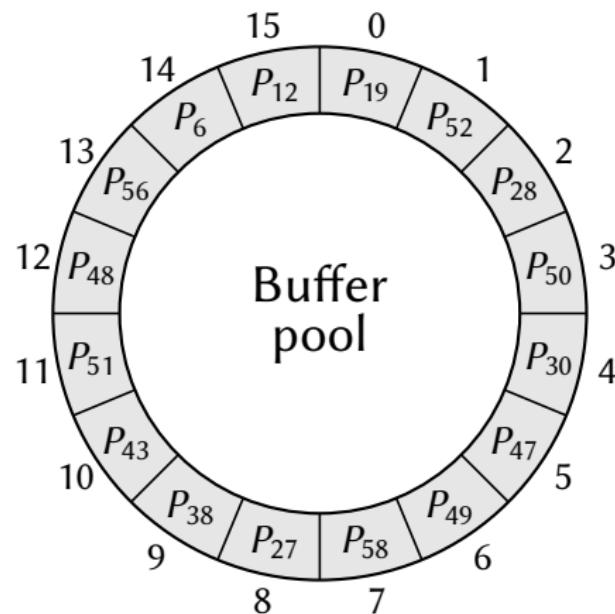
# LOOP Algorithm



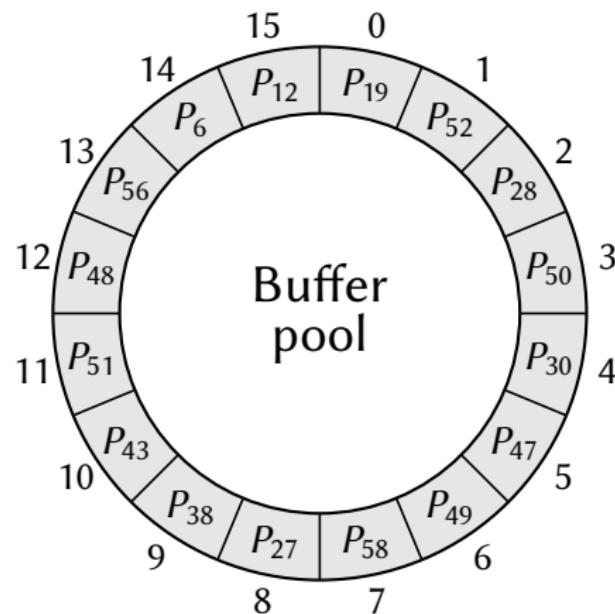
# LOOP Algorithm



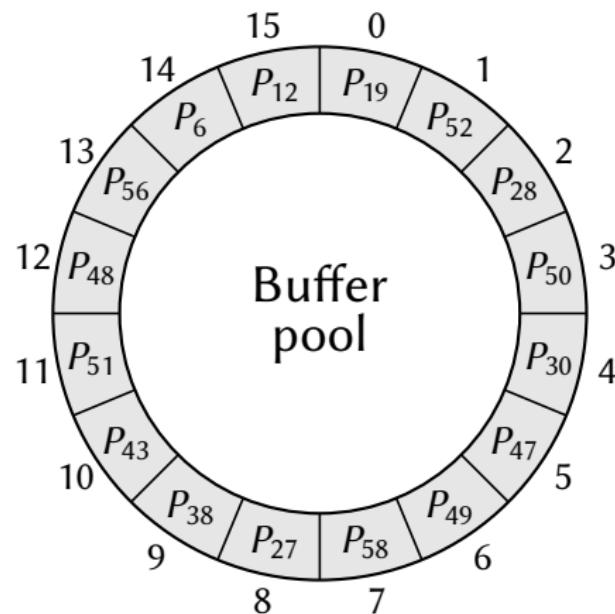
# LOOP Algorithm



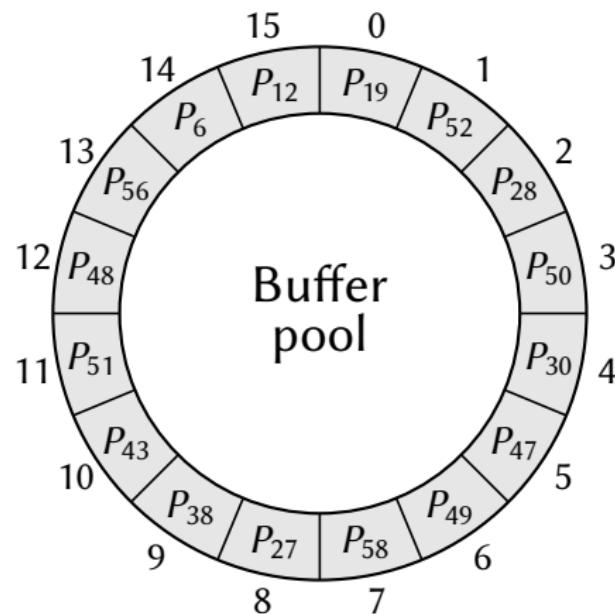
# LOOP Algorithm



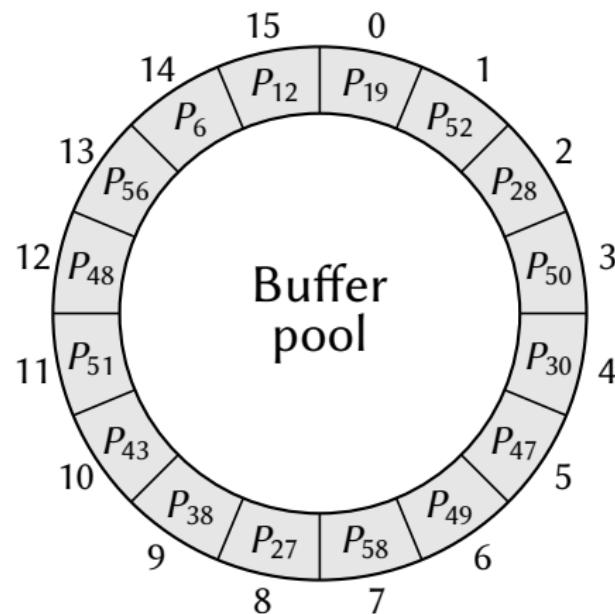
# LOOP Algorithm



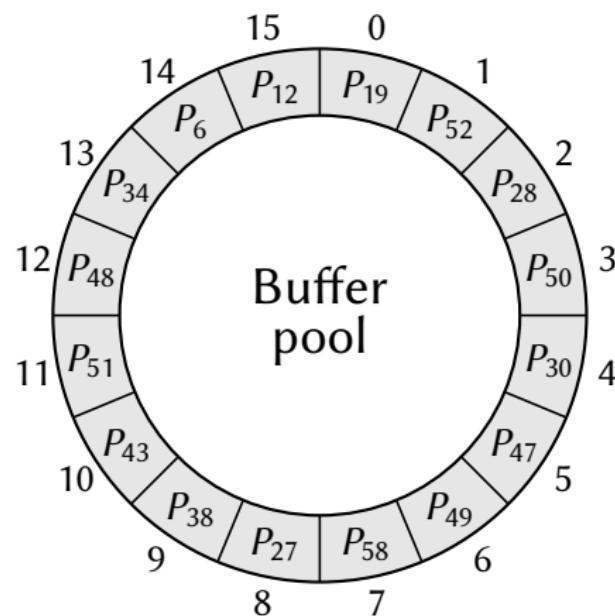
# LOOP Algorithm



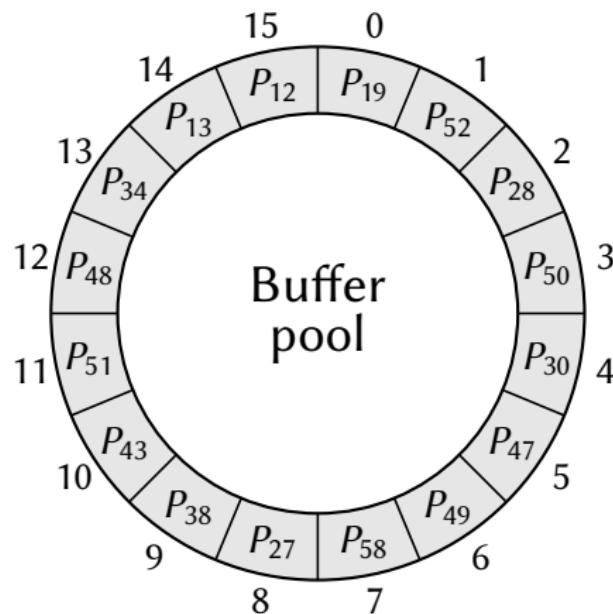
# LOOP Algorithm



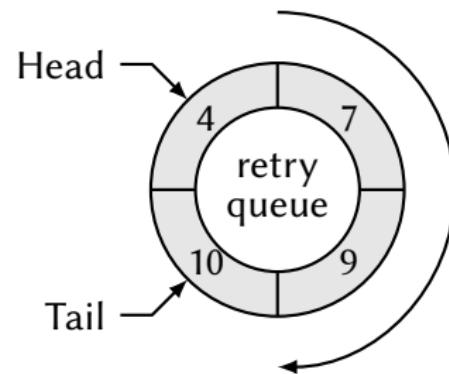
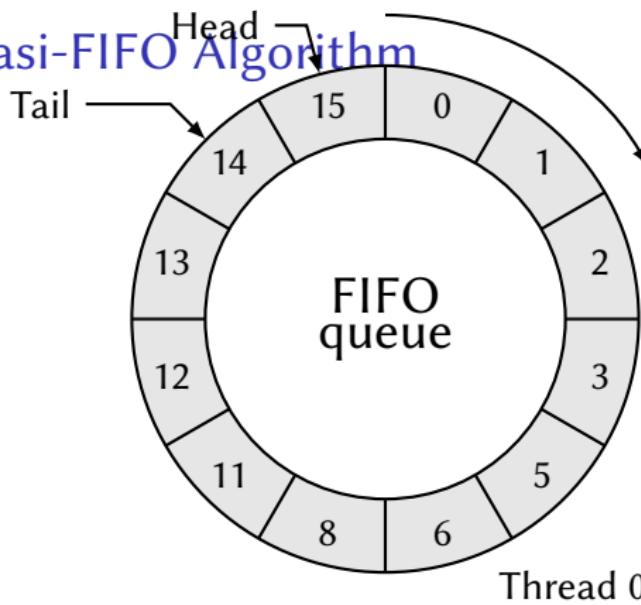
# LOOP Algorithm



# LOOP Algorithm



## Quasi-FIFO Algorithm



currentlyCheckingRetryQueue

currentQueueChecks

15

notExplicitlyEvictedList

0	1	2	3	4	5	6	7	8	9	10	11	12	13	14	15
---	---	---	---	---	---	---	---	---	---	----	----	----	----	----	----

# LRU-K Implementations

# LRU-K Implementations

- ▶ *Hash-Map-Doubly-Linked-List*

# LRU-K Implementations

- ▶ *Hash-Map-Doubly-Linked-List*
  - ▶ Entries in the queue are reference IDs instead of frame indexes

# LRU-K Implementations

- ▶ *Hash-Map-Doubly-Linked-List*
  - ▶ Entries in the queue are reference IDs instead of frame indexes
  - ▶ Ghost references are in a segment at the head of the queue

# LRU-K Implementations

- ▶ *Hash-Map-Doubly-Linked-List*
  - ▶ Entries in the queue are reference IDs instead of frame indexes
  - ▶ Ghost references<sup>1</sup> are in a segment at the head of the queue

---

<sup>1</sup> Ghost references are used for pages with  $< k$  references

# LRU-K Implementations

- ▶ *Hash-Map-Doubly-Linked-List*
  - ▶ Entries in the queue are reference IDs instead of frame indexes
  - ▶ Ghost references<sup>1</sup> are in a segment at the head of the queue
- ▶ *Timestamp-Sorting*

---

<sup>1</sup> Ghost references are used for pages with  $< k$  references

# LRU-K Implementations

- ▶ *Hash-Map-Doubly-Linked-List*
  - ▶ Entries in the queue are reference IDs instead of frame indexes
  - ▶ Ghost references<sup>1</sup> are in a segment at the head of the queue
- ▶ *Timestamp-Sorting*
  - ▶  $k$  timestamps per buffer frame

---

<sup>1</sup> Ghost references are used for pages with  $< k$  references

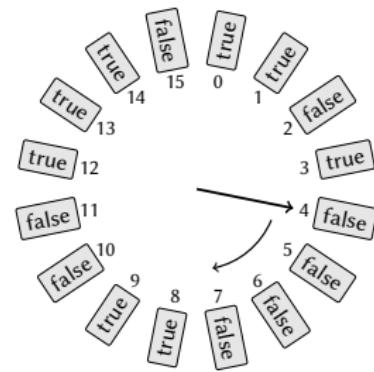
# LRU-K Implementations

- ▶ *Hash-Map-Doubly-Linked-List*
  - ▶ Entries in the queue are reference IDs instead of frame indexes
  - ▶ Ghost references<sup>1</sup> are in a segment at the head of the queue
- ▶ *Timestamp-Sorting*
  - ▶  $k$  timestamps per buffer frame
  - ▶ Ghost references are represented by timestamps very early in the past

---

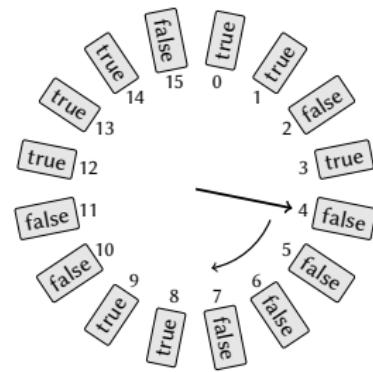
<sup>1</sup> Ghost references are used for pages with  $< k$  references

# CLOCK Variants



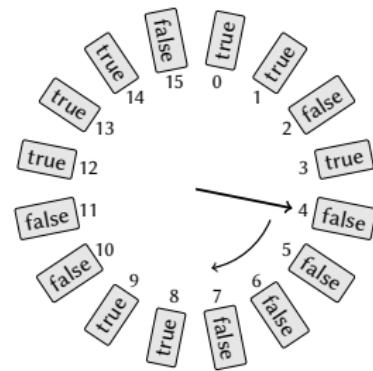
# CLOCK Variants

Fix Set usage-bit on *page fix*



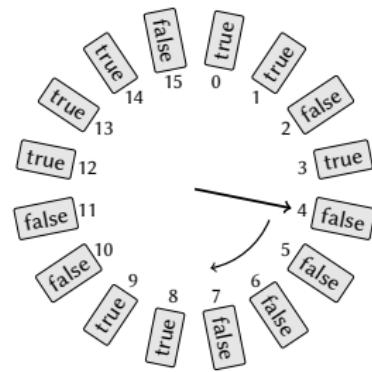
# CLOCK Variants

Fix Set usage-bit on *page fix*  
Unfix Set usage-bit on *page unfix*



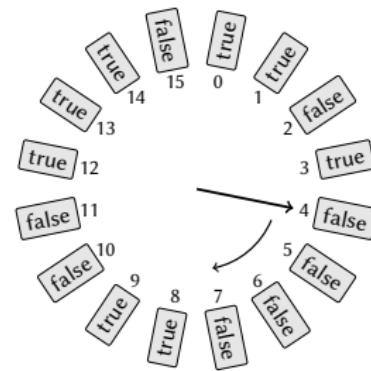
# CLOCK Variants

- Fix Set usage-bit on *page fix*
- Unfix Set usage-bit on *page unfix*
- FixUnfix Set usage-bit on *page fix* and on *page unfix*

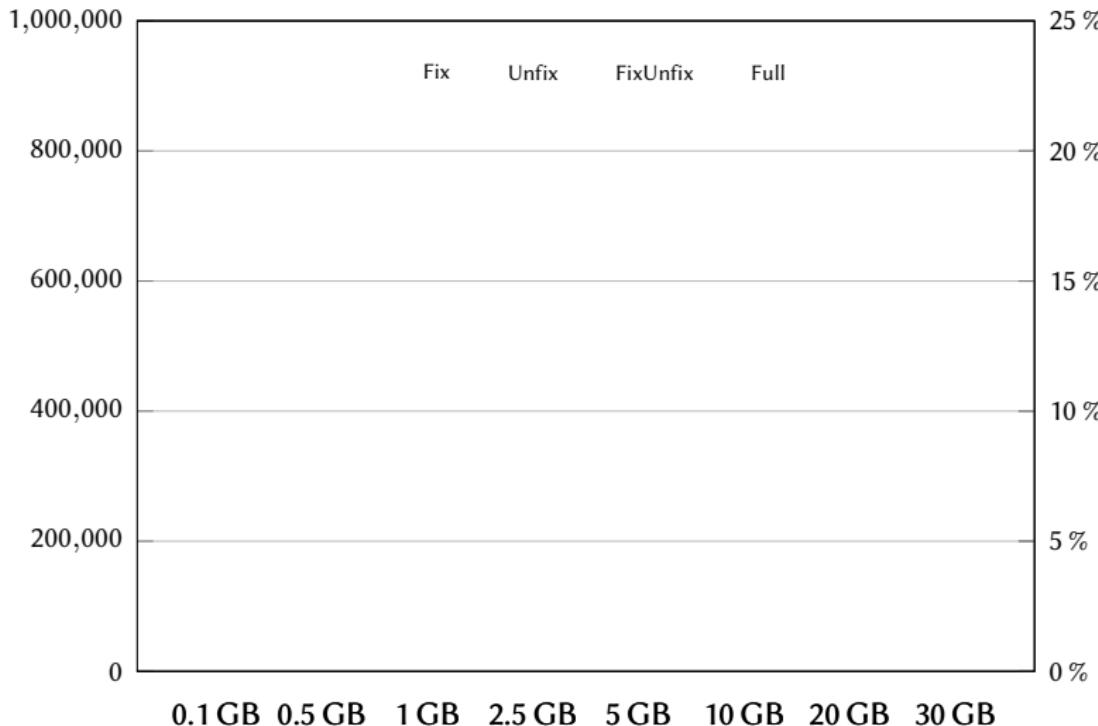


# CLOCK Variants

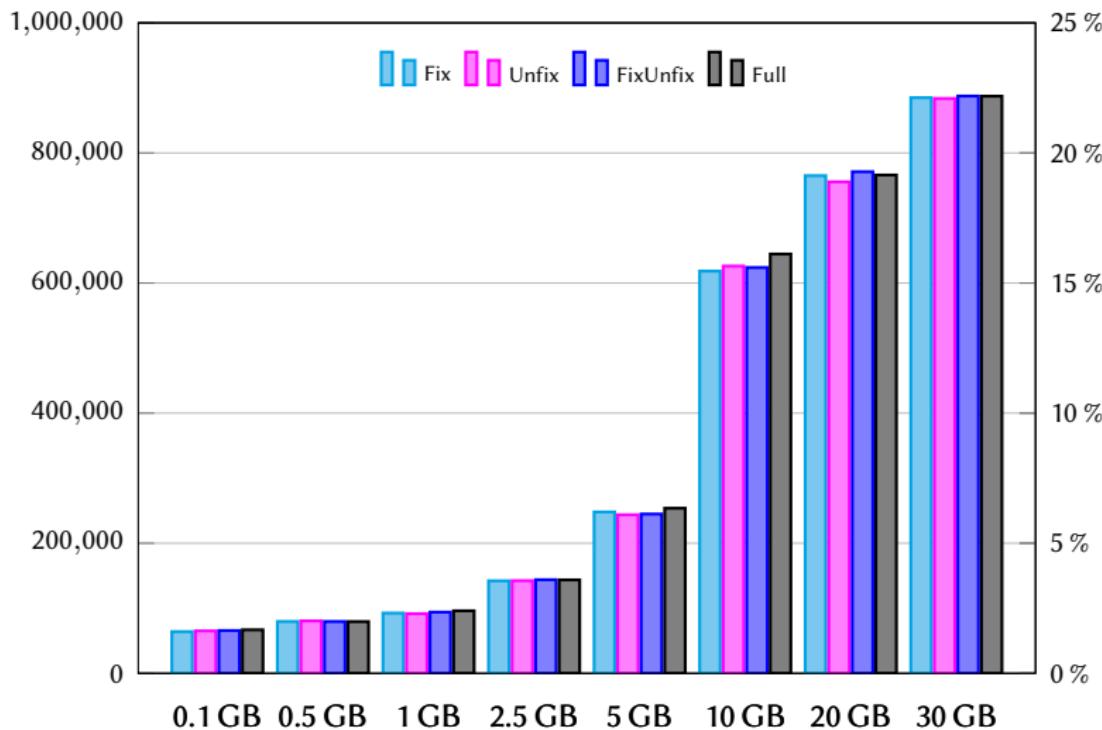
- Fix Set usage-bit on *page fix*
- Unfix Set usage-bit on *page unfix*
- FixUnfix Set usage-bit on *page fix* and on *page unfix*
- Full Like FixUnfix but also set usage-bit when page cannot be evicted



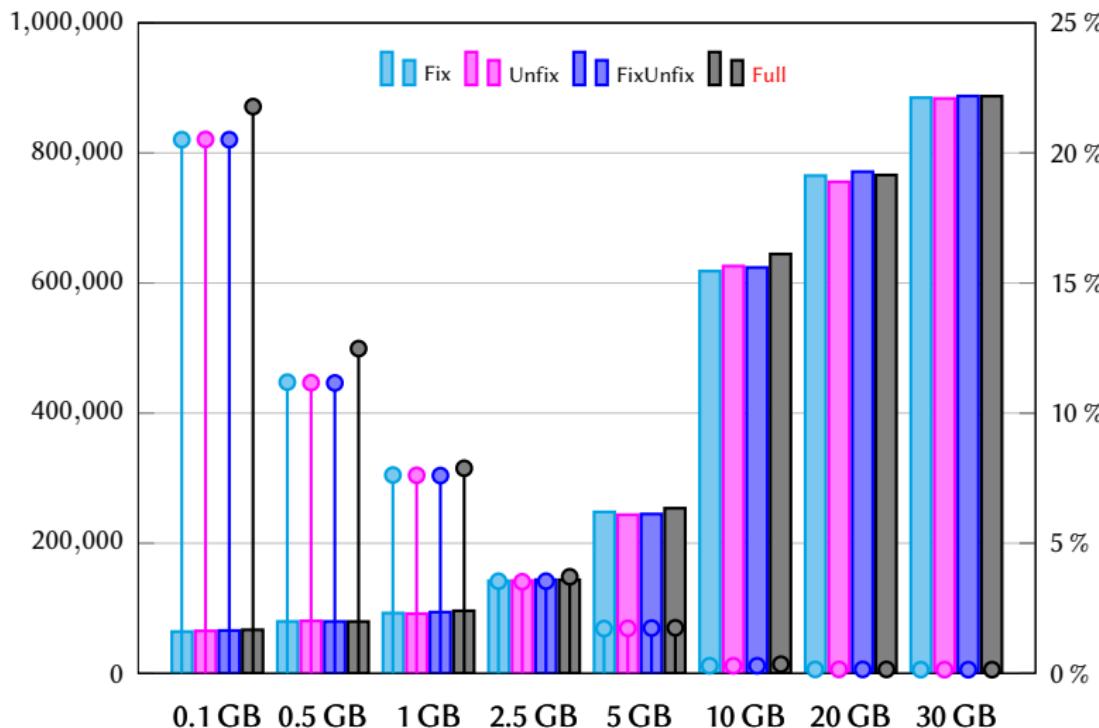
# Performance Evaluation



# Performance Evaluation



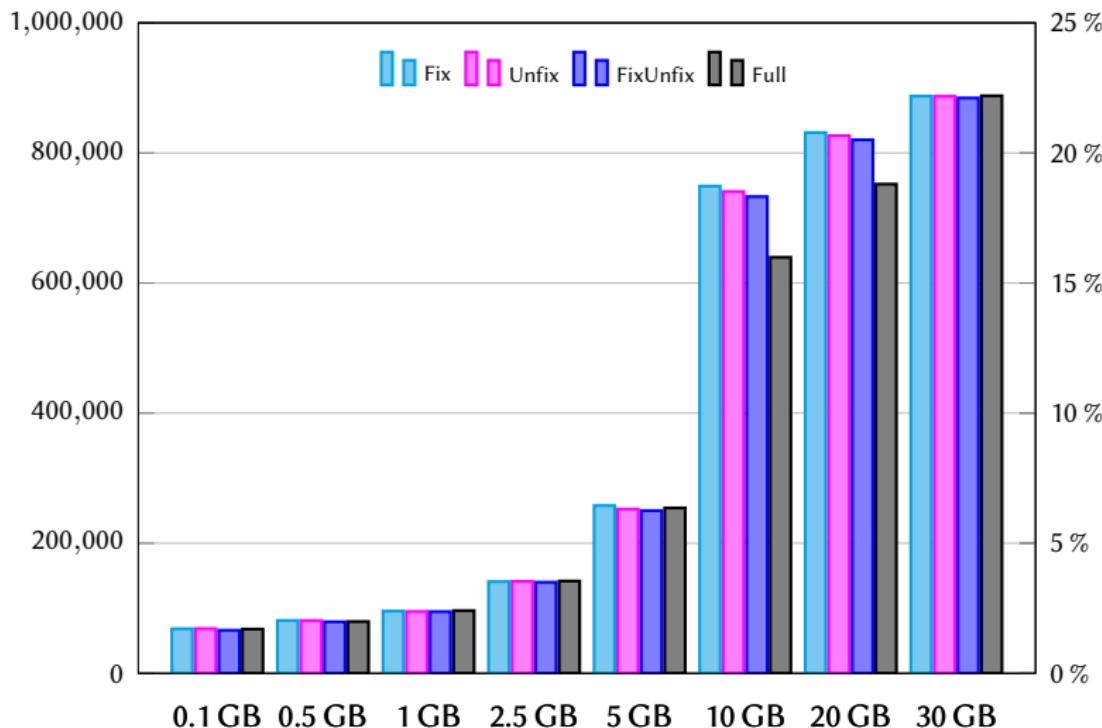
# Performance Evaluation



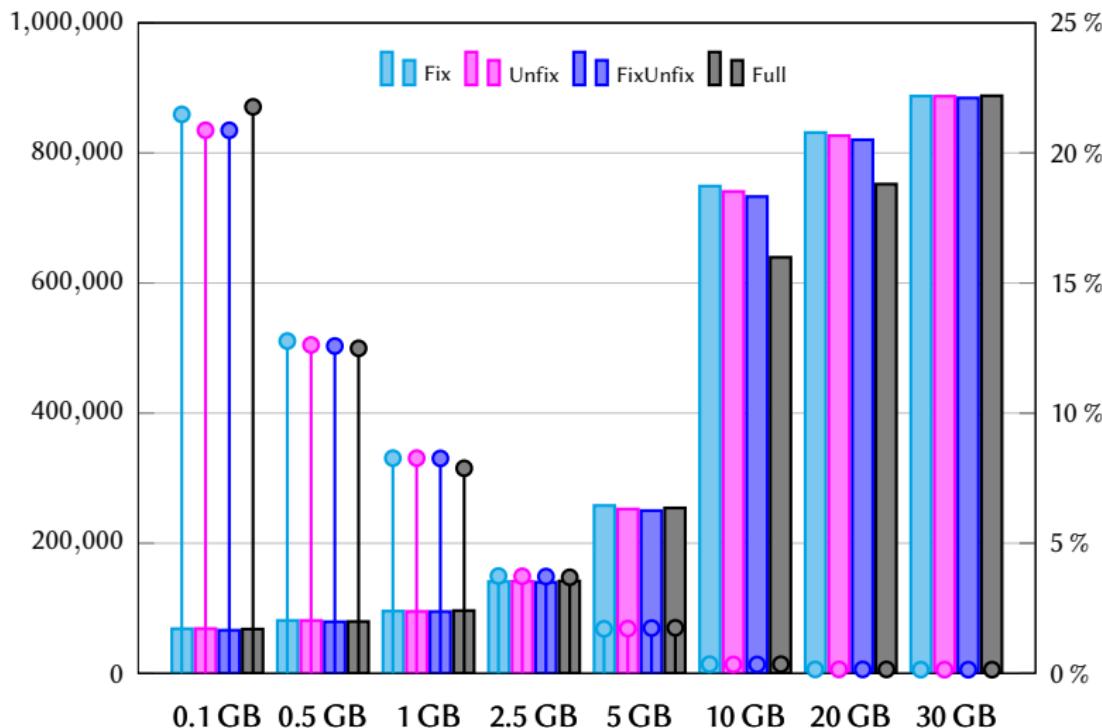
## ZCLOCK Variants

- Fix Set usage-bit on *page fix*
- Unfix Set usage-bit on *page unfix*
- FixUnfix Set usage-bit on *page fix* and on *page unfix*
- Full Like FixUnfix but also set usage-bit when page cannot be evicted

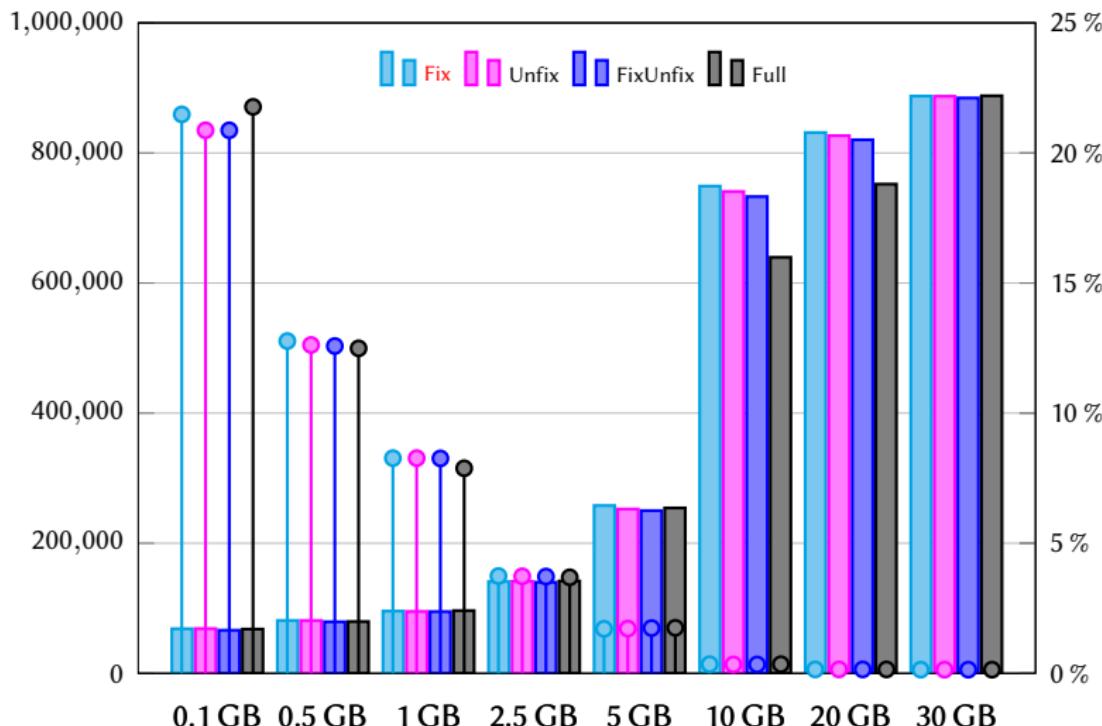
# Performance Evaluation



# Performance Evaluation



# Performance Evaluation



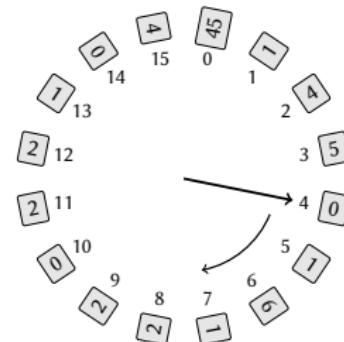
## GCLOCK-V1 Variants

### Usage-Count

$$\text{(initially)} \quad UC(p) = F$$

$$\text{(increase)} \quad UC(p) = UC(p) + R$$

$$\text{(decrease)} \quad UC(p) = UC(p) - 1$$



## GCLOCK-V1 Variants

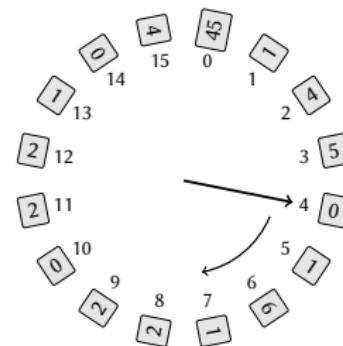
### Usage-Count

$$\text{(initially)} \quad UC(p) = F$$

$$\text{(increase)} \quad UC(p) = UC(p) + R$$

$$\text{(decrease)} \quad UC(p) = UC(p) - 1$$

Fix Increase usage-count on *page fix*



## GCLOCK-V1 Variants

### Usage-Count

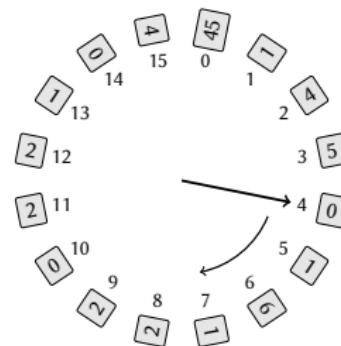
(initially)  $UC(p) = F$

(increase)  $UC(p) = UC(p) + R$

(decrease)  $UC(p) = UC(p) - 1$

**Fix** Increase usage-count on *page fix*

**FixFull** Like **Fix** but also increase usage-count when page cannot be evicted



## GCLOCK-V1 Variants

### Usage-Count

$$\text{(initially)} \quad UC(p) = F$$

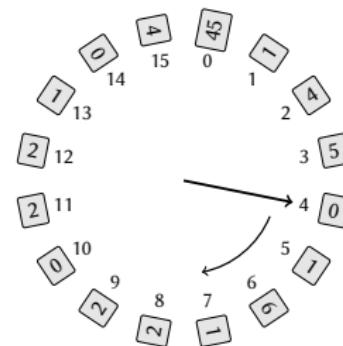
$$\text{(increase)} \quad UC(p) = UC(p) + R$$

$$\text{(decrease)} \quad UC(p) = UC(p) - 1$$

**Fix** Increase usage-count on *page fix*

**FixFull** Like **Fix** but also increase usage-count when page cannot be evicted

**Unfix** Increase usage-count on *page unfix*



# GCLOCK-V1 Variants

## Usage-Count

$$\text{(initially)} \quad UC(p) = F$$

$$\text{(increase)} \quad UC(p) = UC(p) + R$$

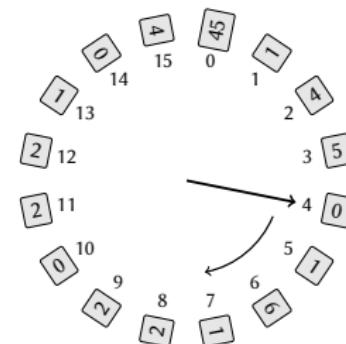
$$\text{(decrease)} \quad UC(p) = UC(p) - 1$$

**Fix** Increase usage-count on *page fix*

**FixFull** Like **Fix** but also increase usage-count when page cannot be evicted

**Unfix** Increase usage-count on *page unfix*

**UnfixFull** Like **Unfix** but also increase usage-count when page cannot be evicted



# GCLOCK-V1 Variants

## Usage-Count

(initially)       $UC(p) = F$

(increase)       $UC(p) = UC(p) + R$

(decrease)       $UC(p) = UC(p) - 1$

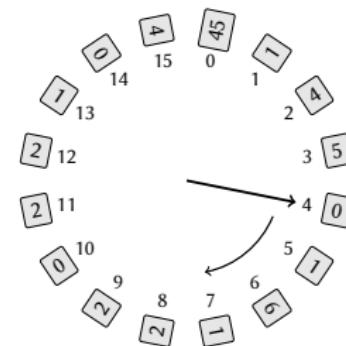
**Fix** Increase usage-count on *page fix*

**FixFull** Like **Fix** but also increase usage-count when page cannot be evicted

**Unfix** Increase usage-count on *page unfix*

**UnfixFull** Like **Unfix** but also increase usage-count when page cannot be evicted

$$F = 25, R = 5$$



# GCLOCK-V1 Variants

## Usage-Count

$$\text{(initially)} \quad UC(p) = F$$

$$\text{(increase)} \quad UC(p) = UC(p) + R$$

$$\text{(decrease)} \quad UC(p) = UC(p) - 1$$

**Fix** Increase usage-count on *page fix*

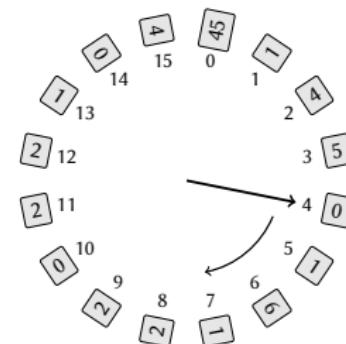
**FixFull** Like **Fix** but also increase usage-count when page cannot be evicted

**Unfix** Increase usage-count on *page unfix*

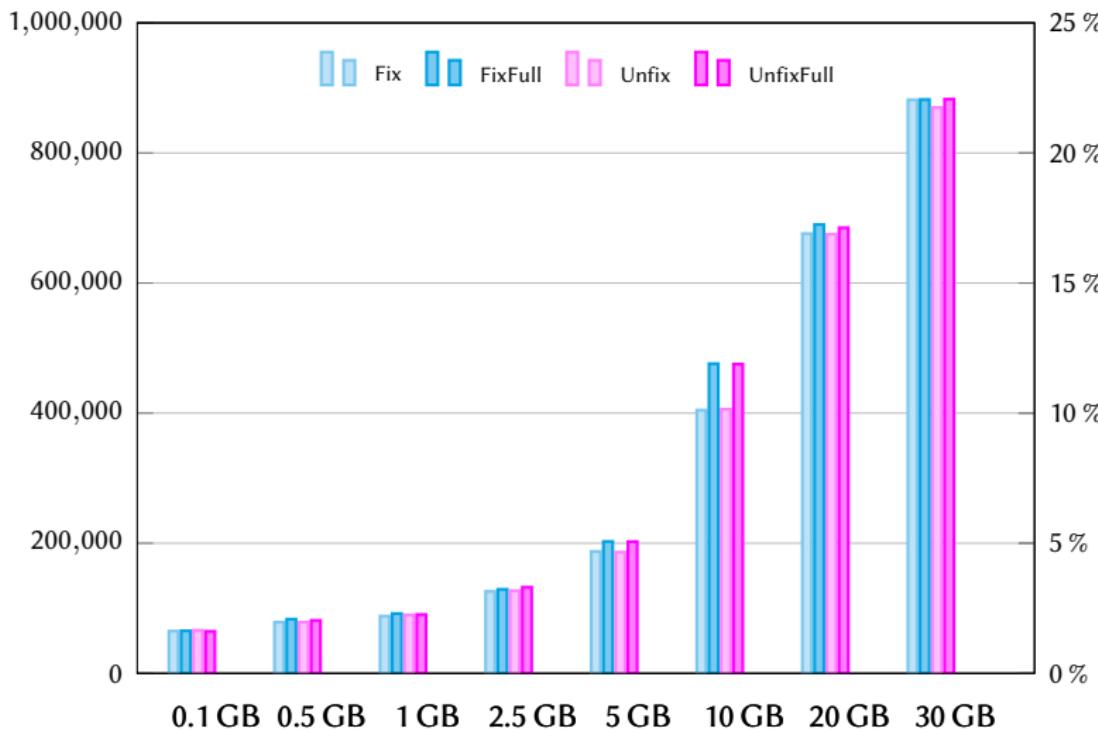
**UnfixFull** Like **Unfix** but also increase usage-count when page cannot be evicted

$$F = 25, R = 5$$

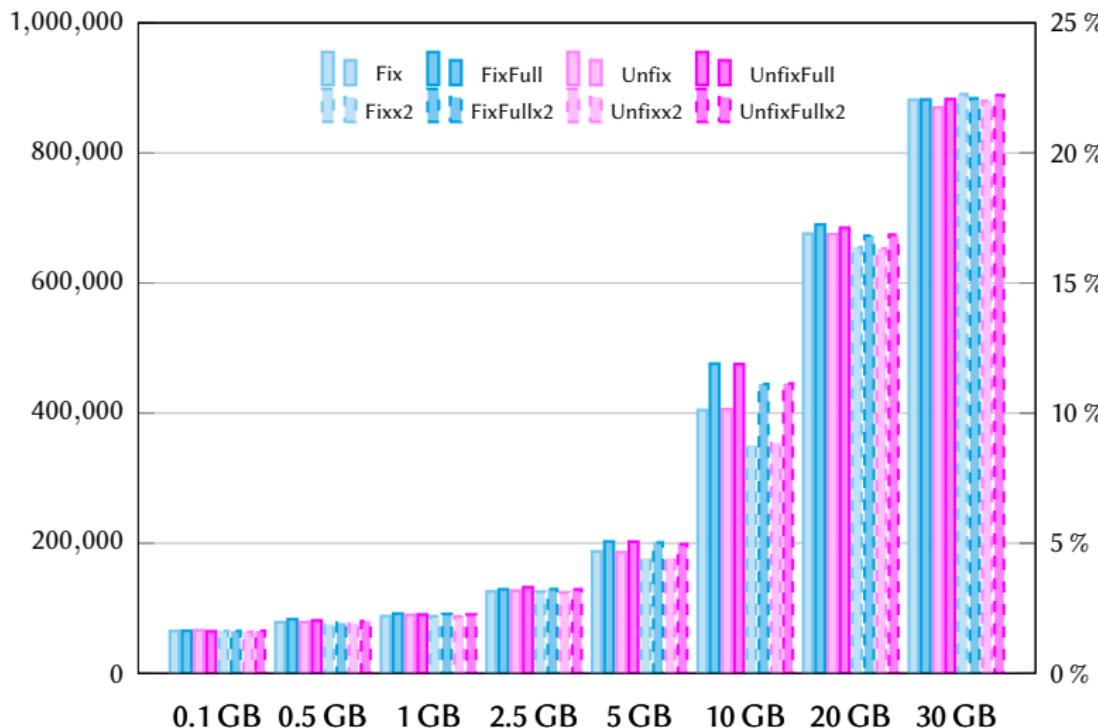
$$x2 \quad F = 50, R = 10$$



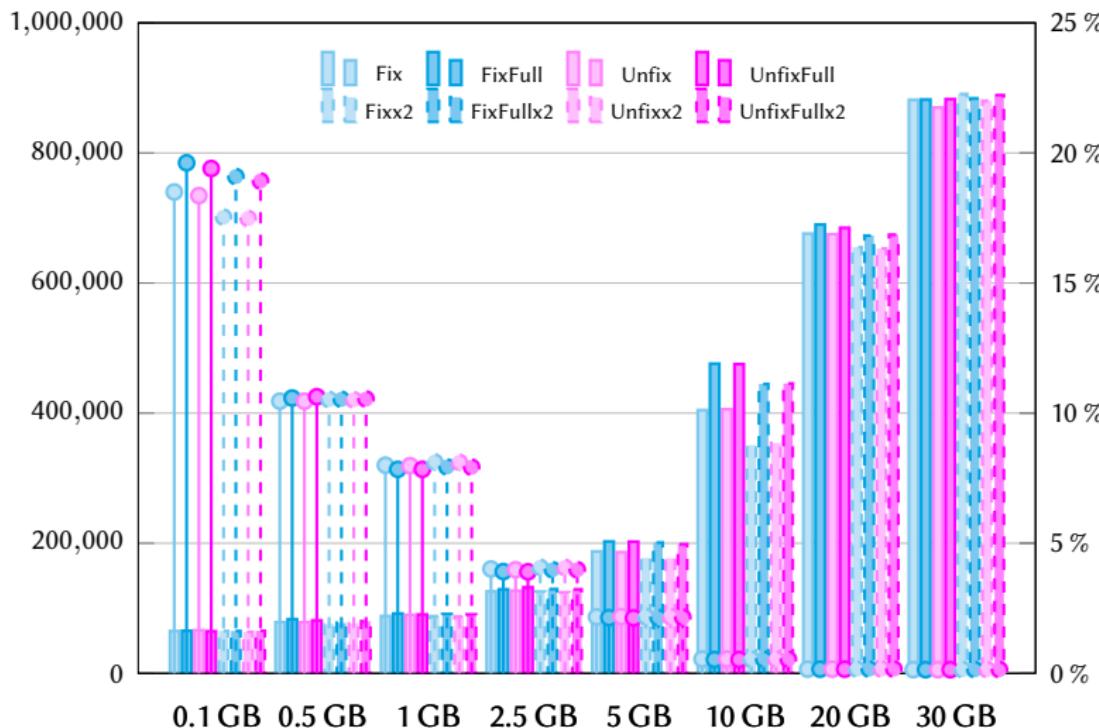
# Performance Evaluation



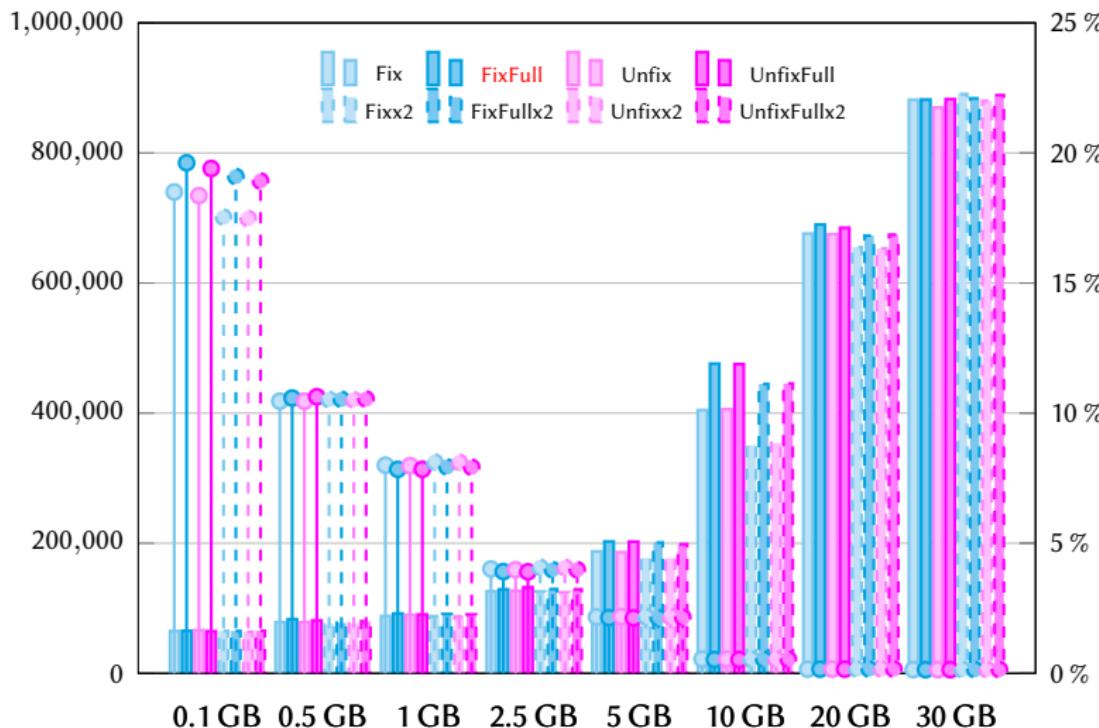
# Performance Evaluation



# Performance Evaluation



# Performance Evaluation



## GCLOCK-V2 Variants

**Fix** Set usage-count on *page fix*

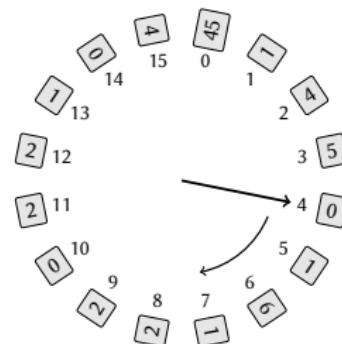
**FixFull** Like **Fix** but also set usage-count when page cannot be evicted

**Unfix** Set usage-count on *page unfix*

**UnfixFull** Like **Unfix** but also set usage-count when page cannot be evicted

$$F = 25, R = 5$$

$$\times 2 \quad F = 50, R = 10$$



# GCLOCK-V2 Variants

## Usage-Count

(initially)  $UC(p) = F$

(set)  $UC(p) = R$

(decrease)  $UC(p) = UC(p) - 1$

**Fix** Set usage-count on *page fix*

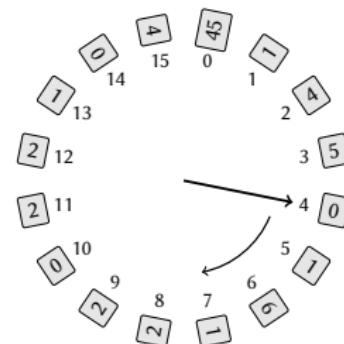
**FixFull** Like **Fix** but also set usage-count when page cannot be evicted

**Unfix** Set usage-count on *page unfix*

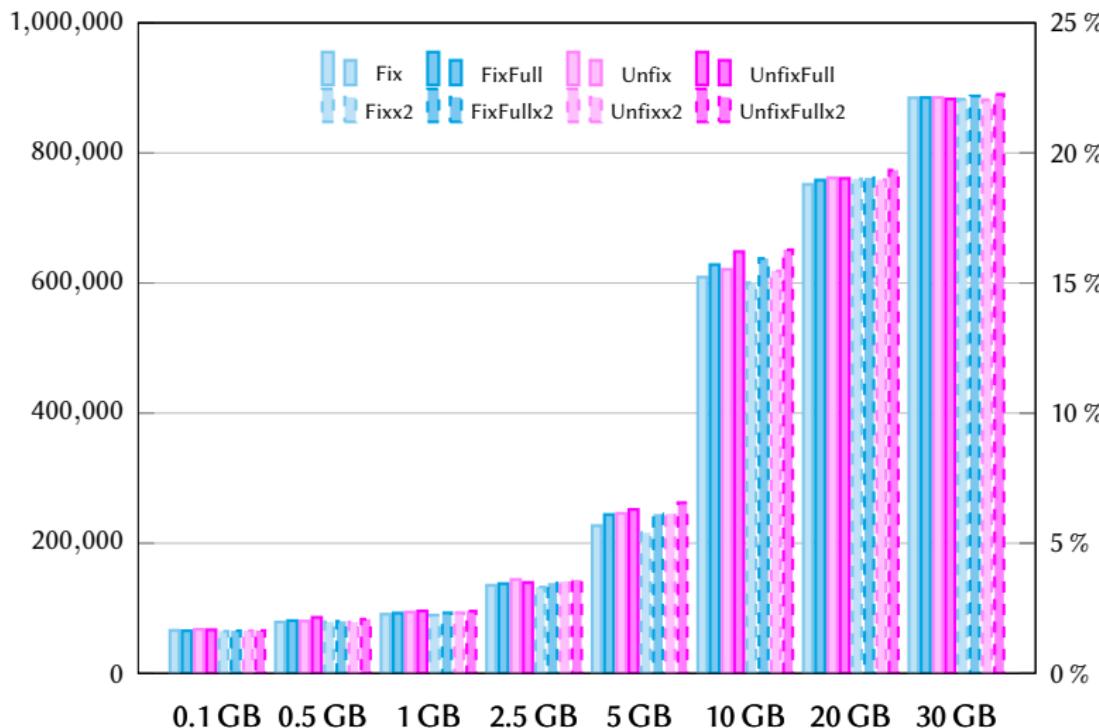
**UnfixFull** Like **Unfix** but also set usage-count when page cannot be evicted

$$F = 25, R = 5$$

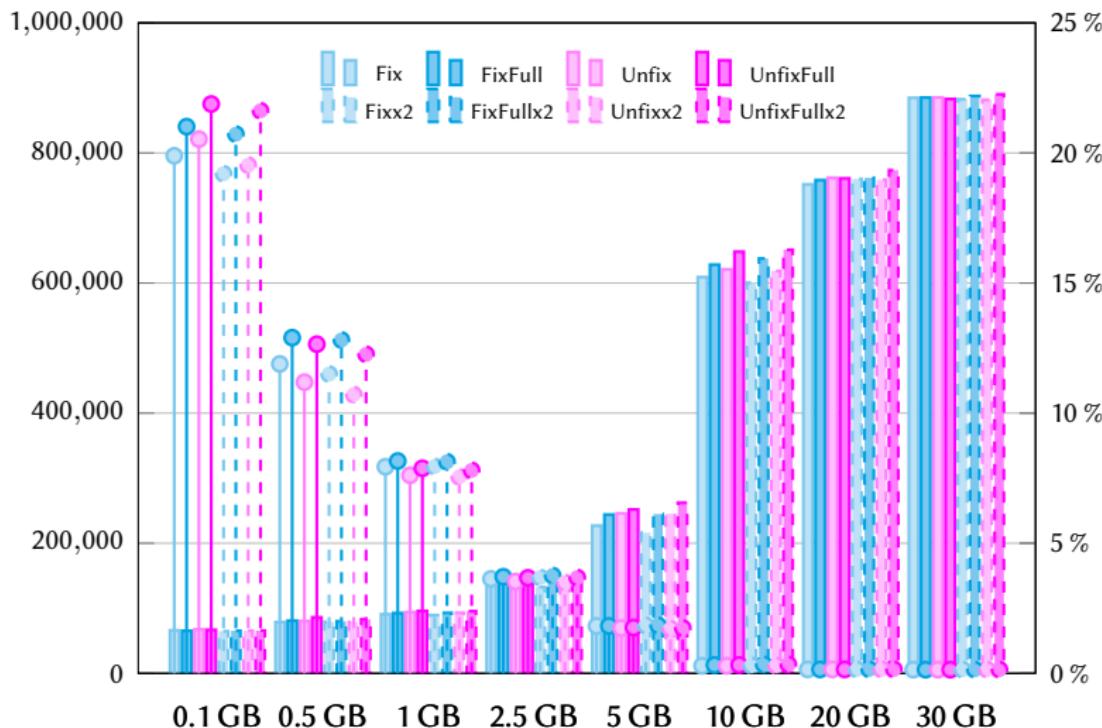
$$\times 2 \quad F = 50, R = 10$$



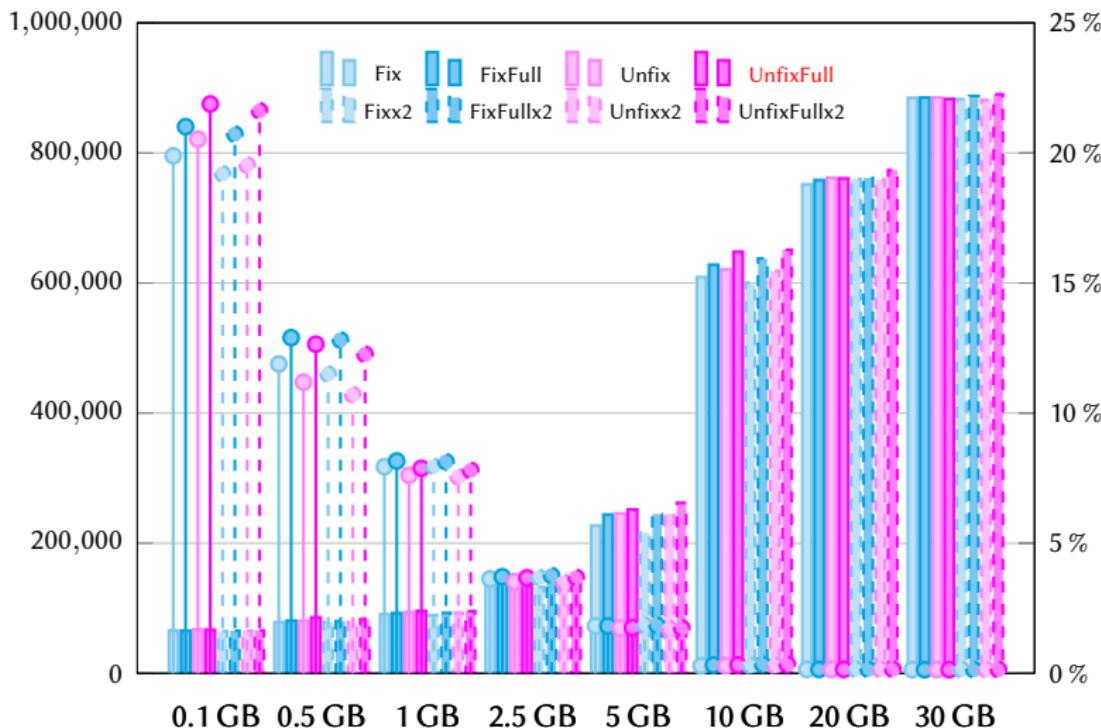
# Performance Evaluation



# Performance Evaluation



# Performance Evaluation



## DGCLOCK-V1 Variants

### Usage-Count

$$\text{(initially)} \quad UC(p) = F_p$$

$$\text{(increase)} \quad UC(p) = UC(p) + R_p$$

$$\text{(decrease)} \quad UC(p) = UC(p) - 1$$

## DGCLOCK-V1 Variants

### Usage-Count

$$\text{(initially)} \quad UC(p) = F_p$$

$$\text{(increase)} \quad UC(p) = UC(p) + R_p$$

$$\text{(decrease)} \quad UC(p) = UC(p) - 1$$

**Fix** Increase usage-count on *page fix*

**FixFull** Like **Fix** but also increase usage-count when page cannot be evicted

## DGCLOCK-V1 Variants

### Usage-Count

$$\text{(initially)} \quad UC(p) = F_p$$

$$\text{(increase)} \quad UC(p) = UC(p) + R_p$$

$$\text{(decrease)} \quad UC(p) = UC(p) - 1$$

Fix Increase usage-count on *page fix*

FixFull Like Fix but also increase usage-count when page cannot be evicted

$$F_{\text{non-B-tree}} = 25, \quad F_{\text{root}} = 25, \quad F_{\text{root-1}} = 10, \quad F_{\text{root-}\geq 2} = 5,$$

$$R_{\text{non-B-tree}} = 5, \quad R_{\text{root}} = 5, \quad R_{\text{root-1}} = 2, \quad R_{\text{root-}\geq 2} = 1$$

## DGCLOCK-V1 Variants

### Usage-Count

$$\text{(initially)} \quad UC(p) = F_p$$

$$\text{(increase)} \quad UC(p) = UC(p) + R_p$$

$$\text{(decrease)} \quad UC(p) = UC(p) - 1$$

Fix Increase usage-count on *page fix*

FixFull Like Fix but also increase usage-count when page cannot be evicted

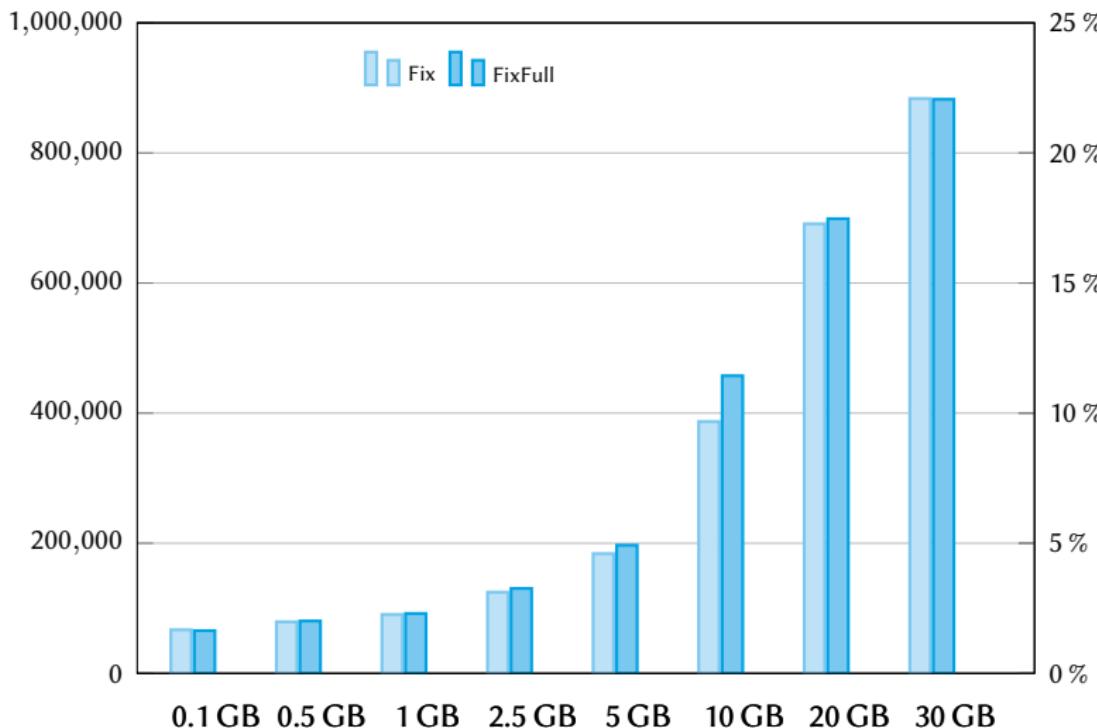
$$F_{\text{non-B-tree}} = 25, F_{\text{root}} = 25, F_{\text{root-1}} = 10, F_{\text{root-}\geq 2} = 5,$$

$$R_{\text{non-B-tree}} = 5, R_{\text{root}} = 5, R_{\text{root-1}} = 2, R_{\text{root-}\geq 2} = 1$$

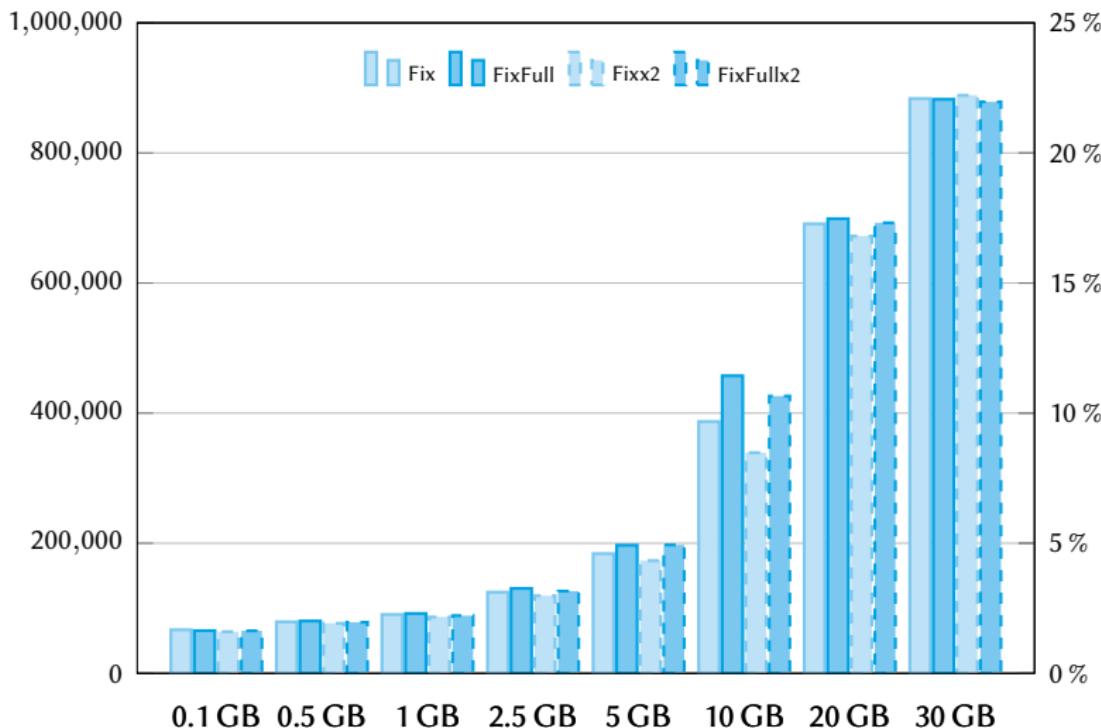
$$\times 2 \quad F_{\text{non-B-tree}} = 50, F_{\text{root}} = 50, F_{\text{root-1}} = 25, F_{\text{root-}\geq 2} = 10,$$

$$R_{\text{non-B-tree}} = 10, R_{\text{root}} = 10, R_{\text{root-1}} = 5, R_{\text{root-}\geq 2} = 2$$

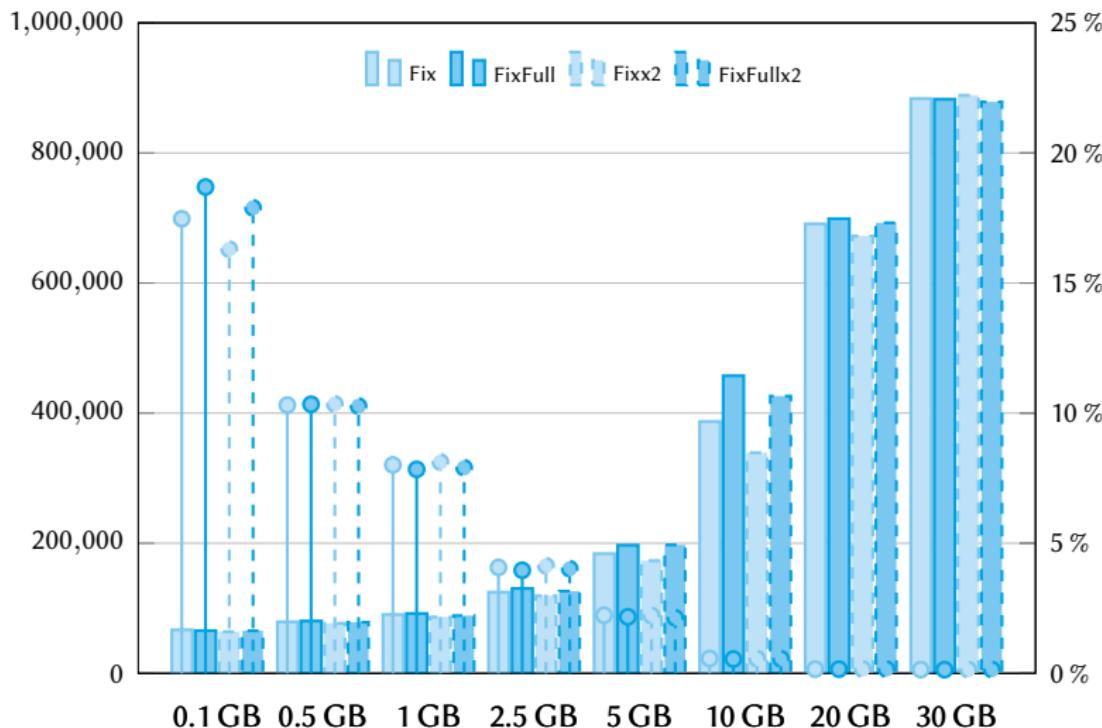
# Performance Evaluation



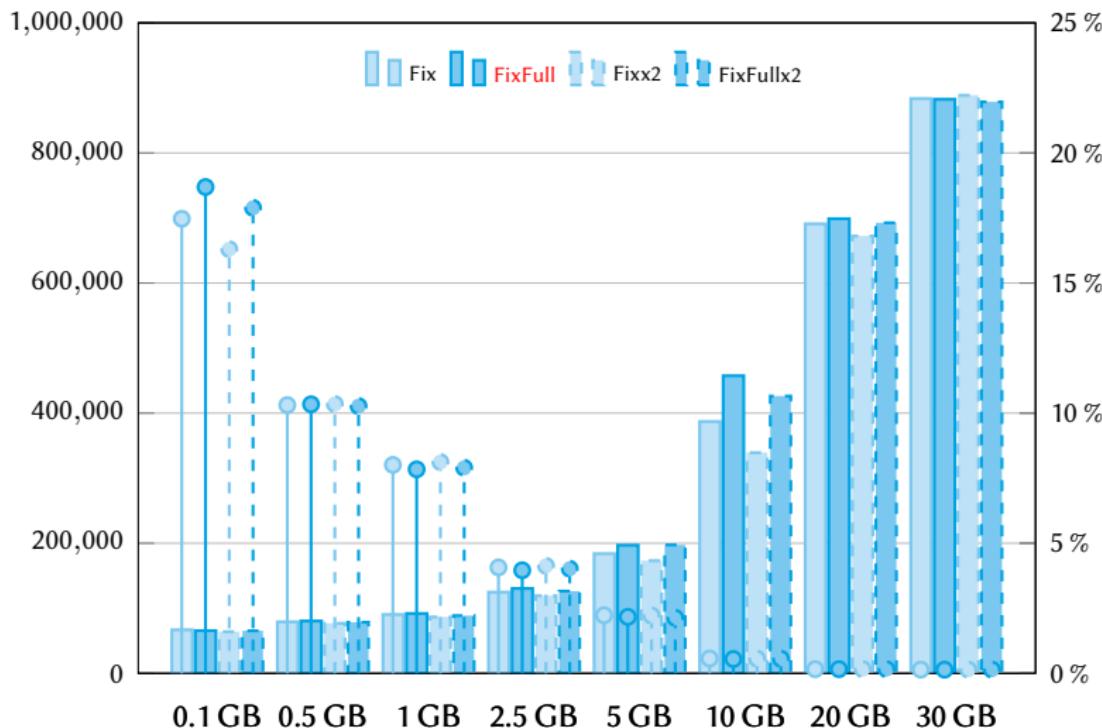
# Performance Evaluation



# Performance Evaluation



# Performance Evaluation



## DGCLOCK-V2 Variants

- Fix Set usage-count on *page fix*
  - FixFull Like Fix but also set usage-count when page cannot be evicted
- $F_{\text{non-B-tree}} = 25, F_{\text{root}} = 25, F_{\text{root-1}} = 10, F_{\text{root-}\geq 2} = 5,$   
 $R_{\text{non-B-tree}} = 5, R_{\text{root}} = 5, R_{\text{root-1}} = 2, R_{\text{root-}\geq 2} = 1$
- x2  $F_{\text{non-B-tree}} = 50, F_{\text{root}} = 50, F_{\text{root-1}} = 25, F_{\text{root-}\geq 2} = 10,$   
 $R_{\text{non-B-tree}} = 10, R_{\text{root}} = 10, R_{\text{root-1}} = 5, R_{\text{root-}\geq 2} = 2$

## DGCLOCK-V2 Variants

### Usage-Count

$$\text{(initially)} \quad UC(p) = F_p$$

$$\text{(set)} \quad UC(p) = R_p$$

$$\text{(decrease)} \quad UC(p) = UC(p) - 1$$

Fix Set usage-count on *page fix*

FixFull Like Fix but also set usage-count when page cannot be evicted

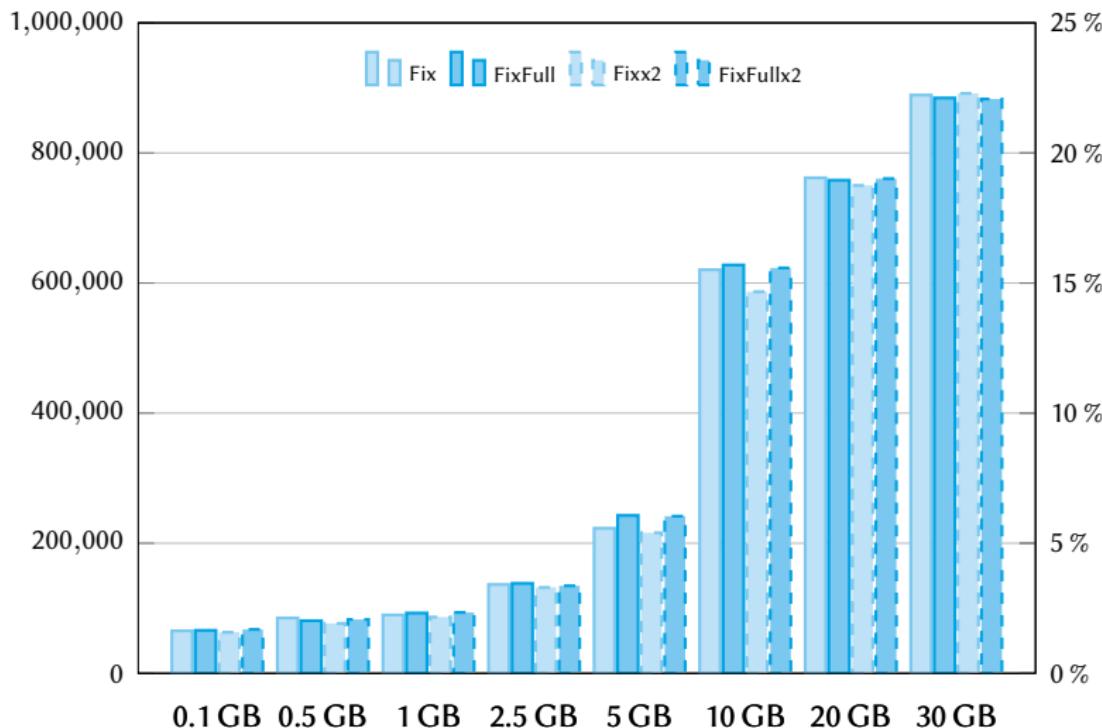
$$F_{\text{non-B-tree}} = 25, F_{\text{root}} = 25, F_{\text{root-1}} = 10, F_{\text{root-}\geq 2} = 5,$$

$$R_{\text{non-B-tree}} = 5, R_{\text{root}} = 5, R_{\text{root-1}} = 2, R_{\text{root-}\geq 2} = 1$$

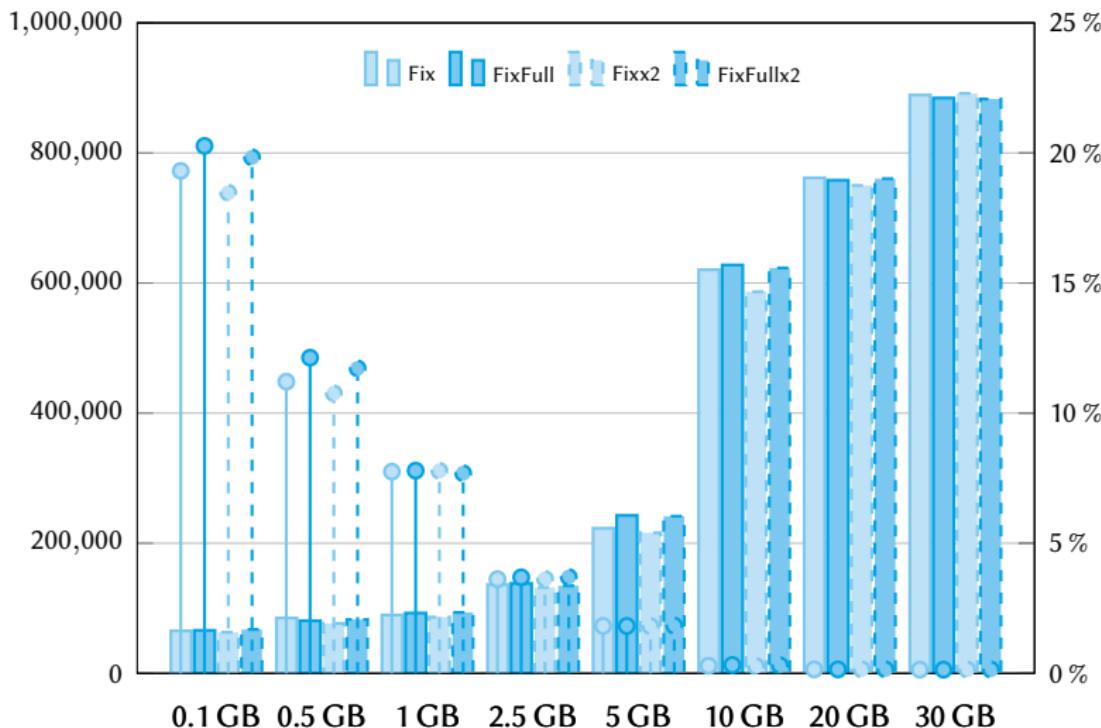
$$\times 2 \quad F_{\text{non-B-tree}} = 50, F_{\text{root}} = 50, F_{\text{root-1}} = 25, F_{\text{root-}\geq 2} = 10,$$

$$R_{\text{non-B-tree}} = 10, R_{\text{root}} = 10, R_{\text{root-1}} = 5, R_{\text{root-}\geq 2} = 2$$

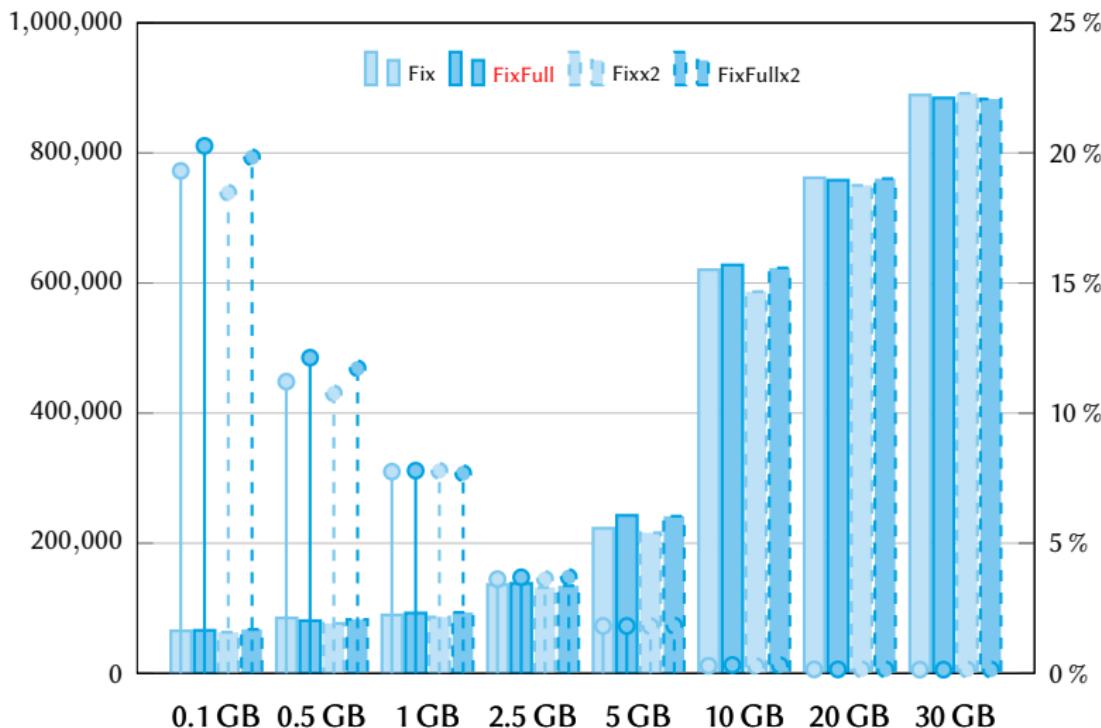
# Performance Evaluation



# Performance Evaluation



# Performance Evaluation



## LRD-V2 Variants

### Reference Counts

$$\begin{array}{ll} \text{(aging)} & UC(p) = UC(p) \circ x \\ \text{(increase)} & RC(p) = RC(p) + 1 \end{array}$$

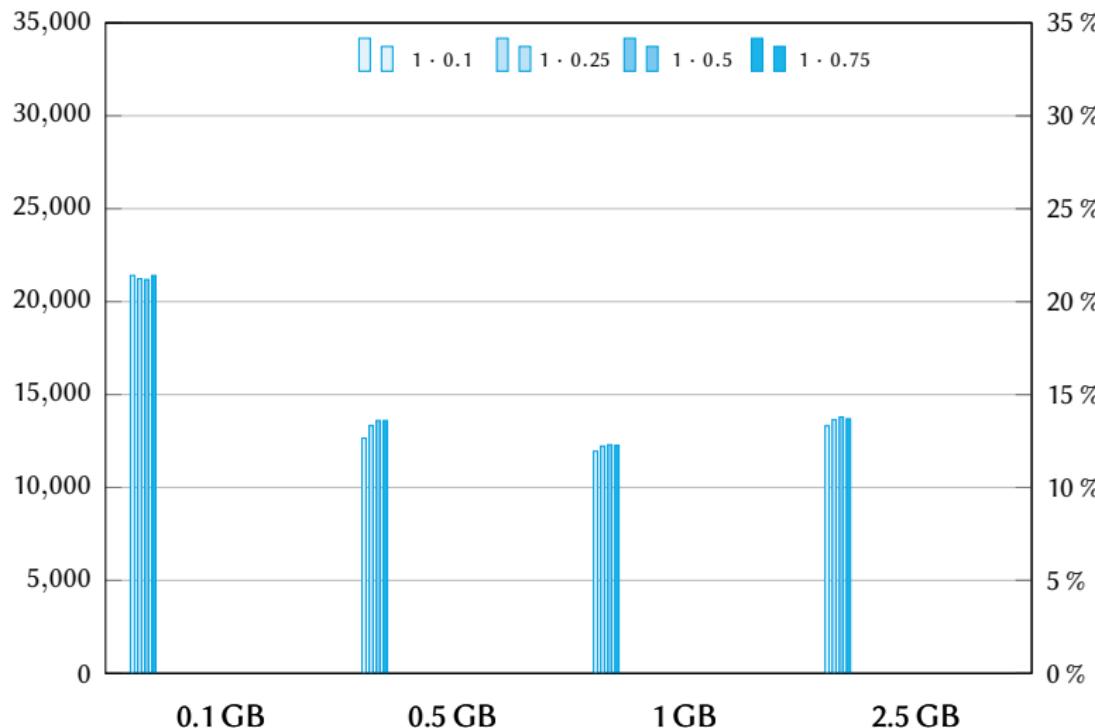
## LRD-V2 Variants

### Reference Counts

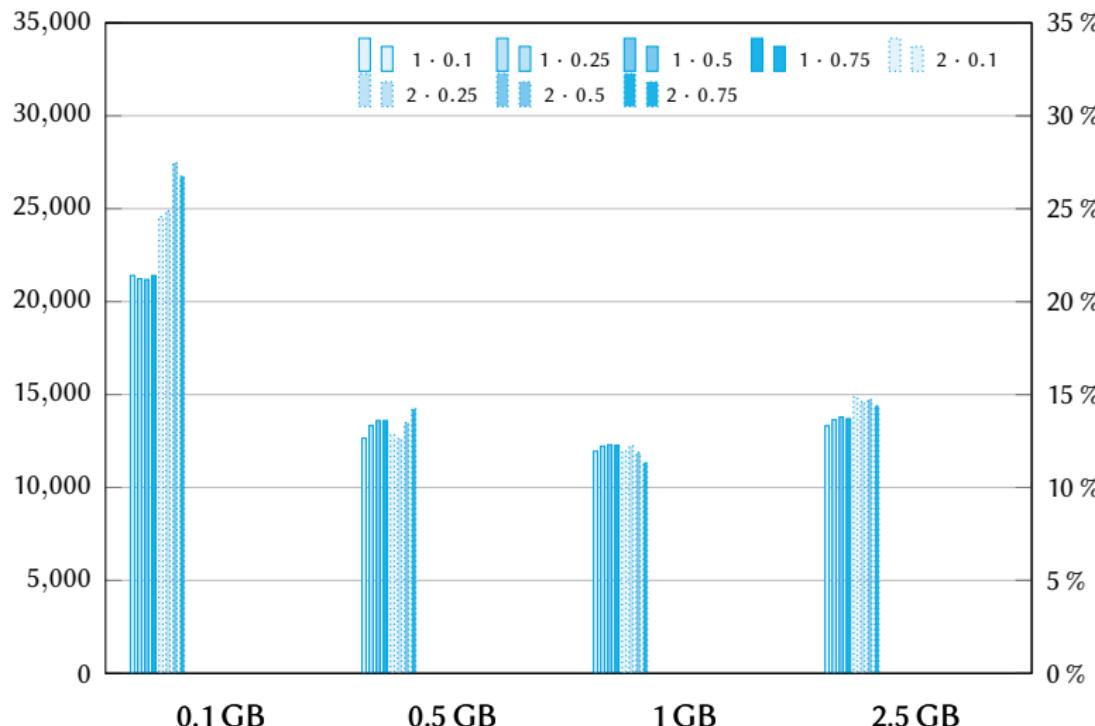
$$\begin{array}{ll} \text{(aging)} & UC(p) = UC(p) \circ x \\ \text{(increase)} & RC(p) = RC(p) + 1 \end{array}$$

$f \circ x$  Perform aging (using  $\circ x$ ) on each reference count  
every  $f \cdot \text{bufferPoolSize}$  global page references

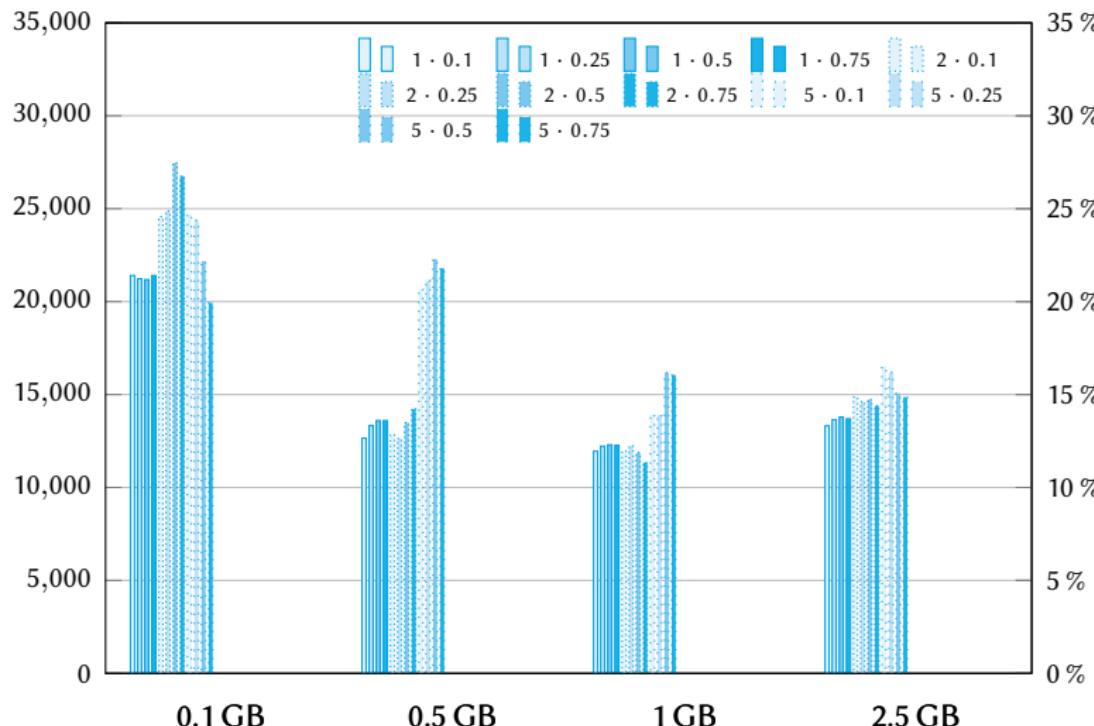
# Performance Evaluation I



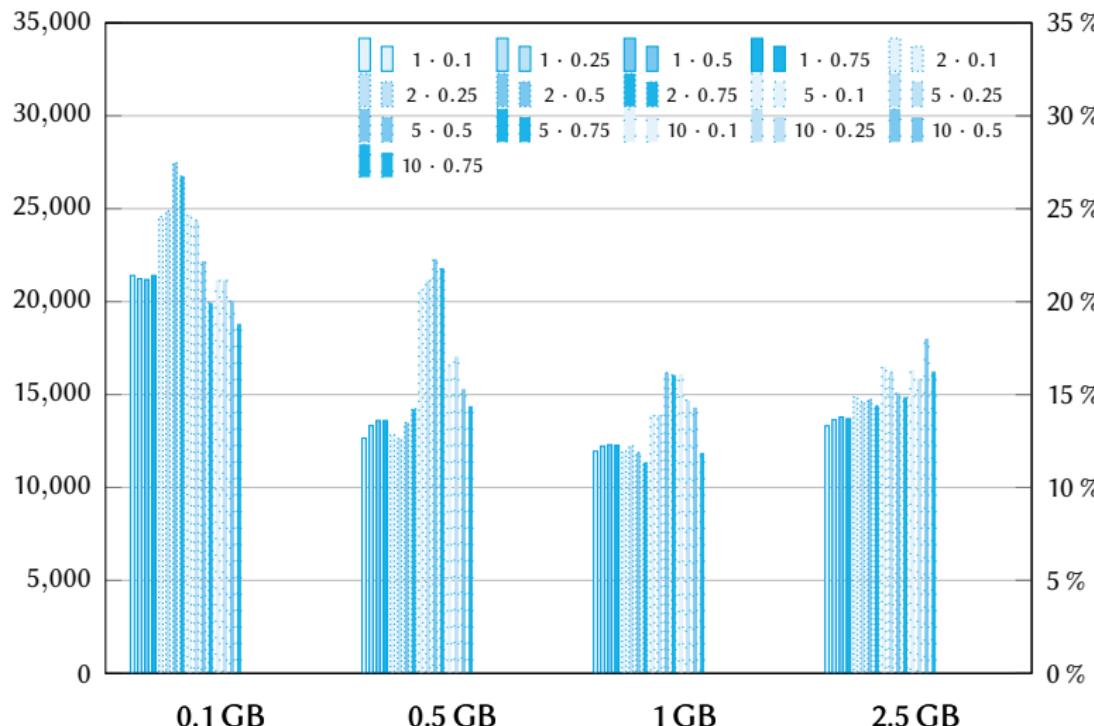
# Performance Evaluation I



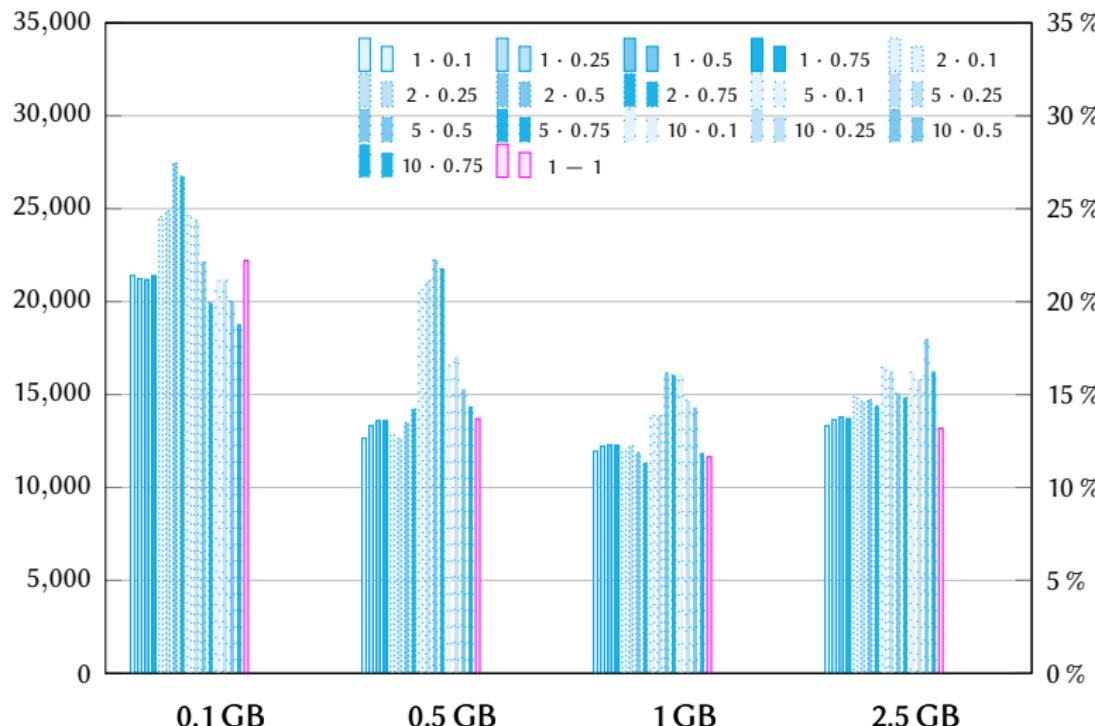
# Performance Evaluation I



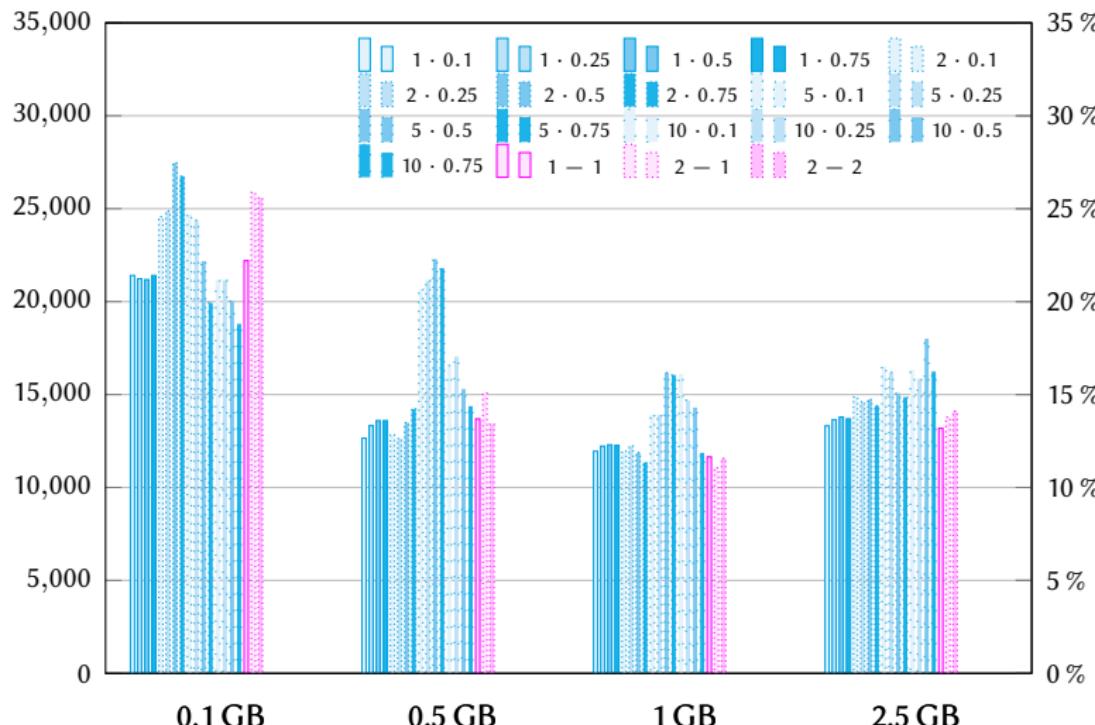
# Performance Evaluation I



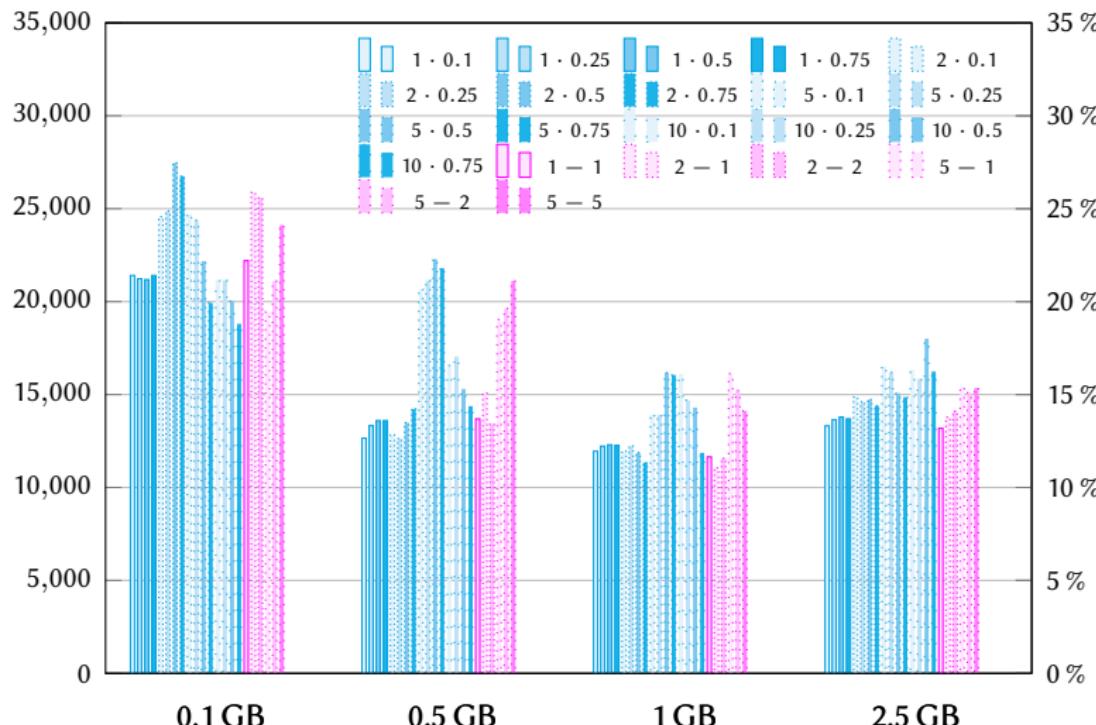
# Performance Evaluation I



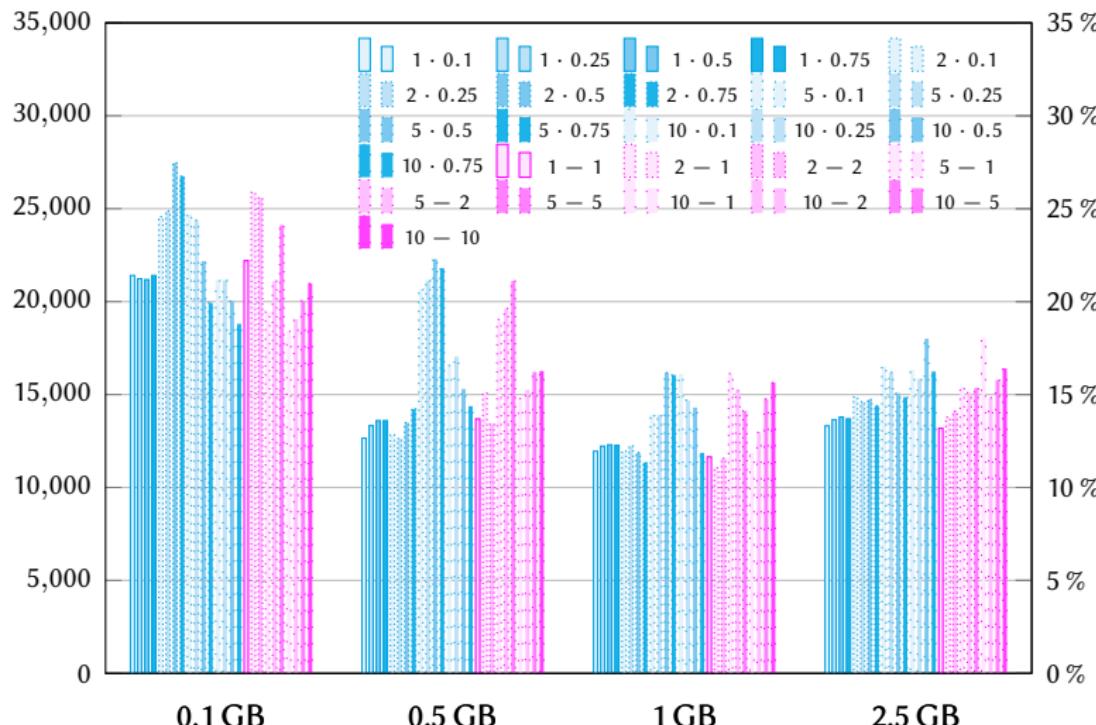
# Performance Evaluation I



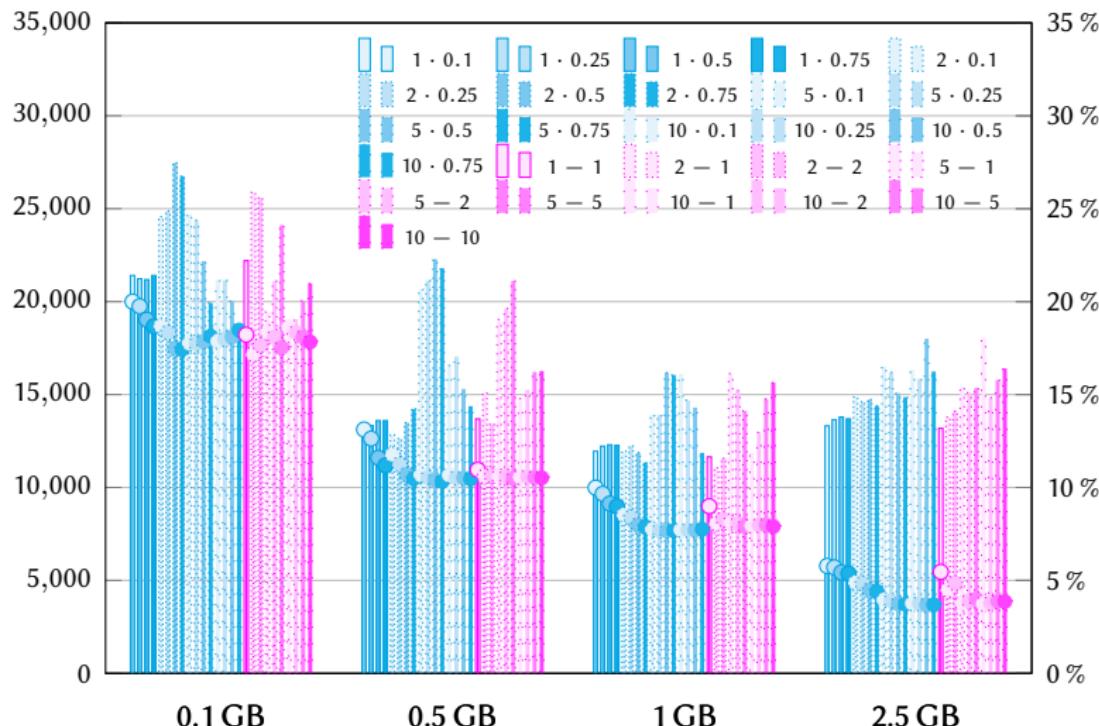
# Performance Evaluation I



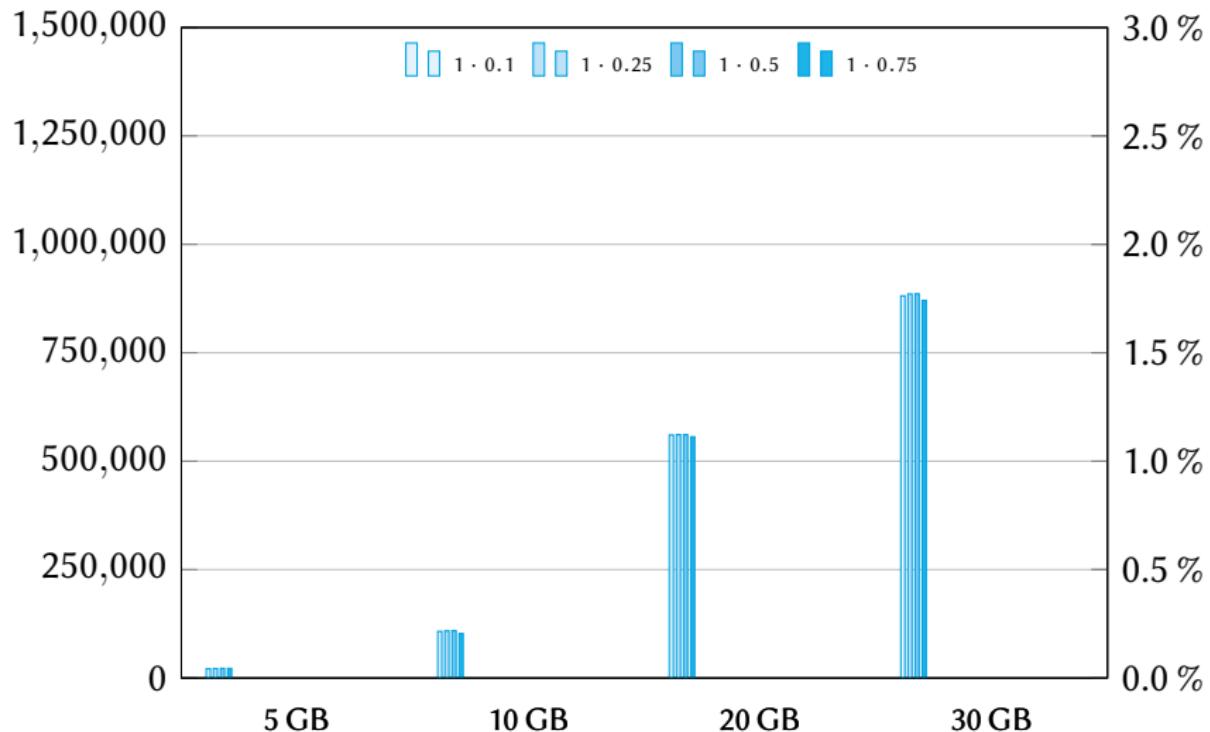
# Performance Evaluation I



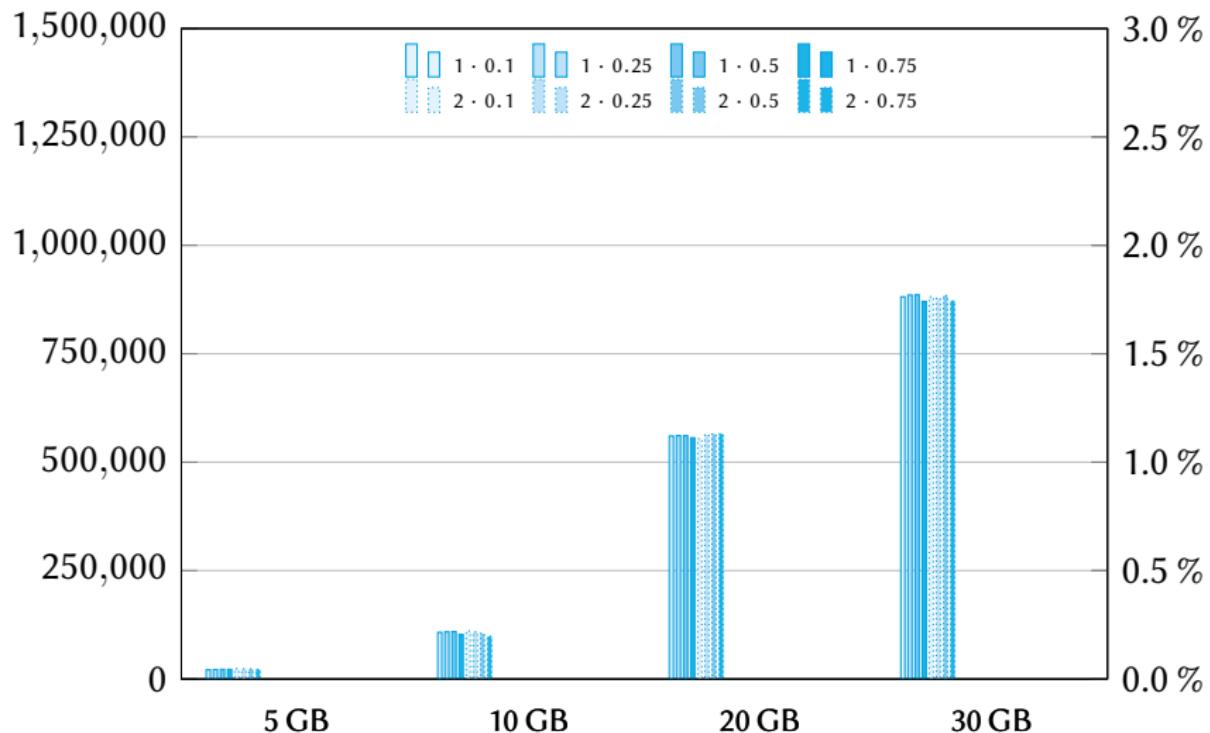
# Performance Evaluation I



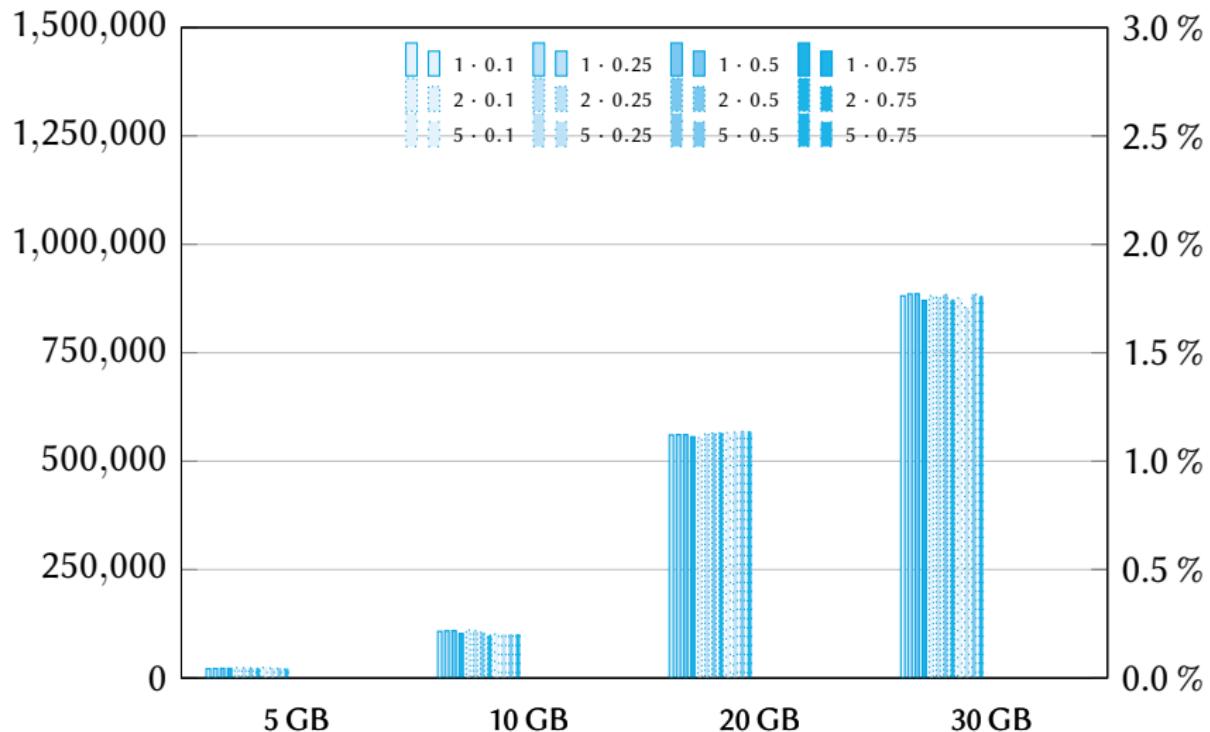
## Performance Evaluation II



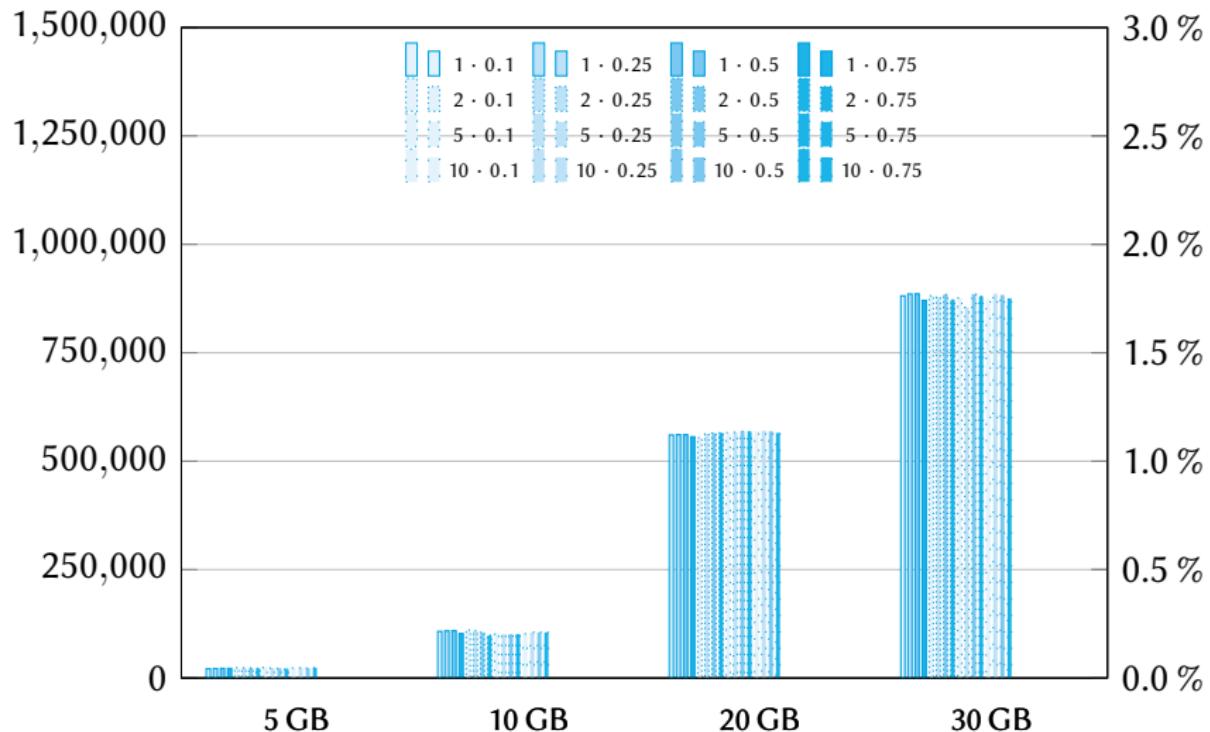
## Performance Evaluation II



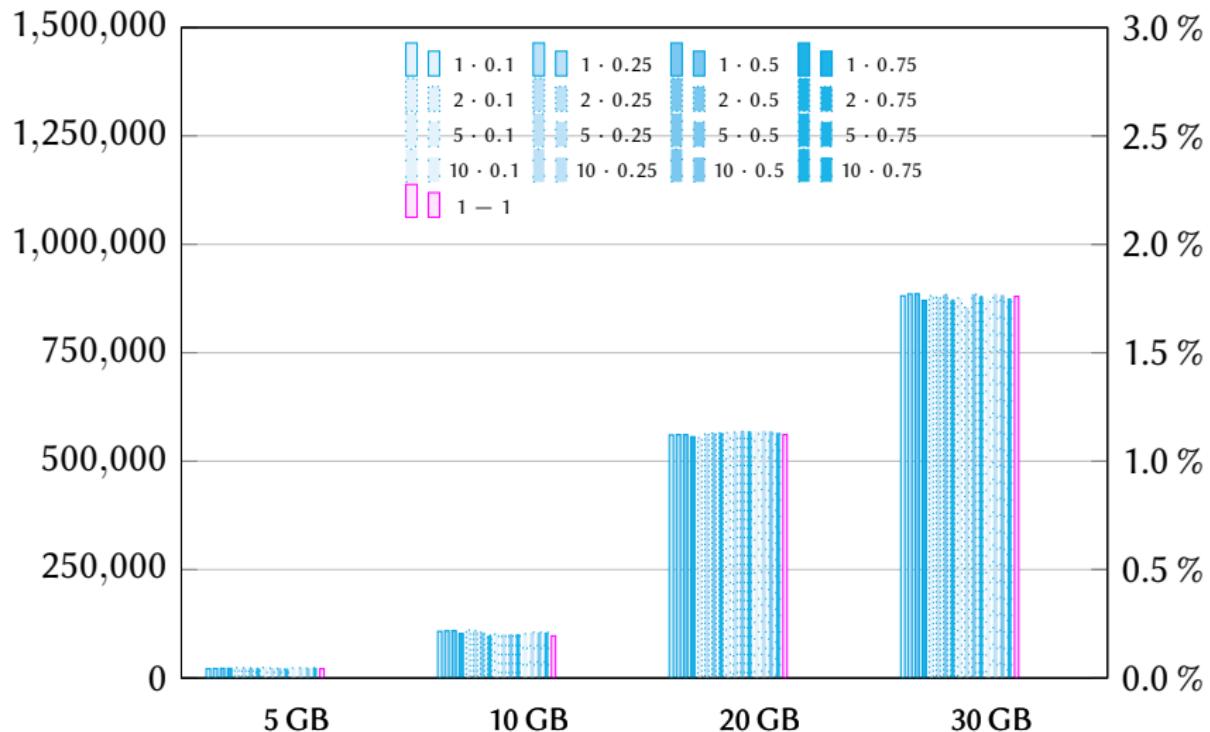
## Performance Evaluation II



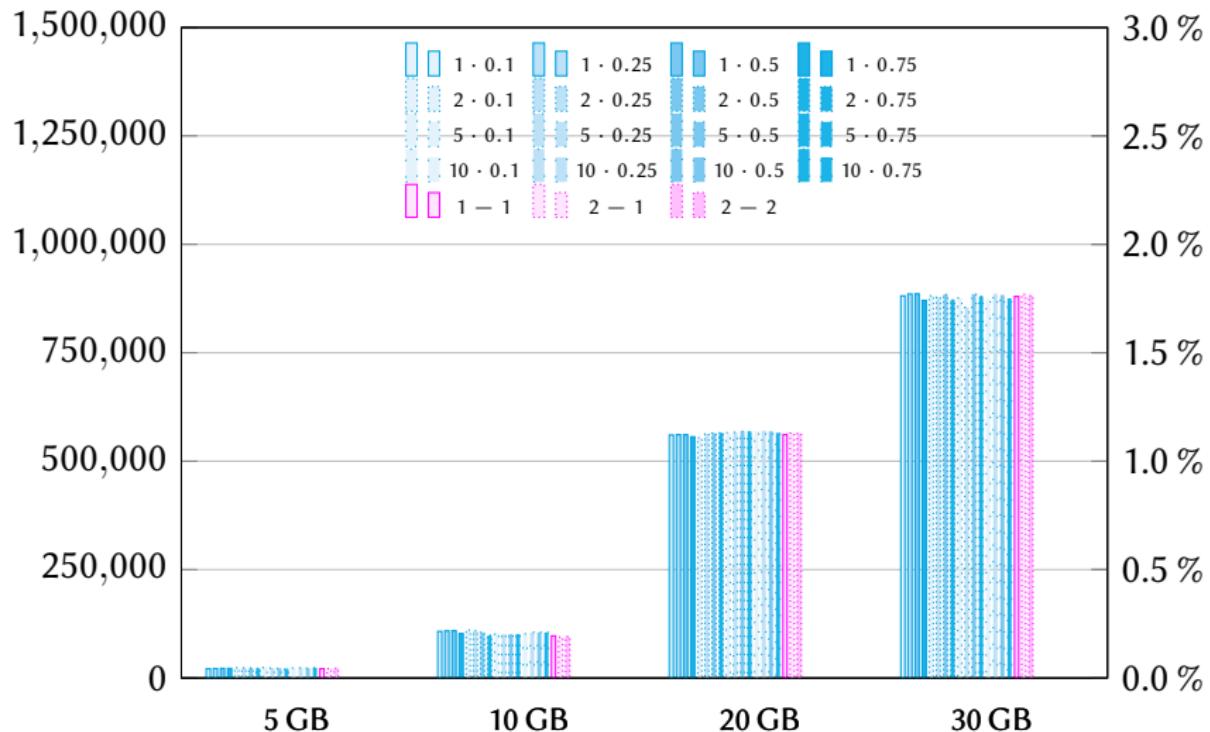
## Performance Evaluation II



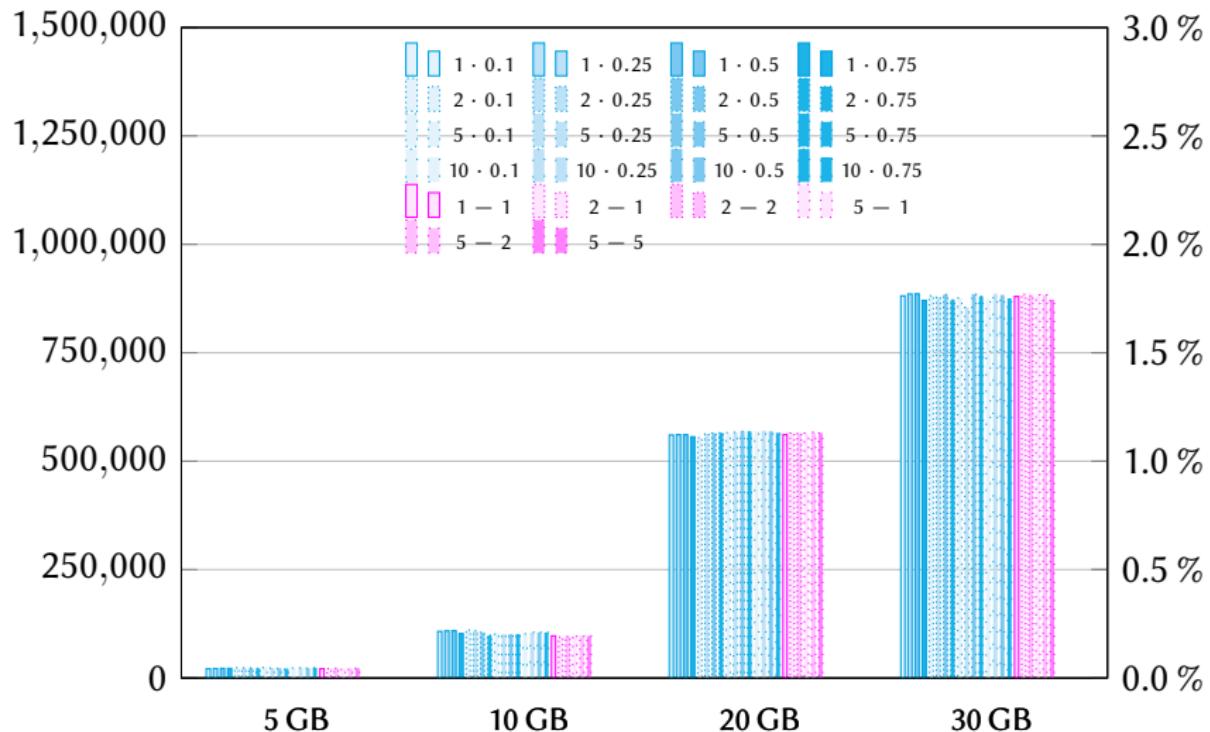
## Performance Evaluation II



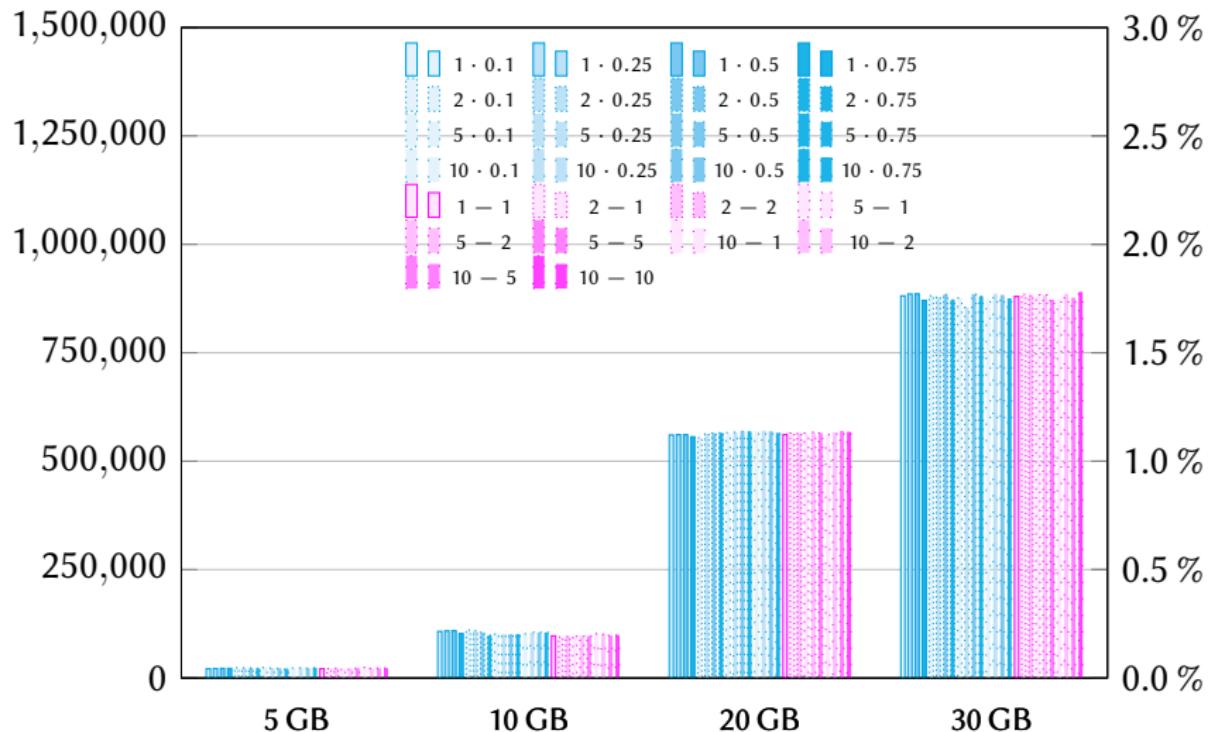
## Performance Evaluation II



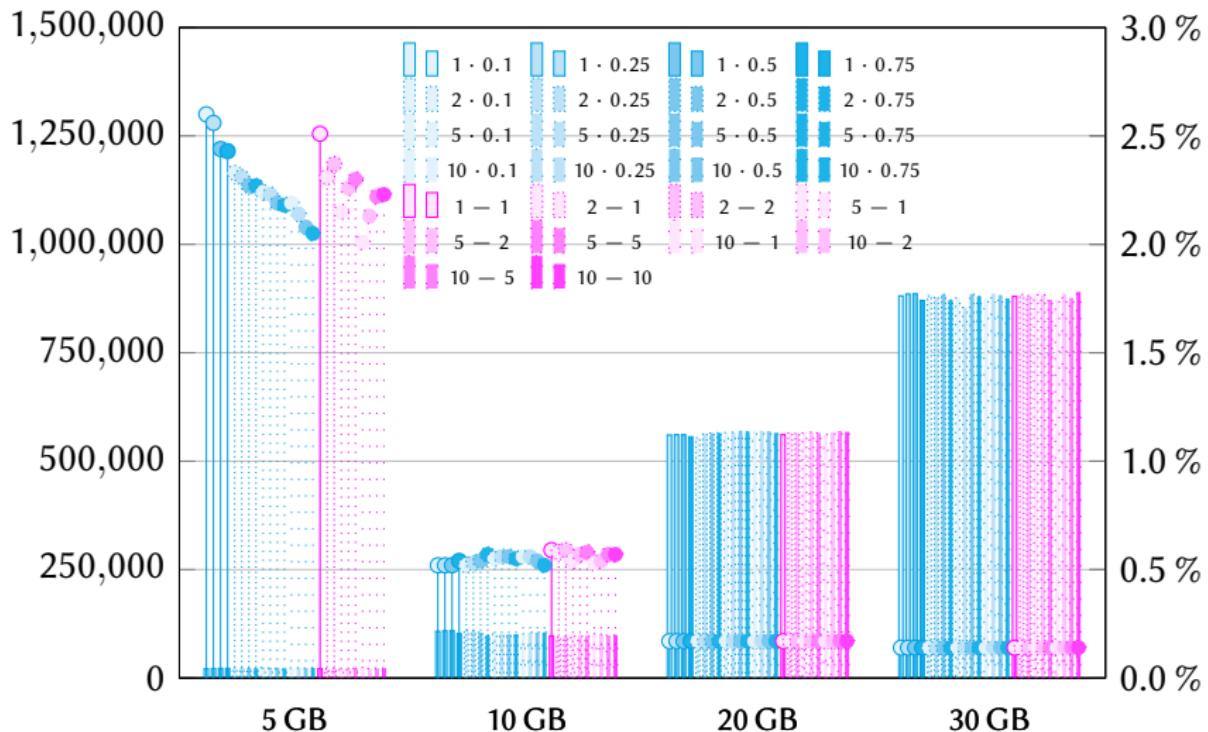
## Performance Evaluation II



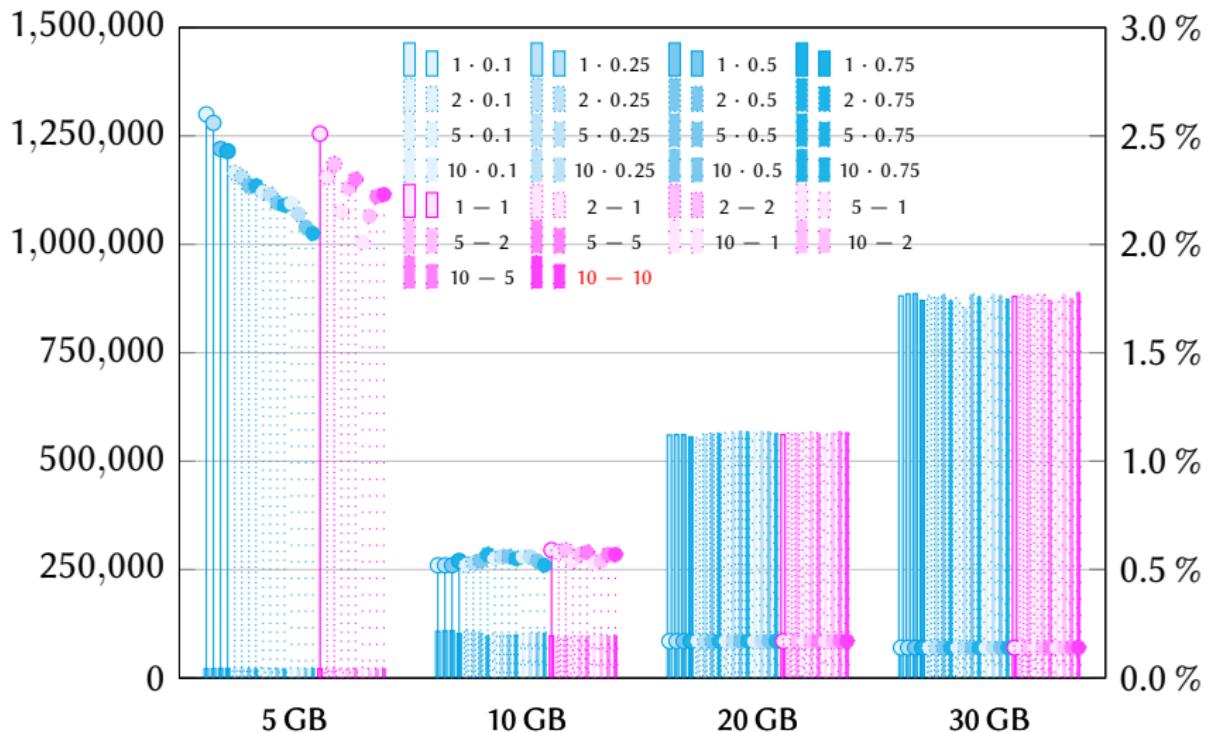
## Performance Evaluation II



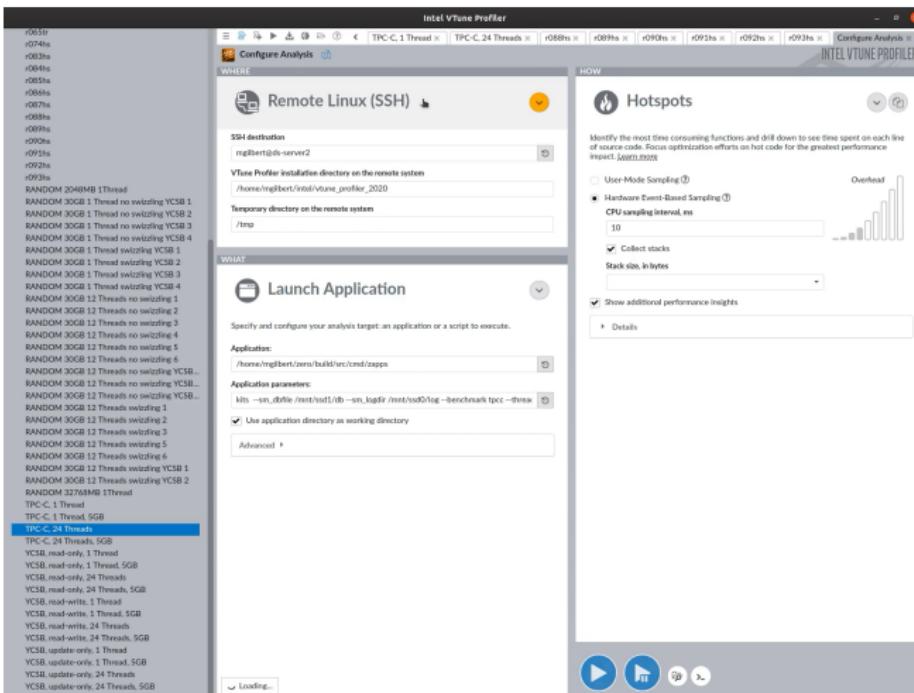
## Performance Evaluation II



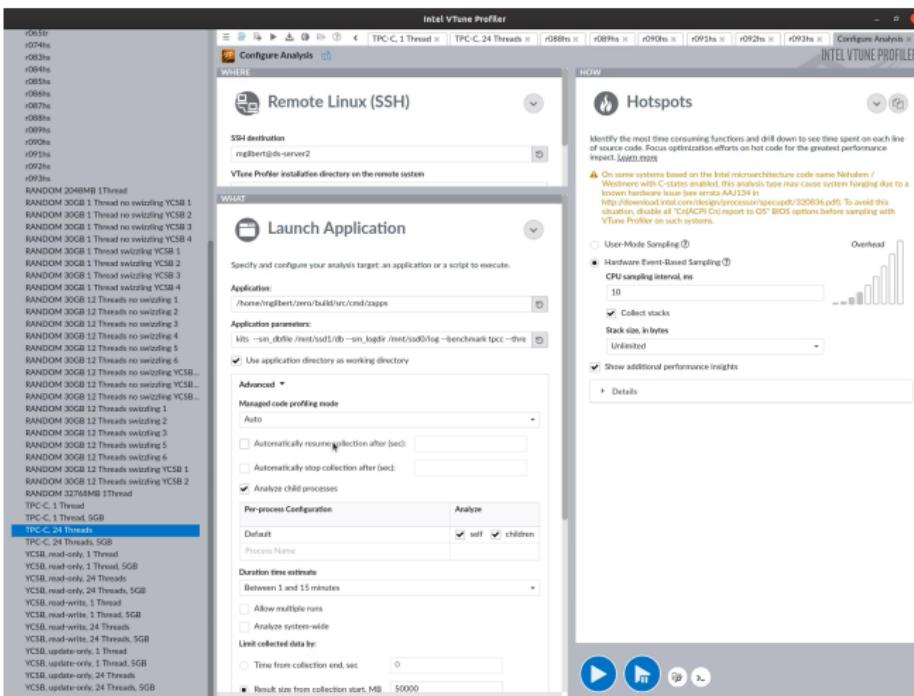
## Performance Evaluation II



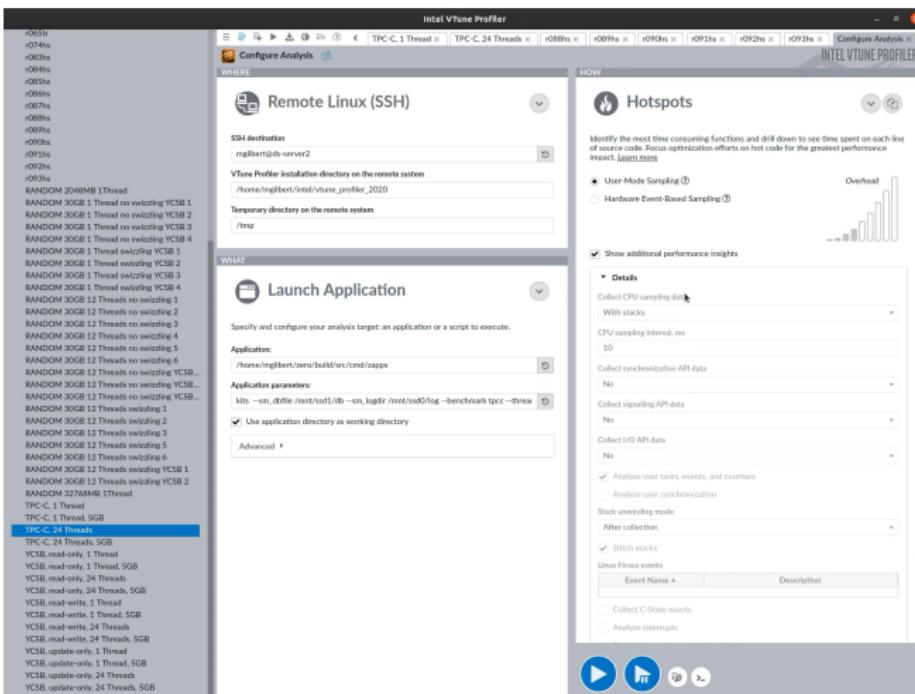
# Analysis Configuration: Where?



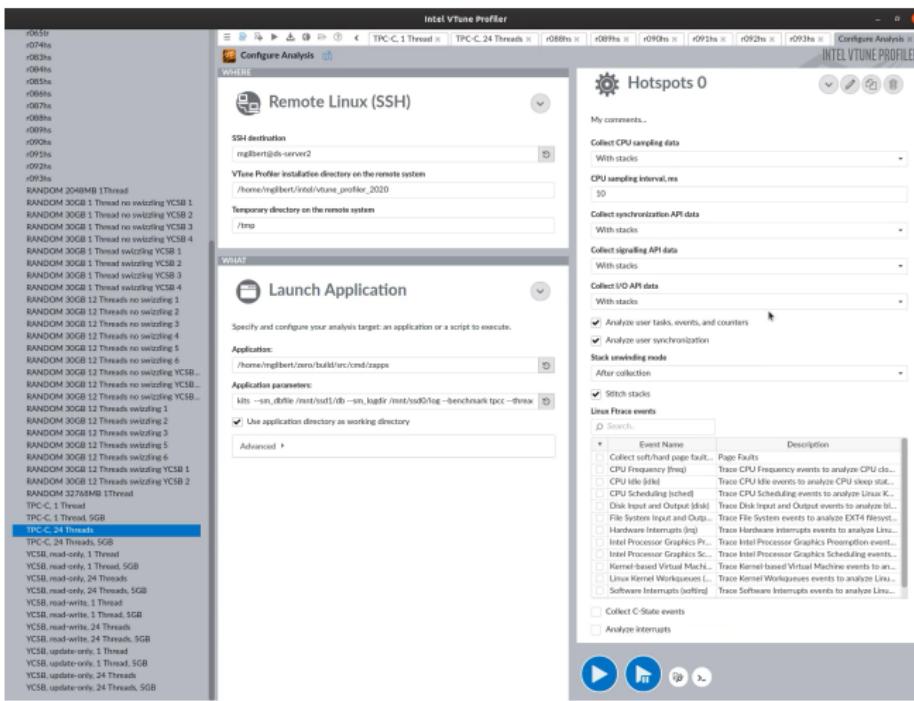
# Analysis Configuration: What?



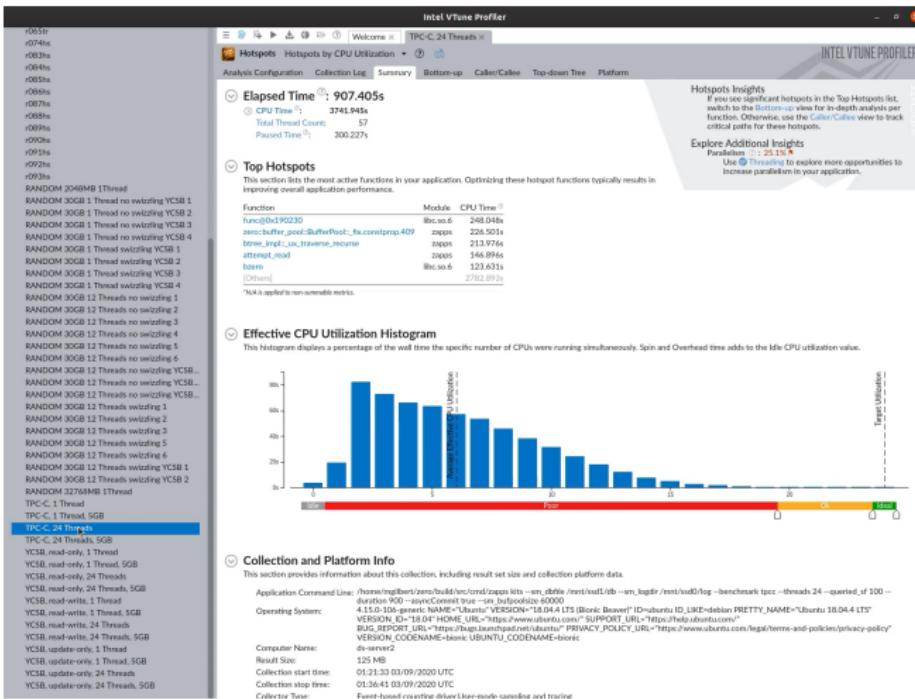
# Analysis Configuration: How?



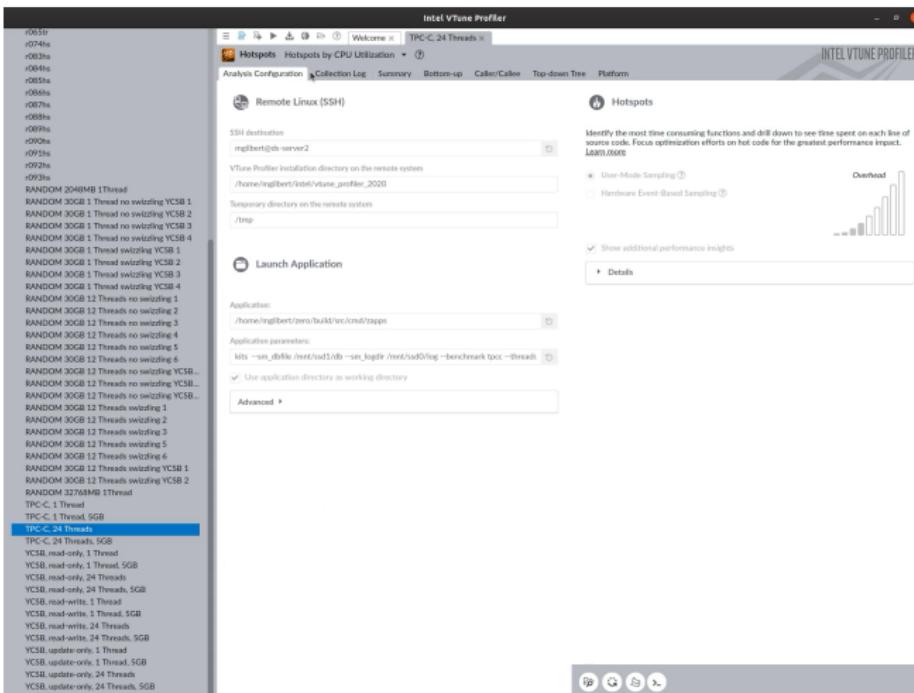
# Analysis Configuration: How??



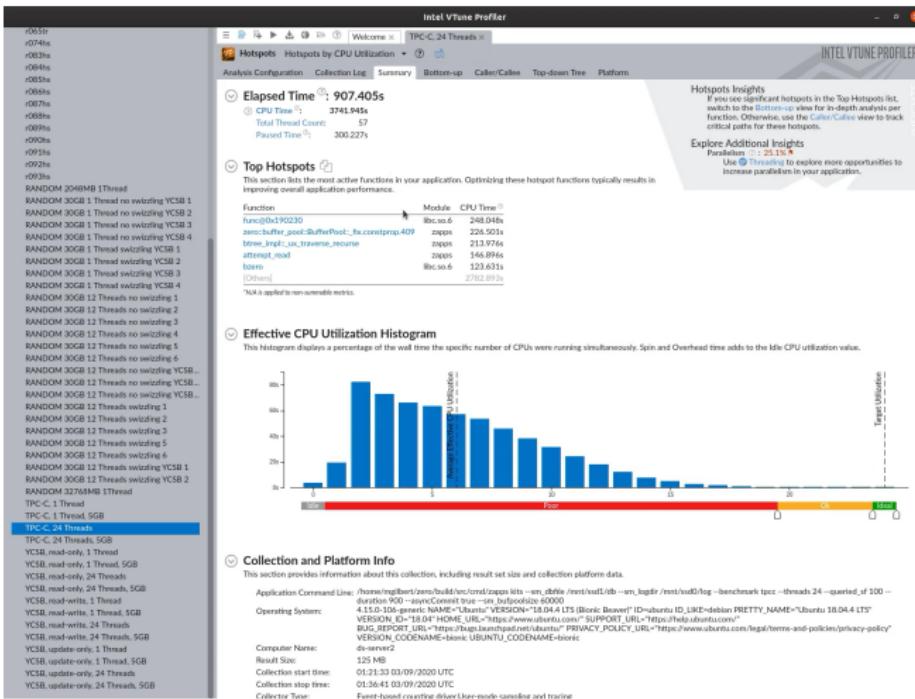
# Analysis: Open



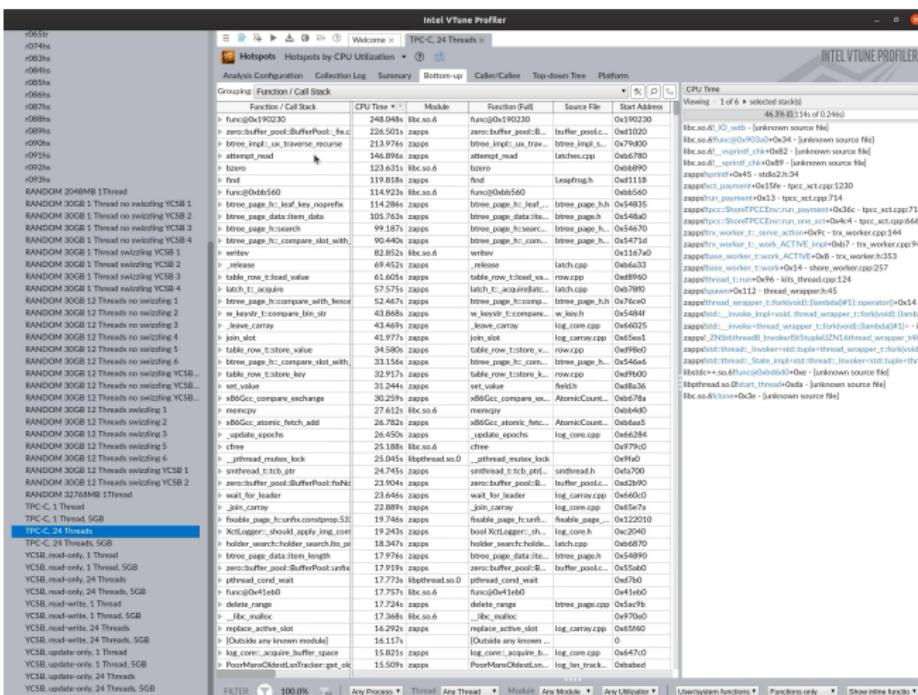
# Analysis: Configuration & Collection Log



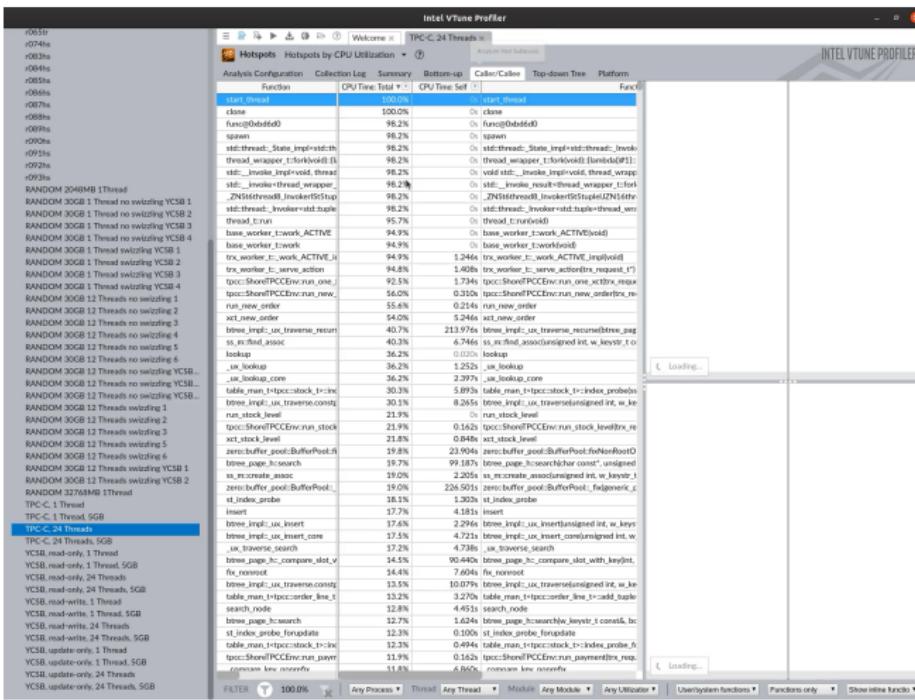
# Analysis: Summary



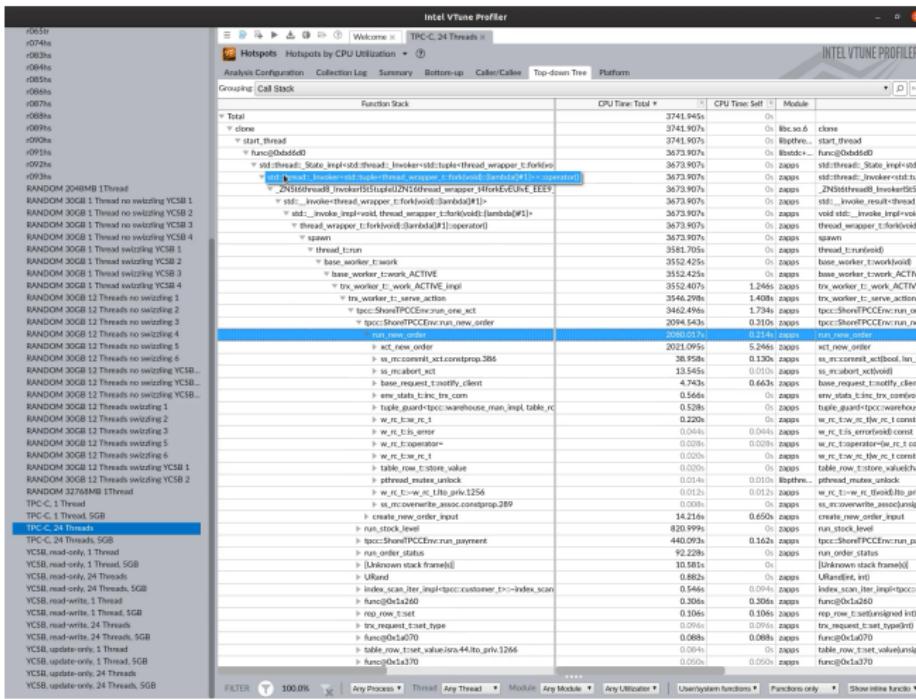
# Analysis: Bottom-Up



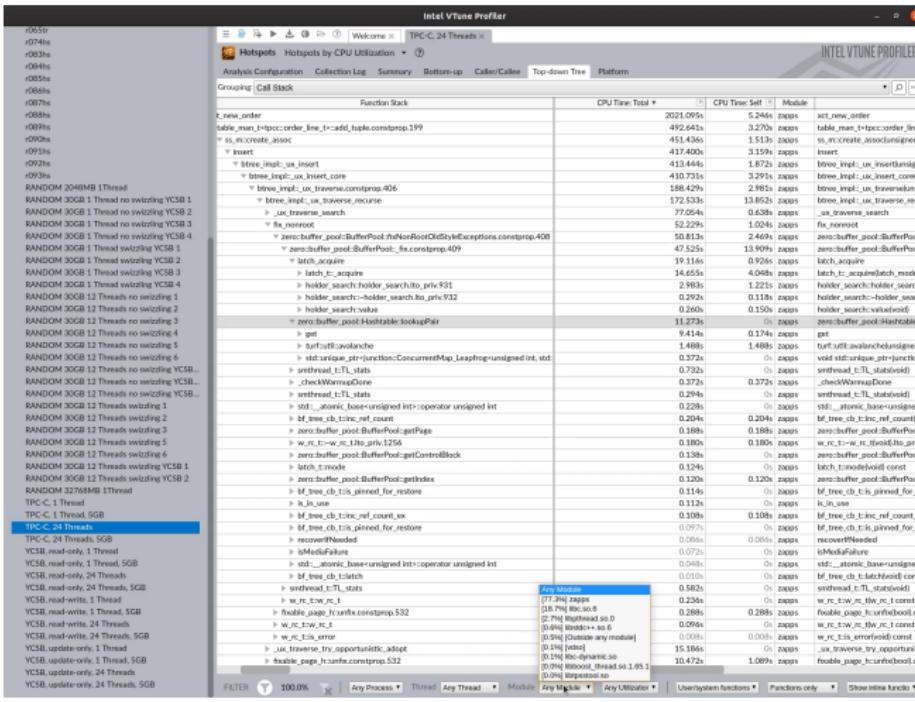
## Analysis: Caller/Callee



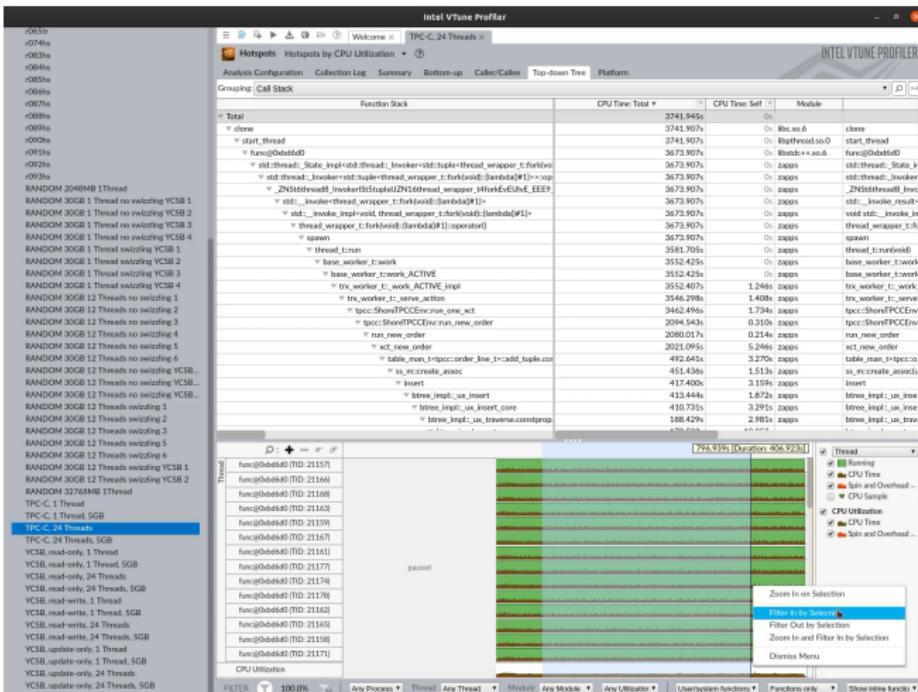
# Analysis: Top-Down Tree



## Analysis: Filter



## Analysis: Filter in Time



# Analysis: Platform

